Introduction to Machine Learning with H_2O

- Introduction (30 mins)
  - From Engineering to Data Science – Why?
  - Machine Learning Basics
  - H_2O Machine Learning Platform
- Tutorial (60 mins)
  - Regression (House Price)
  - Classification (Human Activities)
  - Clustering (Water Treatment Plant)
- Extra: More About H_2O (15 mins)
  - Use Cases
  - H_2O on a Multi-Node Cluster
  - Next-Gen H_2O
- Q & A (15 mins)

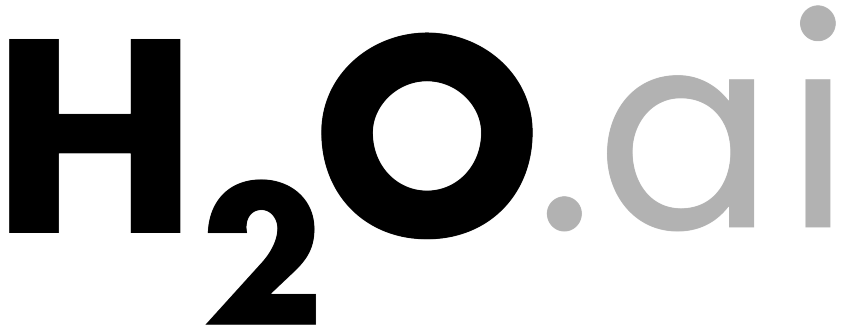

Jo-fai (Joe) Chow
Data Scientist at H_2O.ai
joe@h2o.ai

Download Link: http://bit.ly/h2o_exeter_2017

From Engineering to Data Science

- My Story

About Me

- Civil (Water) Engineer
  - 2005 – 2010
    - Consultant (SEAMS, UK)
      - Utilities / Asset Management
      - Constrained Optimization
  - 2010 – 2014
    - STREAM EngD (Exeter)
      - Infrastructure Design Optimisation
      - Machine Learning + Water Engineering
      - Discovered H_2O in 2014
- Data Scientist
  - 2015 – 2016
    - Virgin Media (UK)
    - Domino Data Lab (Silicon Valley)
  - 2016 – Present
    - H_2O.ai (Silicon Valley)
  - How?
    - bit.ly/joe_kaggle_story

About Me – I ❤ DataViz

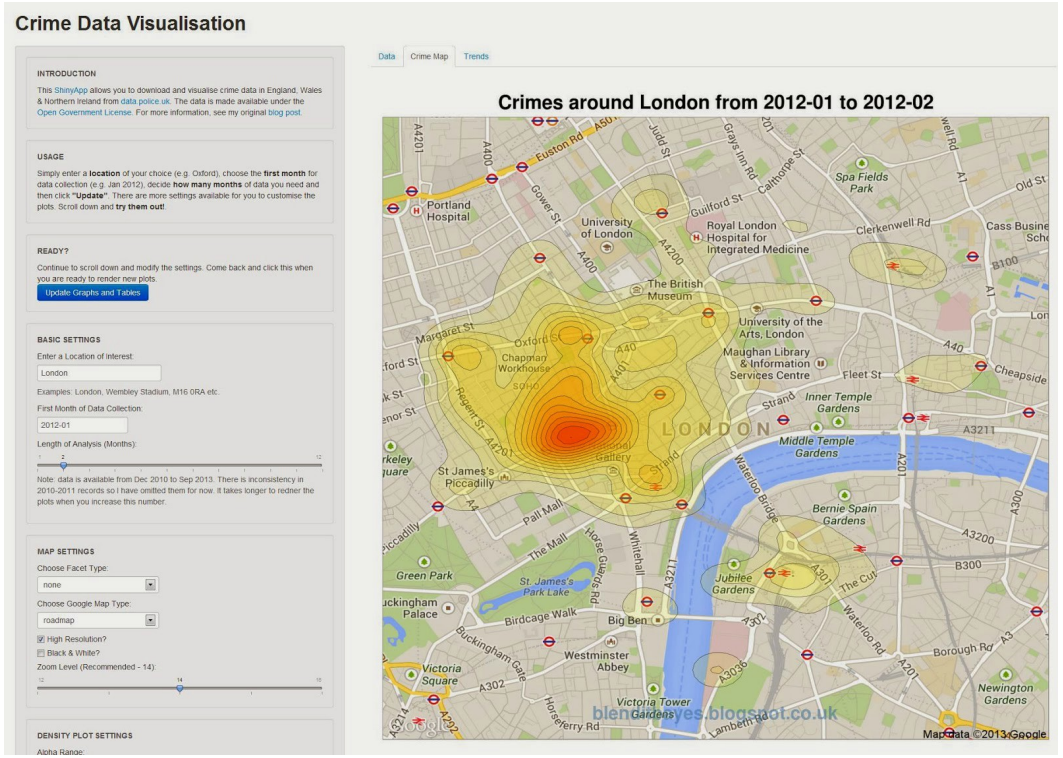

My First Data Viz & Shiny App Experience
CrimeMap (2013)

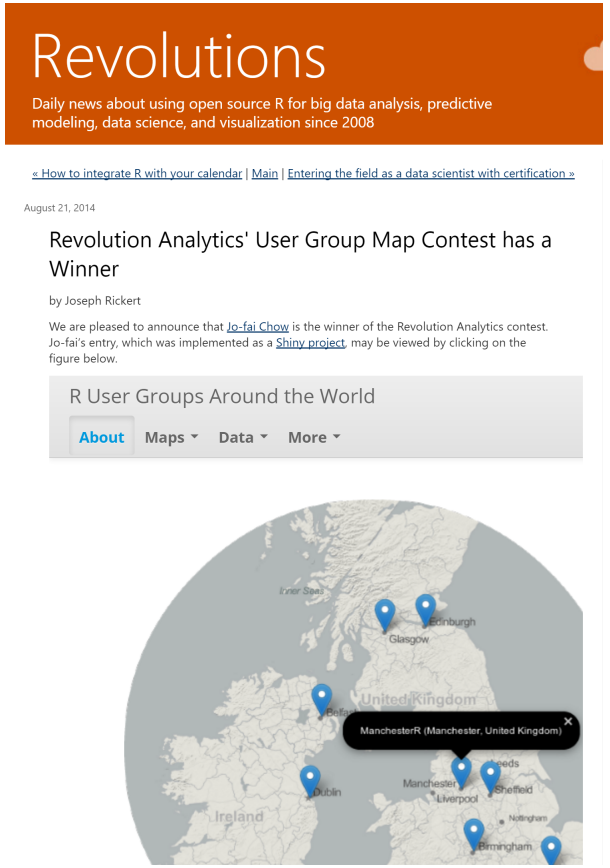

Revolution Analytics' Data Viz Contest
RUGSMAPS (2014)

H2O.ai

useR! 2014

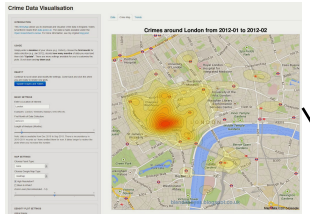

CrimeMap -> poster

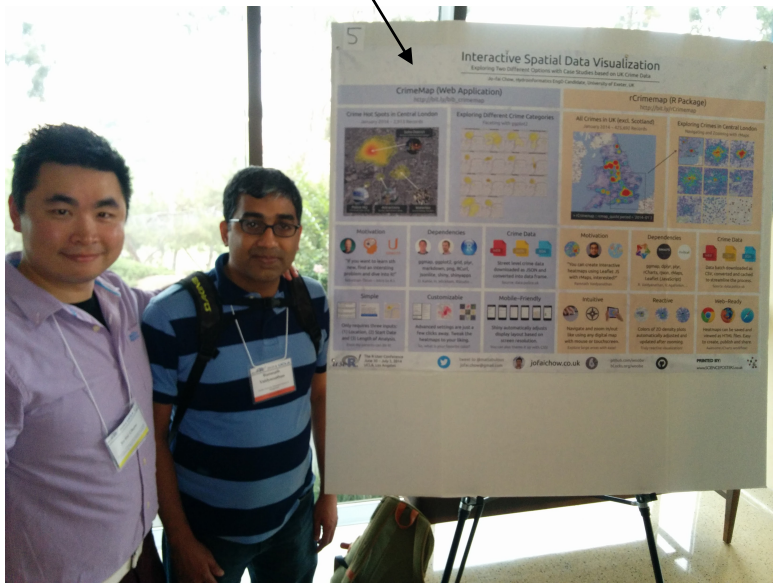

H2O.ai
@h2oai
Following

#linus for #rstat #user2014 #JohnChambers

H2O.ai
@h2oai
Following

#SparkSummit ~~> #useR2014 #h2o

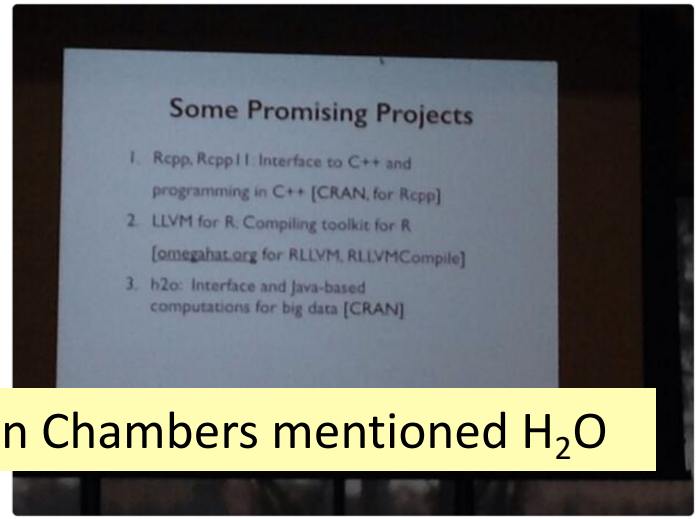

John Chambers mentioned H_2O

H2O.ai @h2oai · 28 Jul 2014
Smoooooth! - if I have to explain it in one word. Oxdata made this really easy for #R users. r-bloggers.com/things-to-try-… #Thanks #JoFaiChow

Things to try after useR! – Part 1: Deep Learning wit…
Annual R User Conference 2014The useR! 2014 conference was a mind-blowing experience. Hundreds of R enthusiasts and the beautiful UCLA campus, I am rea...
r-bloggers.com

1 5 6

Jo-fai (Joe) Chow
@matlabulous

Replying to @h2oai

Hi @srisatish @ArnoCandel and every1 @hexadata thx 4 making and open-sourcing the powerful #H2O shd hv tried it during (not after) #user2014

LIKES 2

1:41 PM - 28 Jul 2014

… so I gave it a try

H2O.ai

About Me – I ❤ Kaggle

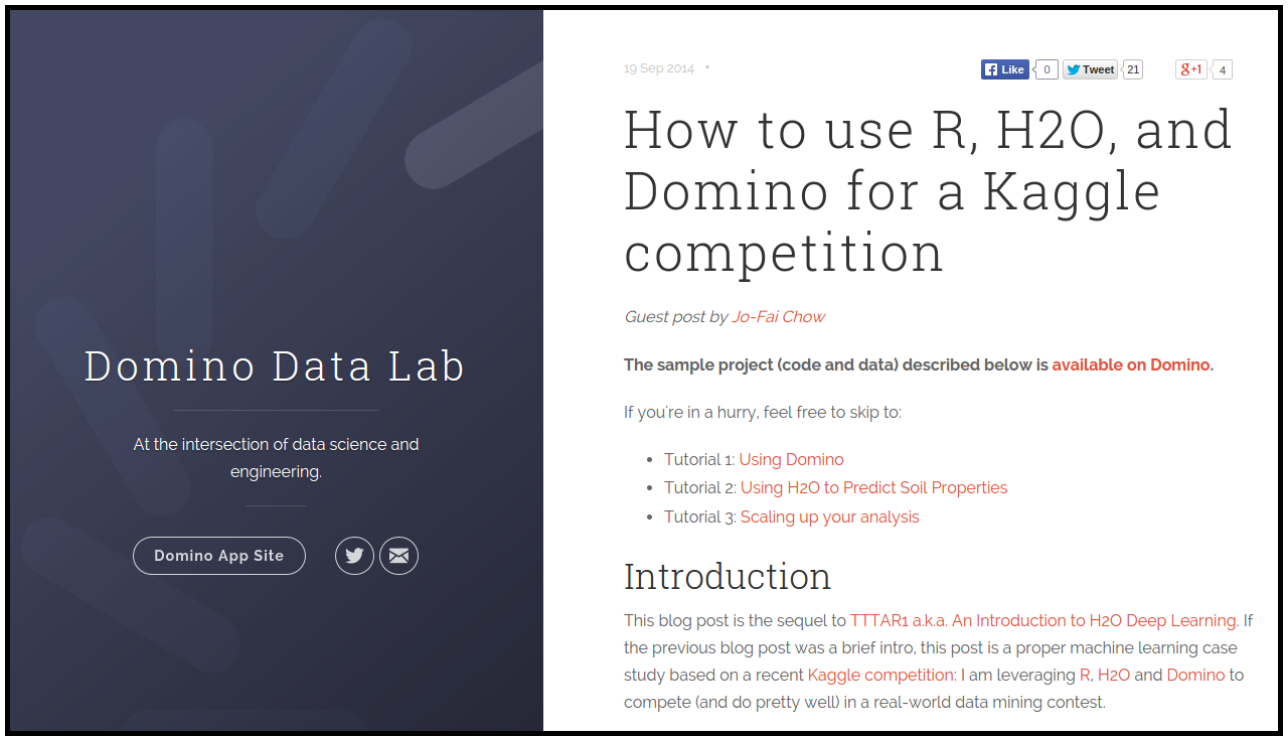

R + H_2O + Domino for Kaggle

Guest Blog Post for Domino & H_2O (2014)

- The Long Story
  - bit.ly/joe_kaggle_story

World Cup 2014

- Machine Learning vs My Friends
  - Team performance data from web
  - Simple machine learning models
  - Objective: Predict Correct Score
  - Me : 10 out of 64 (15.6%)
  - Friends' Avg : 4 out of 64 (6.3%)

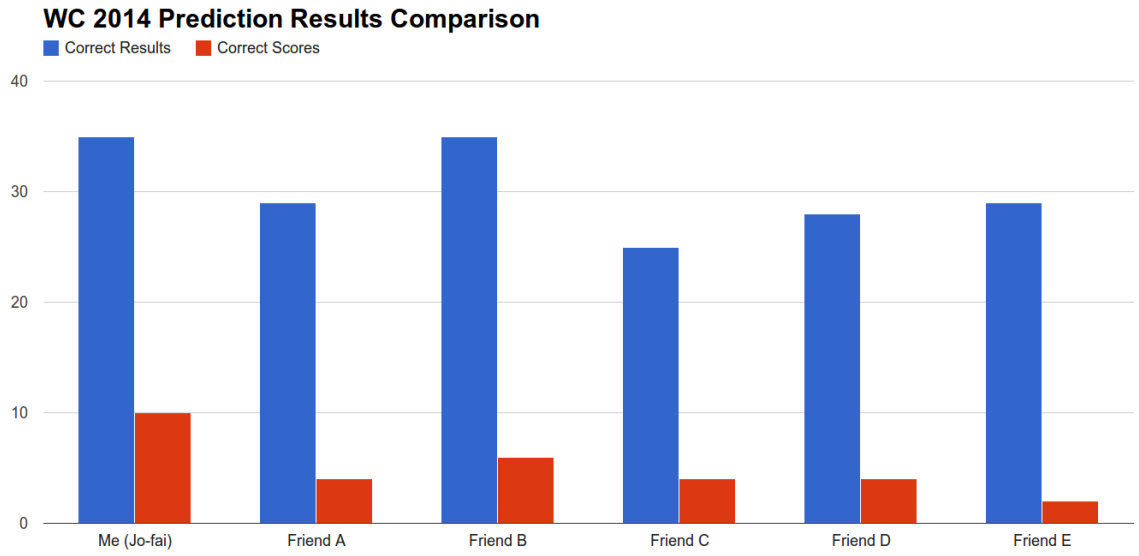

From Engineering to Data Science

- Knowledge Driven Decision Making
  - Write code to deal with data based on knowledge and assumptions.
  - Strong domain knowledge.

- Data Driven Decision Making
  - Define a problem and collect data related to the problem.
  - Write code to show data to computers and let computers (algorithms) find patterns in data.
  - Domain knowledge is helpful but not necessary.

Higgs Boson Machine Learning Challenge

Use the ATLAS experiment to identify the Higgs boson

$13,000 · 1,785 teams · 3 years ago

Overview | Data | Discussion | Leaderboard | Rules | Team | My Submissions | Late Submission

https://www.kaggle.com/c/higgs-boson

Overview

- Description
- Evaluation
- Prizes
- About The Sponsors
- Timeline
- Winners

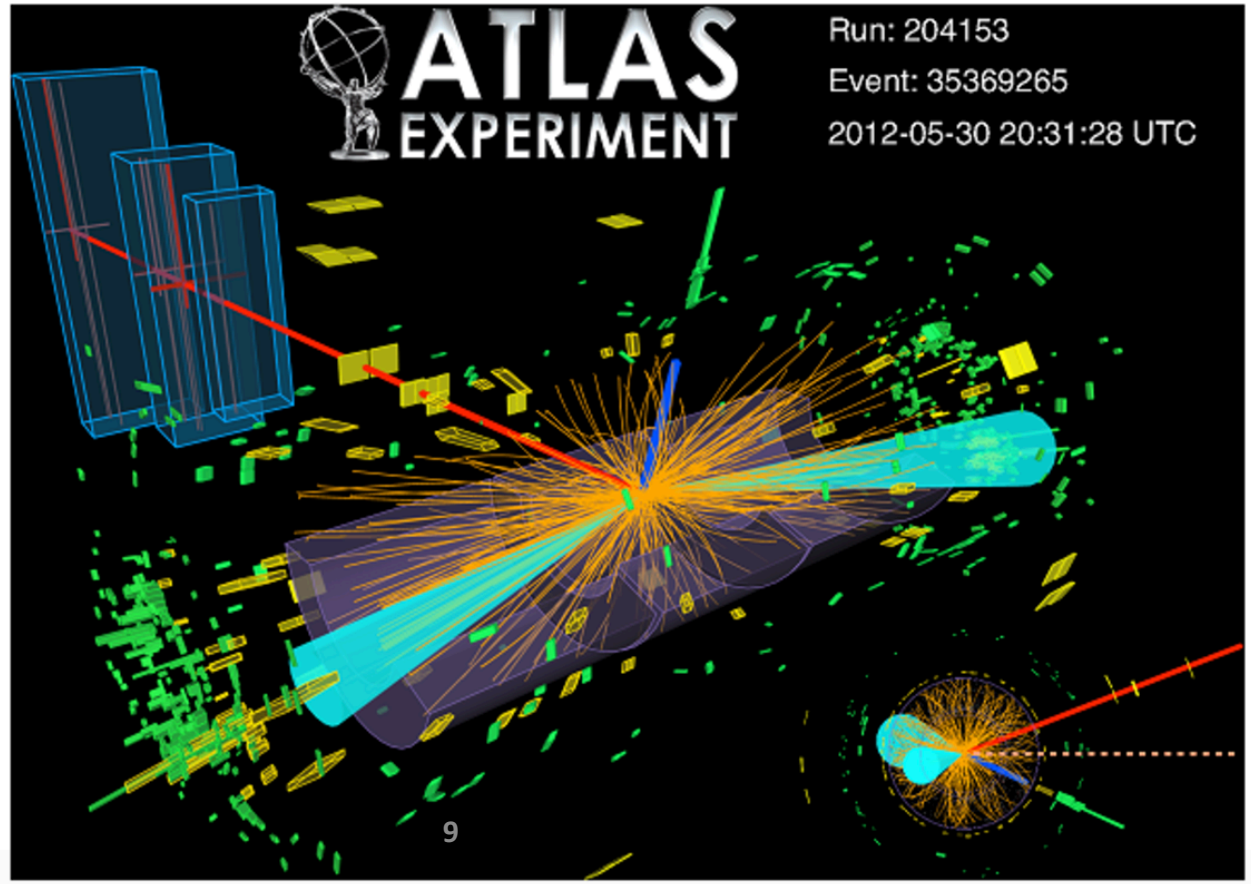

What is Machine Learning?

- Supervised Learning
- Unsupervised Learning

Supervised Learning

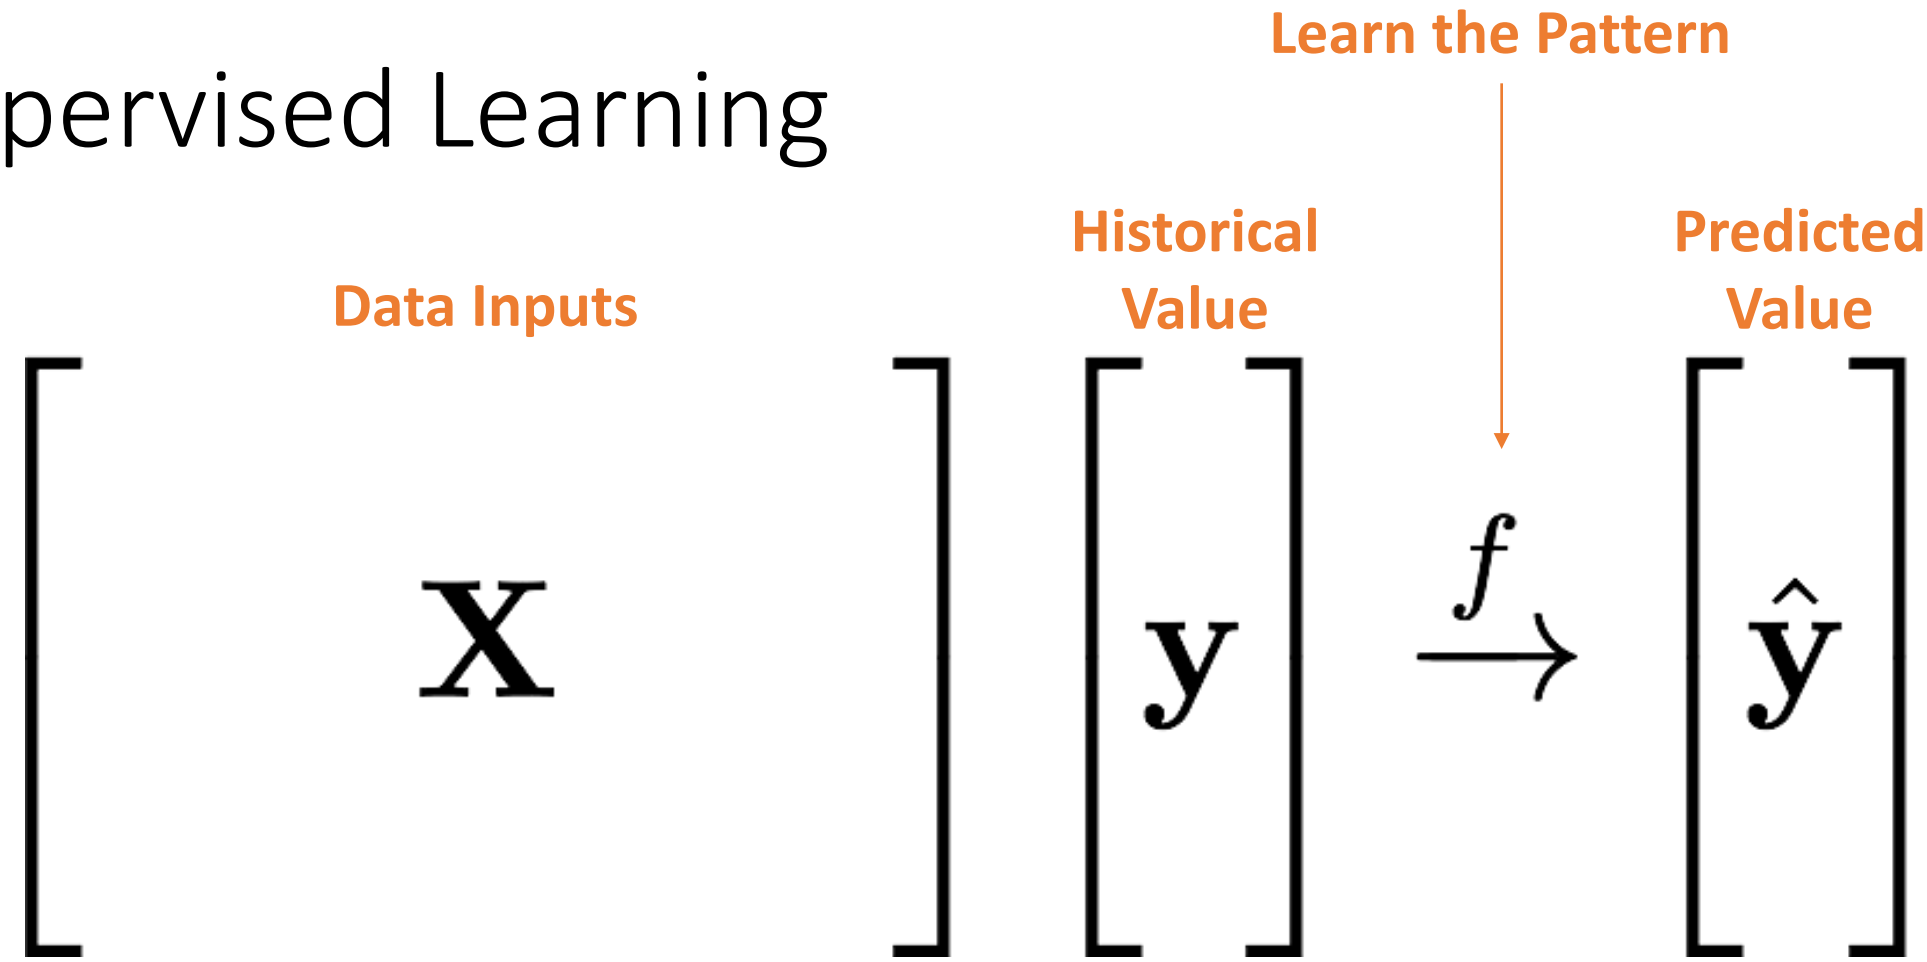

Supervised Learning Example

Machine Learning:
Learn Patterns
from Data

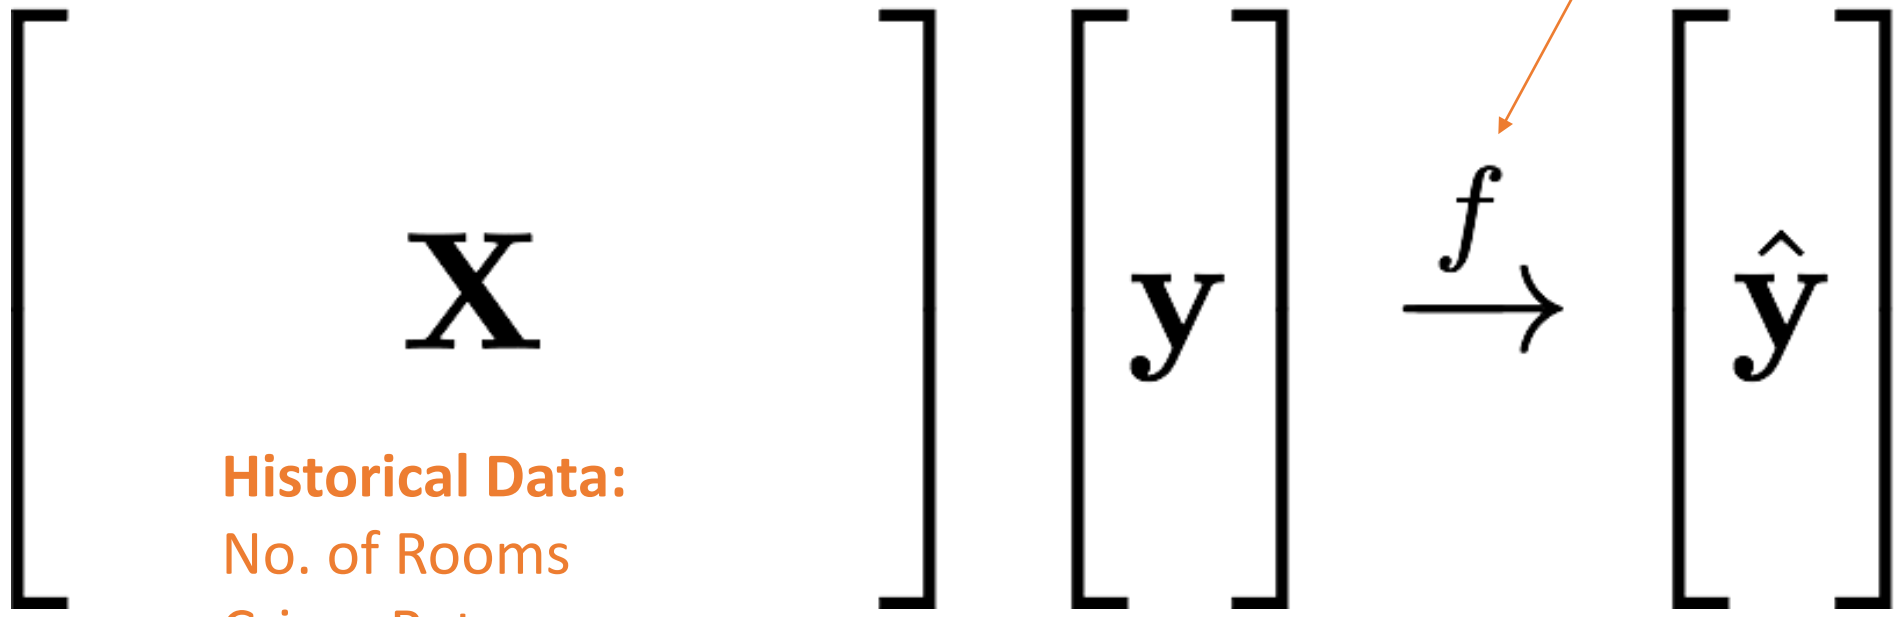

Historical Data:
No. of Rooms
Crime Rate
Pupil-Teacher Ratio
…

Target:
House Value

Predicted Value
(for evaluation)

H2O.ai

Supervised Learning Example

Patterns Learned
from Historical Data

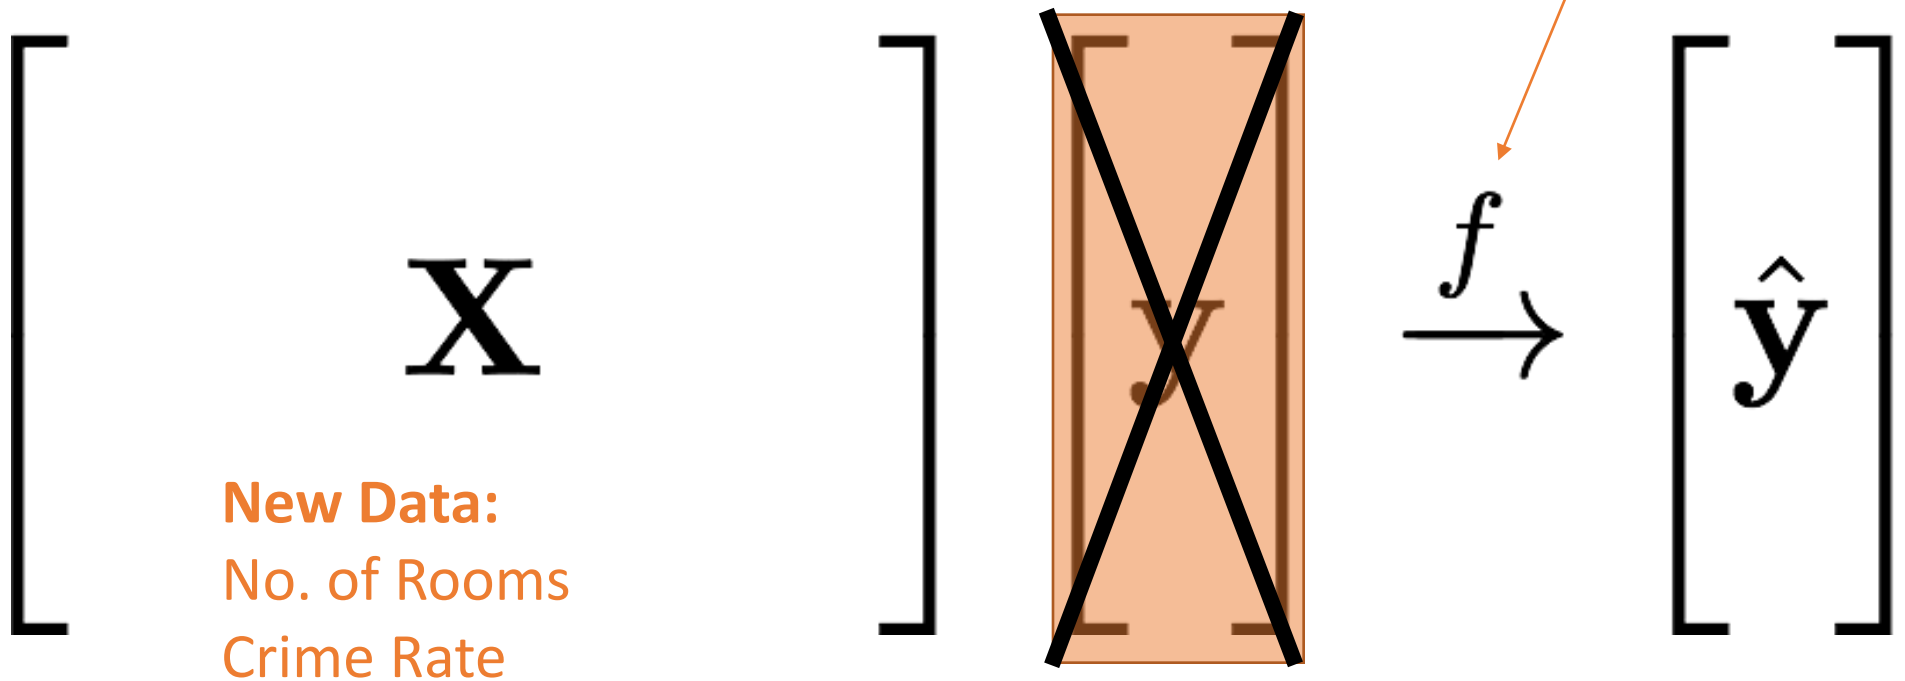

New Data:
No. of Rooms
Crime Rate
Pupil-Teacher Ratio
…

Target:
Unknown

Predicted Value
(for decision making)

H2O.ai

Supervised Learning – You Already Have Target Data

Regression:
How much will a customers spend?

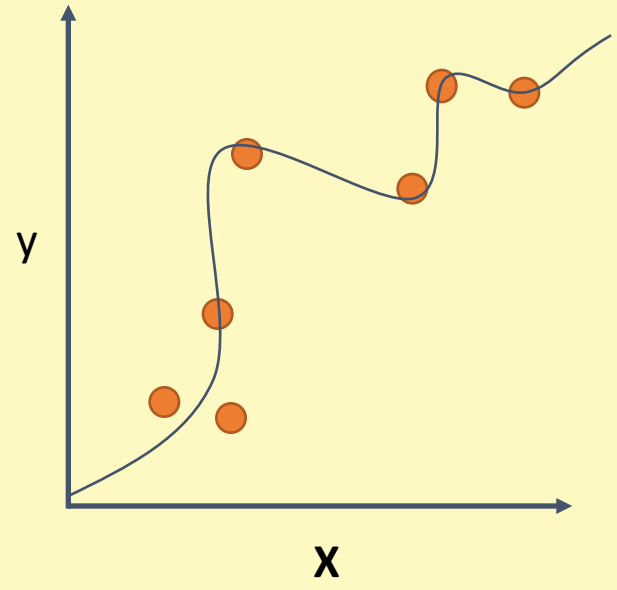

H_2O algos:
Penalized Linear Models
Random Forest
Gradient Boosting
Neural Networks
Stacked Ensembles

Classification:
Will a customer make a purchase? Yes or No

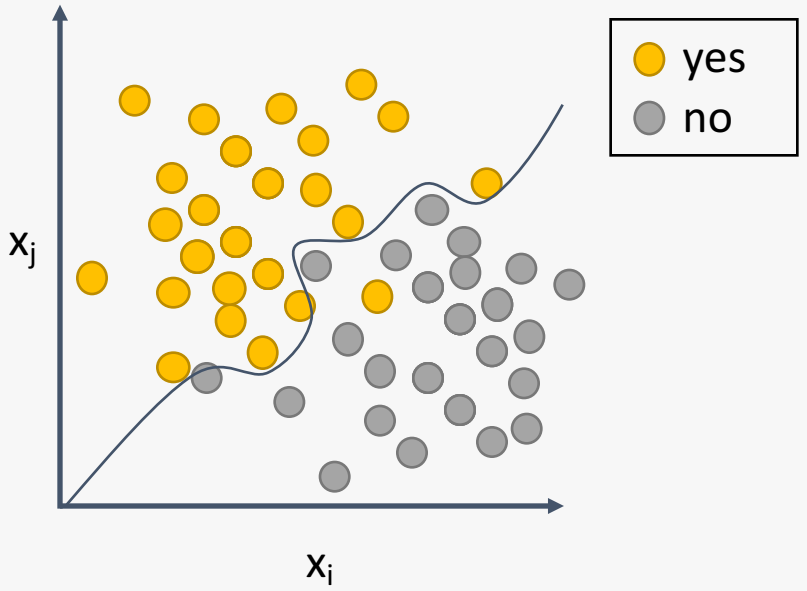

H_2O algos:
Penalized Linear Models
Naïve Bayes
Random Forest
Gradient Boosting
Neural Networks
Stacked Ensembles

H2O.ai

Unsupervised Learning – Discover Hidden Patterns

Clustering:
Customer Segmentation

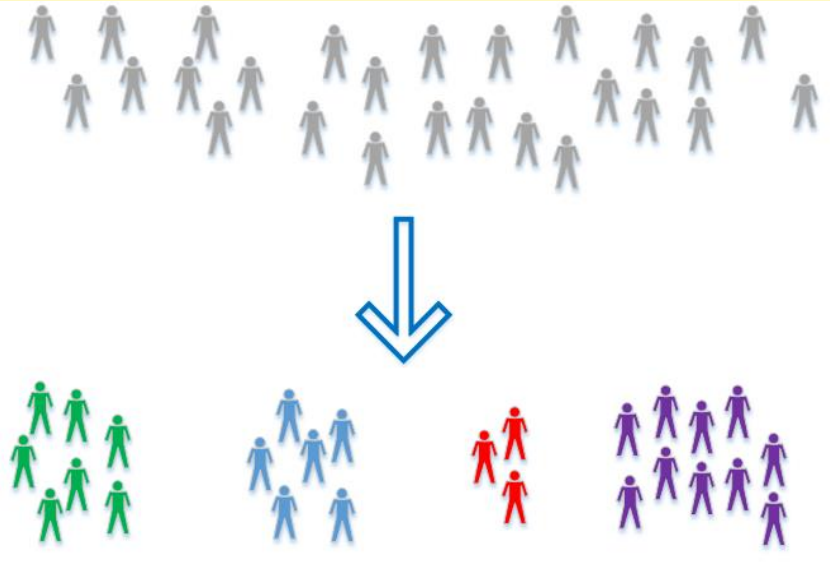

H_2O algos:
K-Means
Generalised Low Rank Model

Dimensionality Reduction:
Linear Transformation of Variables

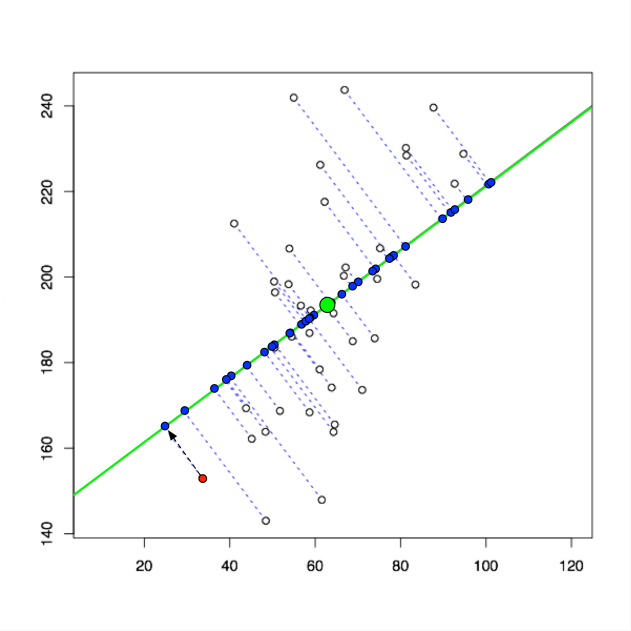

H_2O algos:
Principal Component Analysis
Generalised Low Rank Model

H2O.ai

About H_2O.ai

Company Overview

Founded	2011 Venture-backed, debuted in 2012
Products	• **H_2O Open Source In-Memory AI Prediction Engine** • Sparkling Water (H_2O + Spark) • Deep Water (H_2O + Other Deep Learning Frameworks) • Driverless AI (Next-Gen H_2O)
Mission	Operationalize Data Science, and provide a platform for users to build beautiful data products
Team	75+ employees • Distributed Systems Engineers doing Machine Learning • World-class visualization designers
Headquarters	Mountain View, CA






Alan Lewis


Amy Vu-Tran


Angela Bartz

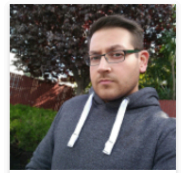
Anmol Bal


Apoorv Chaudhary

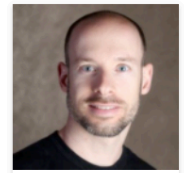
Arno Candel


Leland Wilkinson


Magnus Stensmo

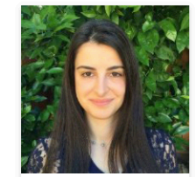
Maral Mandjian


Marios Michailidis


Mark Chan


Mark Landry


Ashrith Barthur


Avkash Chauhan


Avni Wadhwa


Beth Payne


Branden Murray


Cari Andrews


Mateusz Dymczyk


Matt Dowle


Megan Kurka

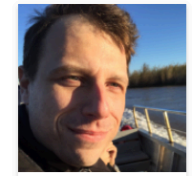
Michal Kurka


Michal Malohlava


Navdeep Gill


Das Narayanan


Dmitry Larko


Erin LeDell


Fonda Ingram


Ian Gomez


Jacqueline Scott


Nidhi Mehta


Nikhil Shekhar


Nishant Kalonia


Pasha Stetsenko


Patrick Hall


Patrick Rice


Jag Bining


Jakub Hava


Jeff Fohl


Jeff Gambera

Jo-Fai Chow

Jon Mckinney

Prithvi Prabhu

Ravi Purushotma

Raymond Peck

Sanjay Joshi

SriSatish Ambati

Terone Ward

Jon Oslowski

Josephine Wang

Justin Loyola

Karen Hayrapetyan

Kerry O'Shea

Lauren DiPerna

Tom Aliotti

Tom Kraljevic

Tomas Nykodym

Venkatesh Yadav

Vinod Iyengar

Wen Phan

Scientific Advisory Council

Dr. Trevor Hastie

- John A. Overdeck Professor of Mathematics, Stanford University
- PhD in Statistics, Stanford University
- Co-author, *The Elements of Statistical Learning: Prediction, Inference and Data Mining*
- Co-author with John Chambers, *Statistical Models in S*
- Co-author, *Generalized Additive Models*

Dr. Robert Tibshirani

- Professor of Statistics and Health Research and Policy, Stanford University
- PhD in Statistics, Stanford University
- Co-author, *The Elements of Statistical Learning: Prediction, Inference and Data Mining*
- Author, *Regression Shrinkage and Selection via the Lasso*
- Co-author, *An Introduction to the Bootstrap*

Dr. Steven Boyd

- Professor of Electrical Engineering and Computer Science, Stanford University
- PhD in Electrical Engineering and Computer Science, UC Berkeley
- Co-author, *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*
- Co-author, *Linear Matrix Inequalities in System and Control Theory*
- Co-author, *Convex Optimization*

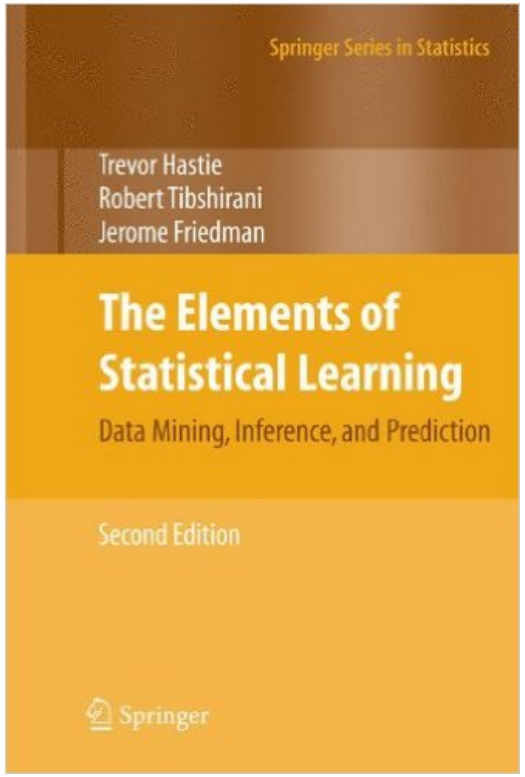

H2O.ai

H2O Community & Fortune 100 customers

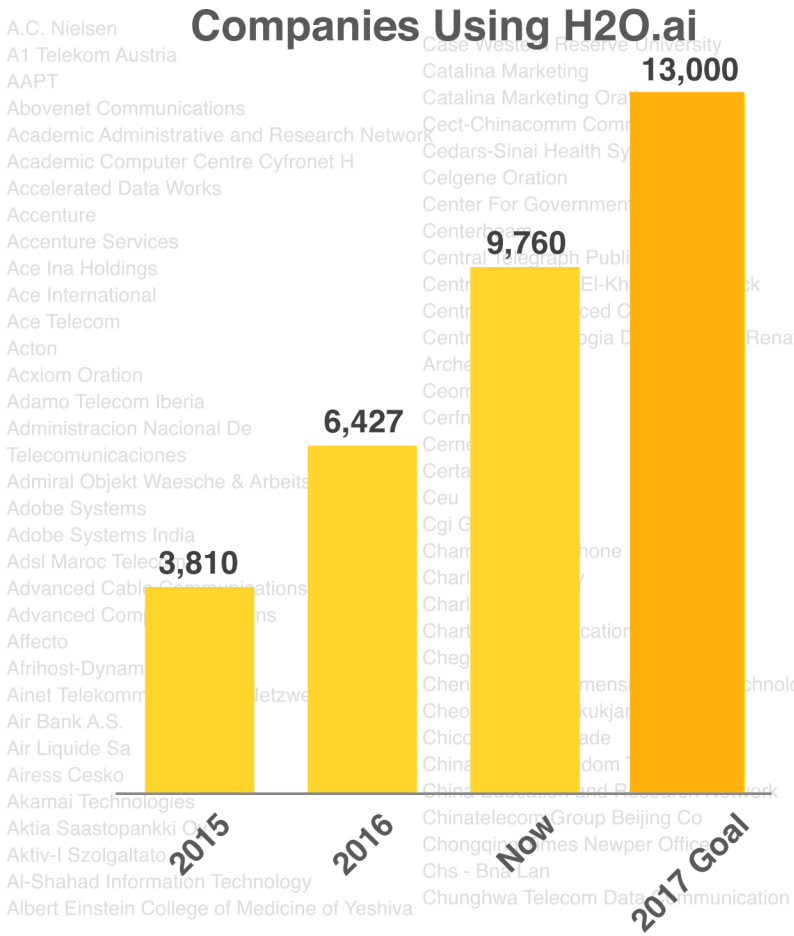

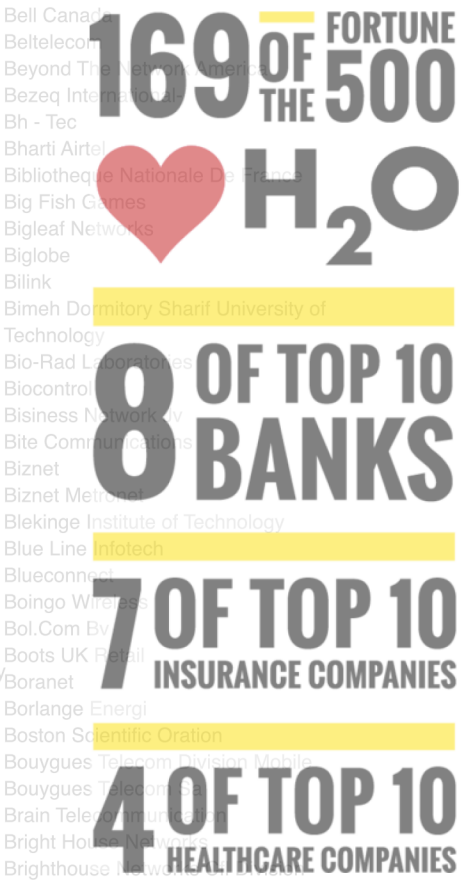

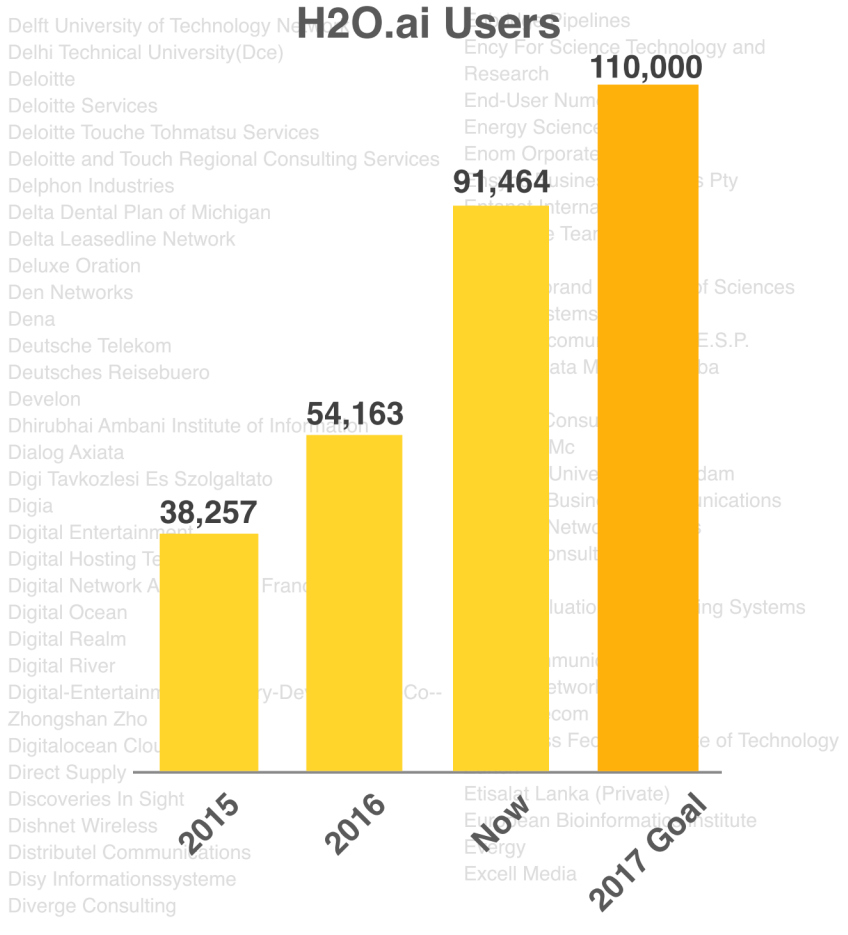

Select Reference Customers:

"Overall customer satisfaction is very high." - Gartner

CapitalOne PayPal Experian WELLS FARGO ADP COMCAST at&t KAISER PERMANENTE PROGRESSIVE Nielsen Catalina Booking.com macy's H2O.ai

pwc | Global

PwC Global › About us › › Stories about our people, clients and communities

Harnessing the power of AI to transform the detection of fraud and error

Setting the scene

PwC has invested significantly in pioneering the use of artificial intelligence for the audit and has partnered with H2O.ai, a leading Silicon Valley-based AI company.

Following 18 months of development, the first outcome of this partnership is PwC's GL.ai, the first module of PwC's Audit.ai - a revolutionary bot that does what humans can't. Its AI analyses billions of different data points in seconds and applies judgement to detect anomalies in general ledger transactions.

"The reason this is such a brilliant tool is the ability to look at different risks in context at the same time. For example, it would be uneconomical for an auditor to look at every single user's pattern of activity and decide what was unusual. With GL.ai, the algorithms do it for us."

Laura Needham partner, PwC UK

http://www.pwc.com/gx/en/about/stories-from-across-the-world/harnessing-the-power-of-ai-to-transform-the-detection-of-fraud-and-error.html

H2O.ai

Community Expansion

Meetup

H2O.ai

56,536	32	50	44	18
members	interested	Meetups	cities	countries

15 Meetups a month

Oct 17 – Amsterdam
Nov 17 – London

#AroundTheWorldWithH2Oai

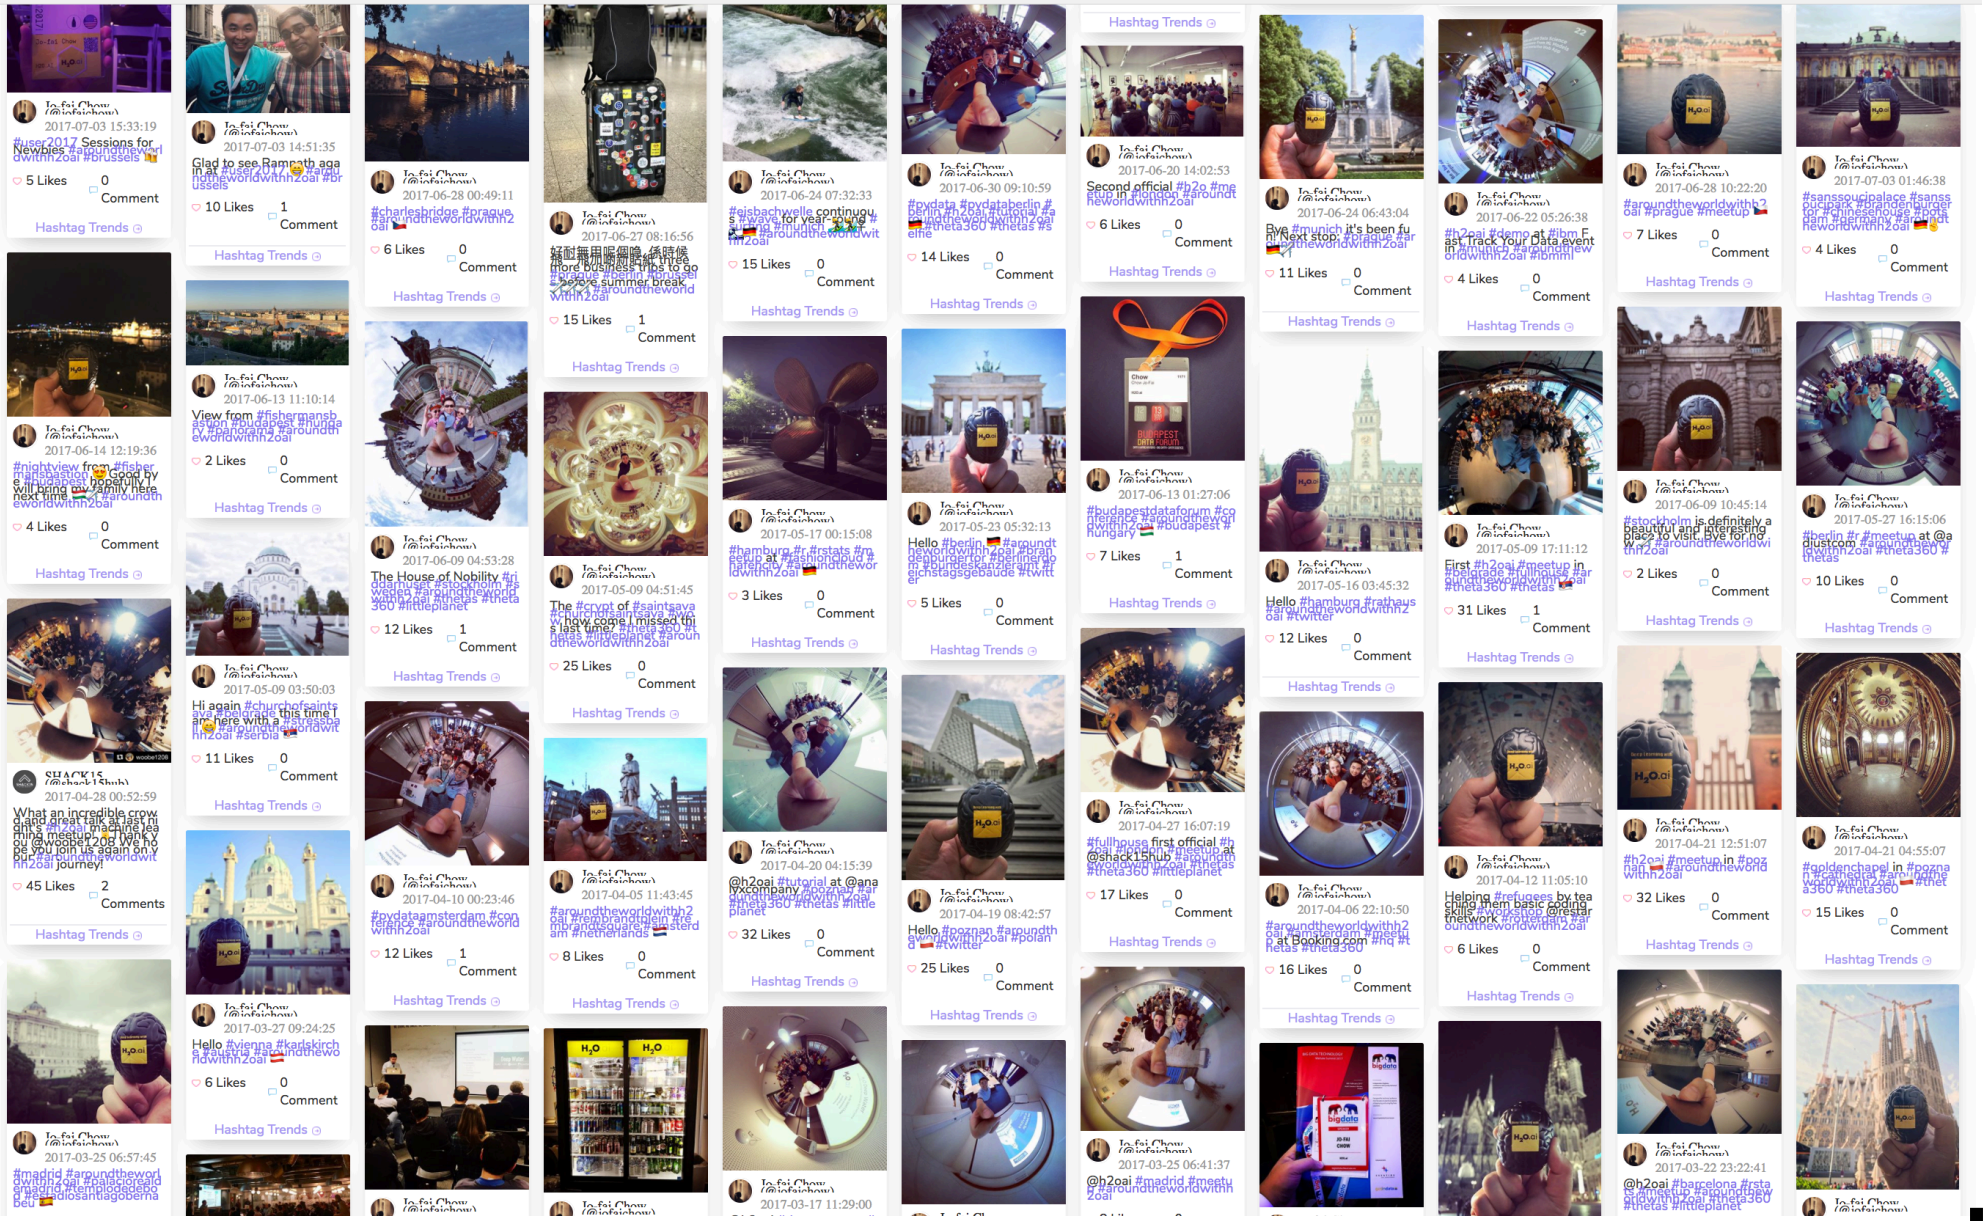

H_2O Machine Learning Platform

High Level Architecture

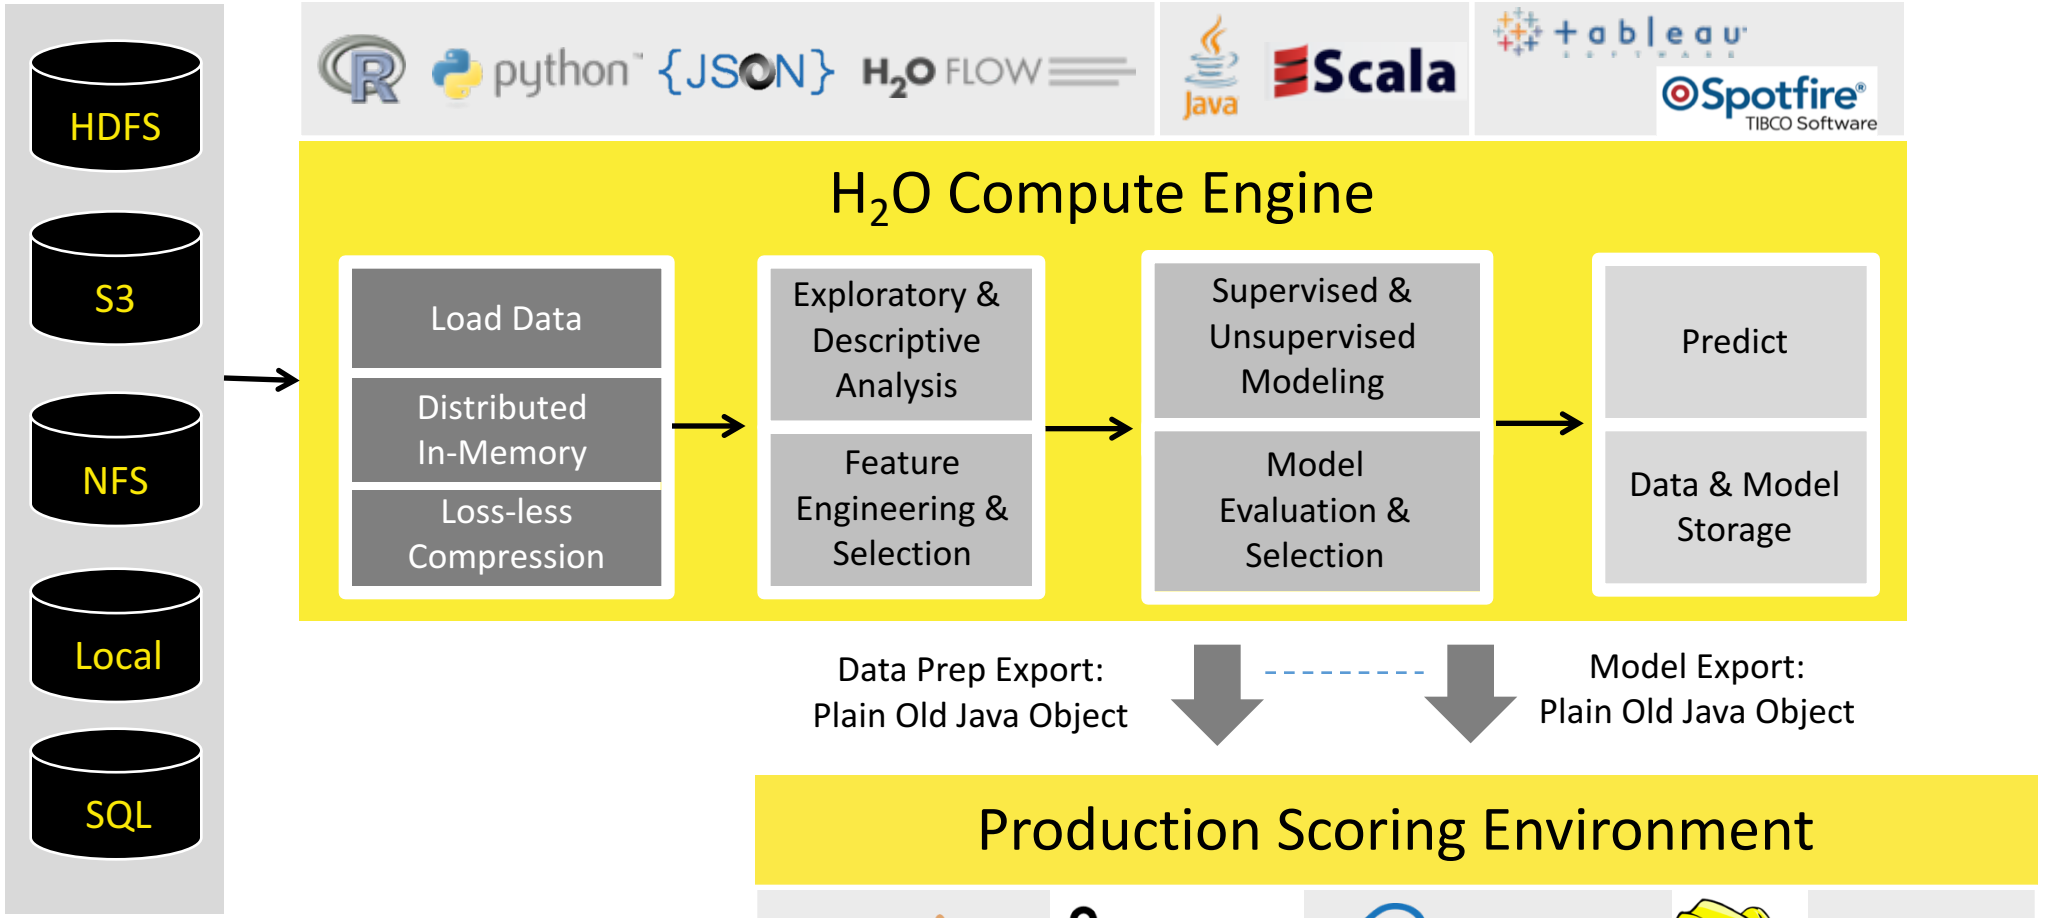

High Level Architecture

Import Data from Multiple Sources

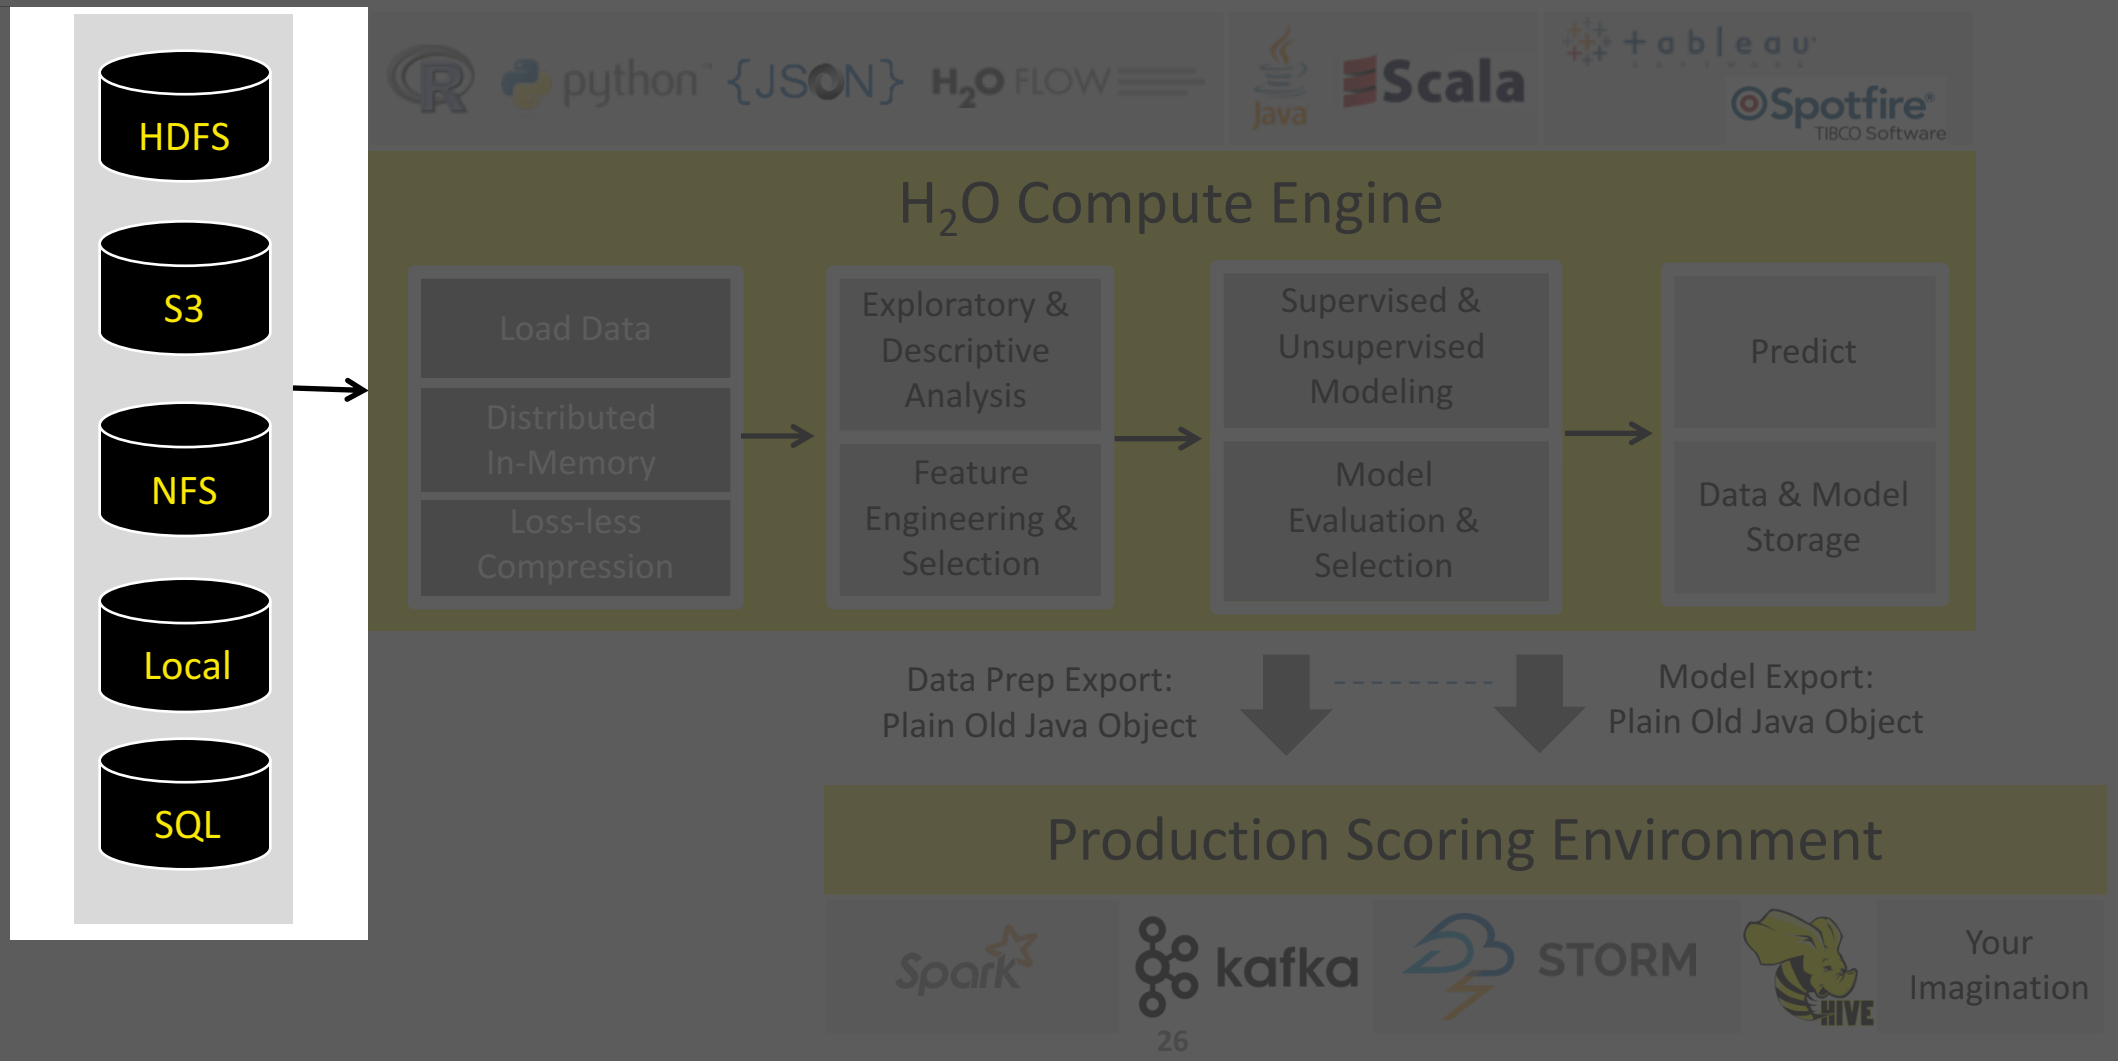

High Level Architecture

Fast, Scalable & Distributed Compute Engine Written in Java

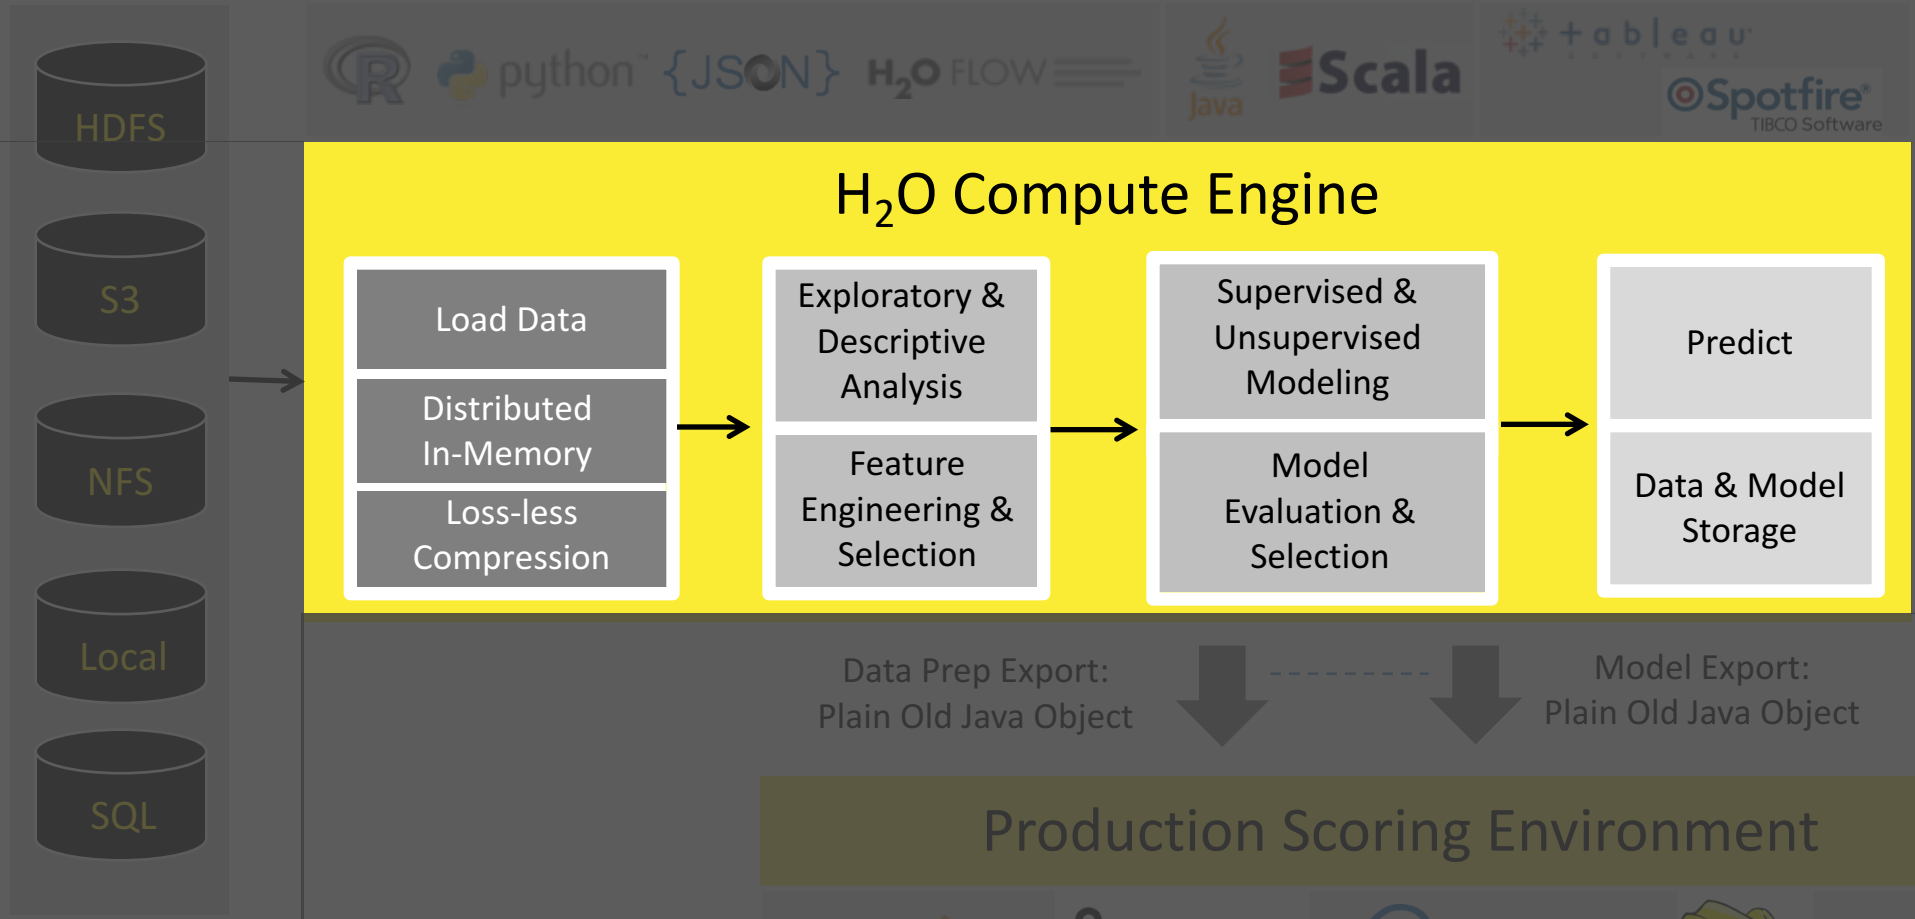

Algorithms Overview

Supervised Learning

Statistical Analysis	• **Generalized Linear Models**: Binomial, Gaussian, Gamma, Poisson and Tweedie • **Naïve Bayes**
Ensembles	• **Distributed Random Forest**: Classification or regression models • **Gradient Boosting Machine**: Produces an ensemble of decision trees with increasing refined approximations
Deep Neural Networks	• **Deep learning**: Create multi-layer feed forward neural networks starting with an input layer followed by multiple layers of nonlinear transformations

Unsupervised Learning

Clustering	• **K-means**: Partitions observations into k clusters/groups of the same spatial size. Automatically detect optimal k
Dimensionality Reduction	• **Principal Component Analysis**: Linearly transforms correlated variables to independent components • **Generalized Low Rank Models:** extend the idea of PCA to handle arbitrary data consisting of numerical, Boolean, categorical, and missing data
Anomaly Detection	• **Autoencoders**: Find outliers using a nonlinear dimensionality reduction using deep learning

High Level Architecture

Multiple Interfaces

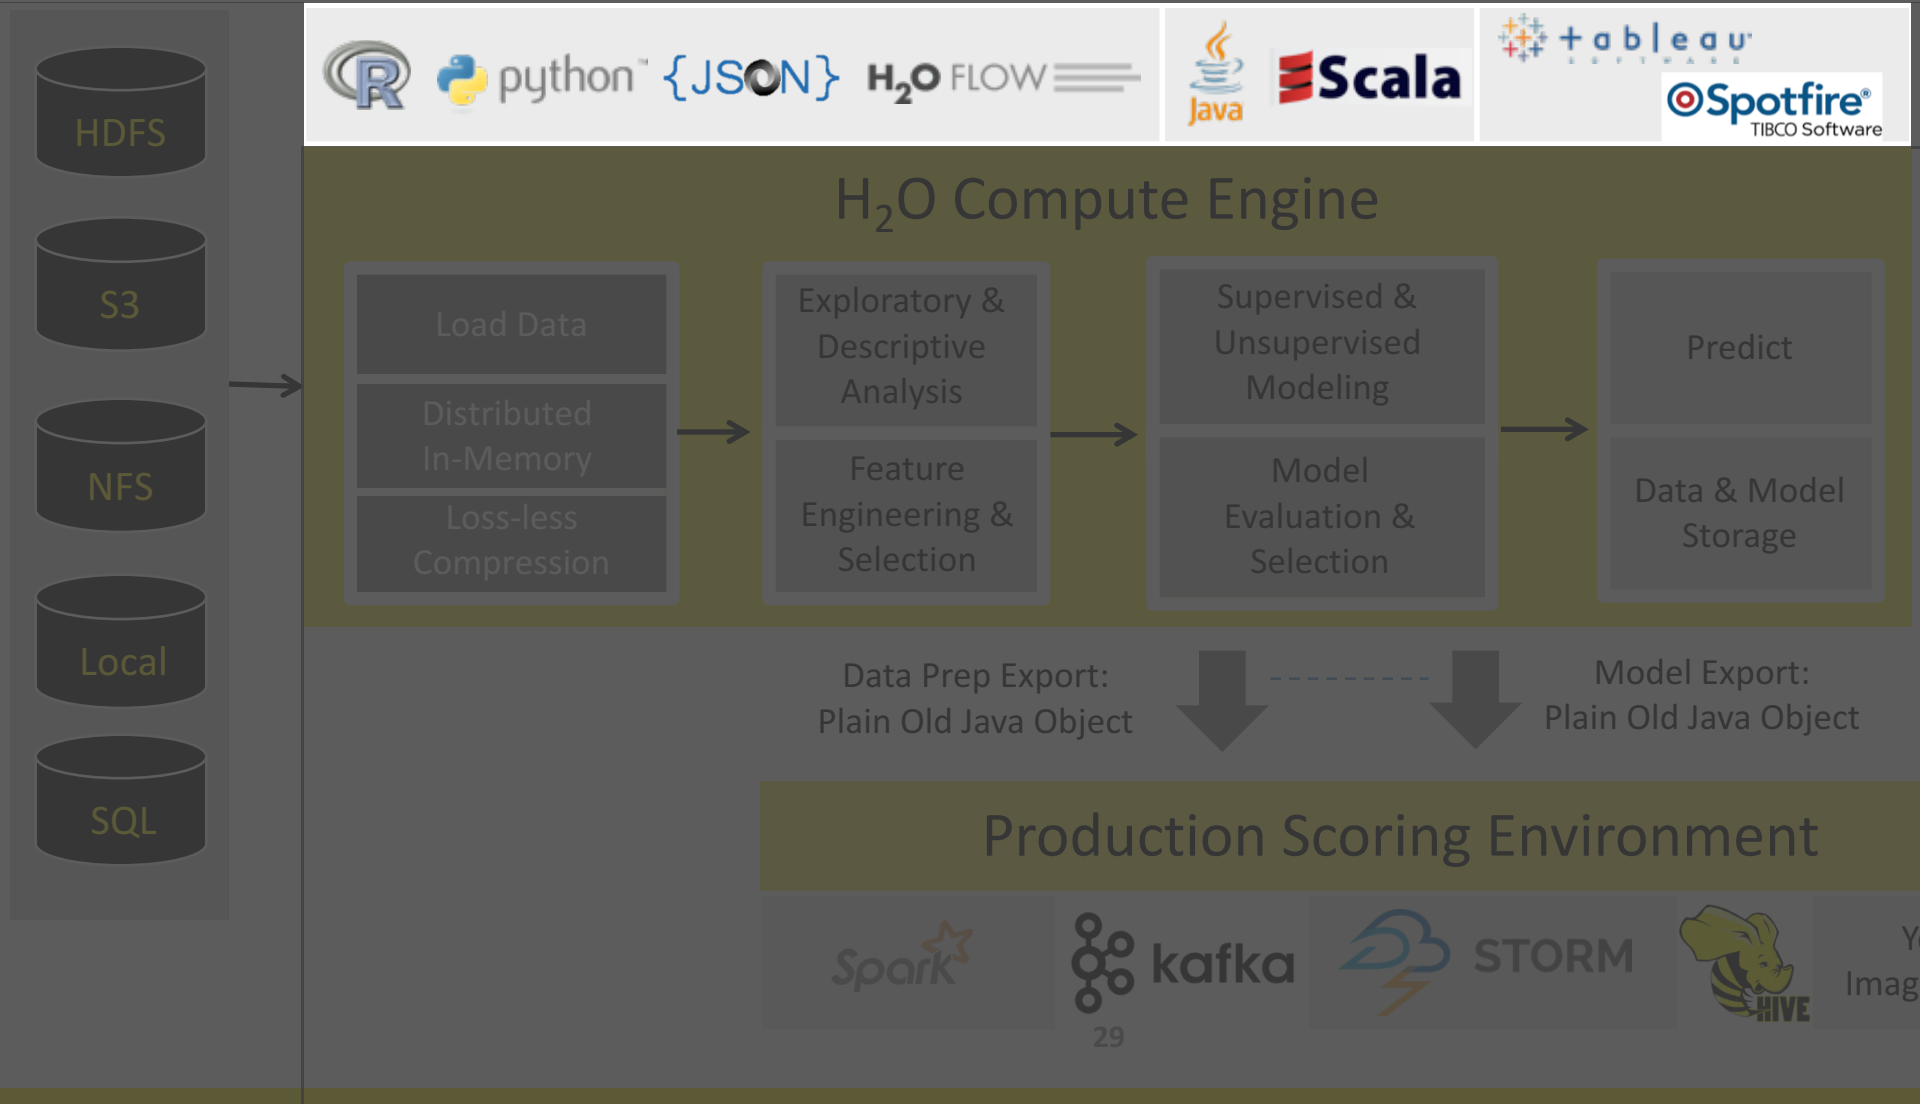

H_2O Flow (Web)

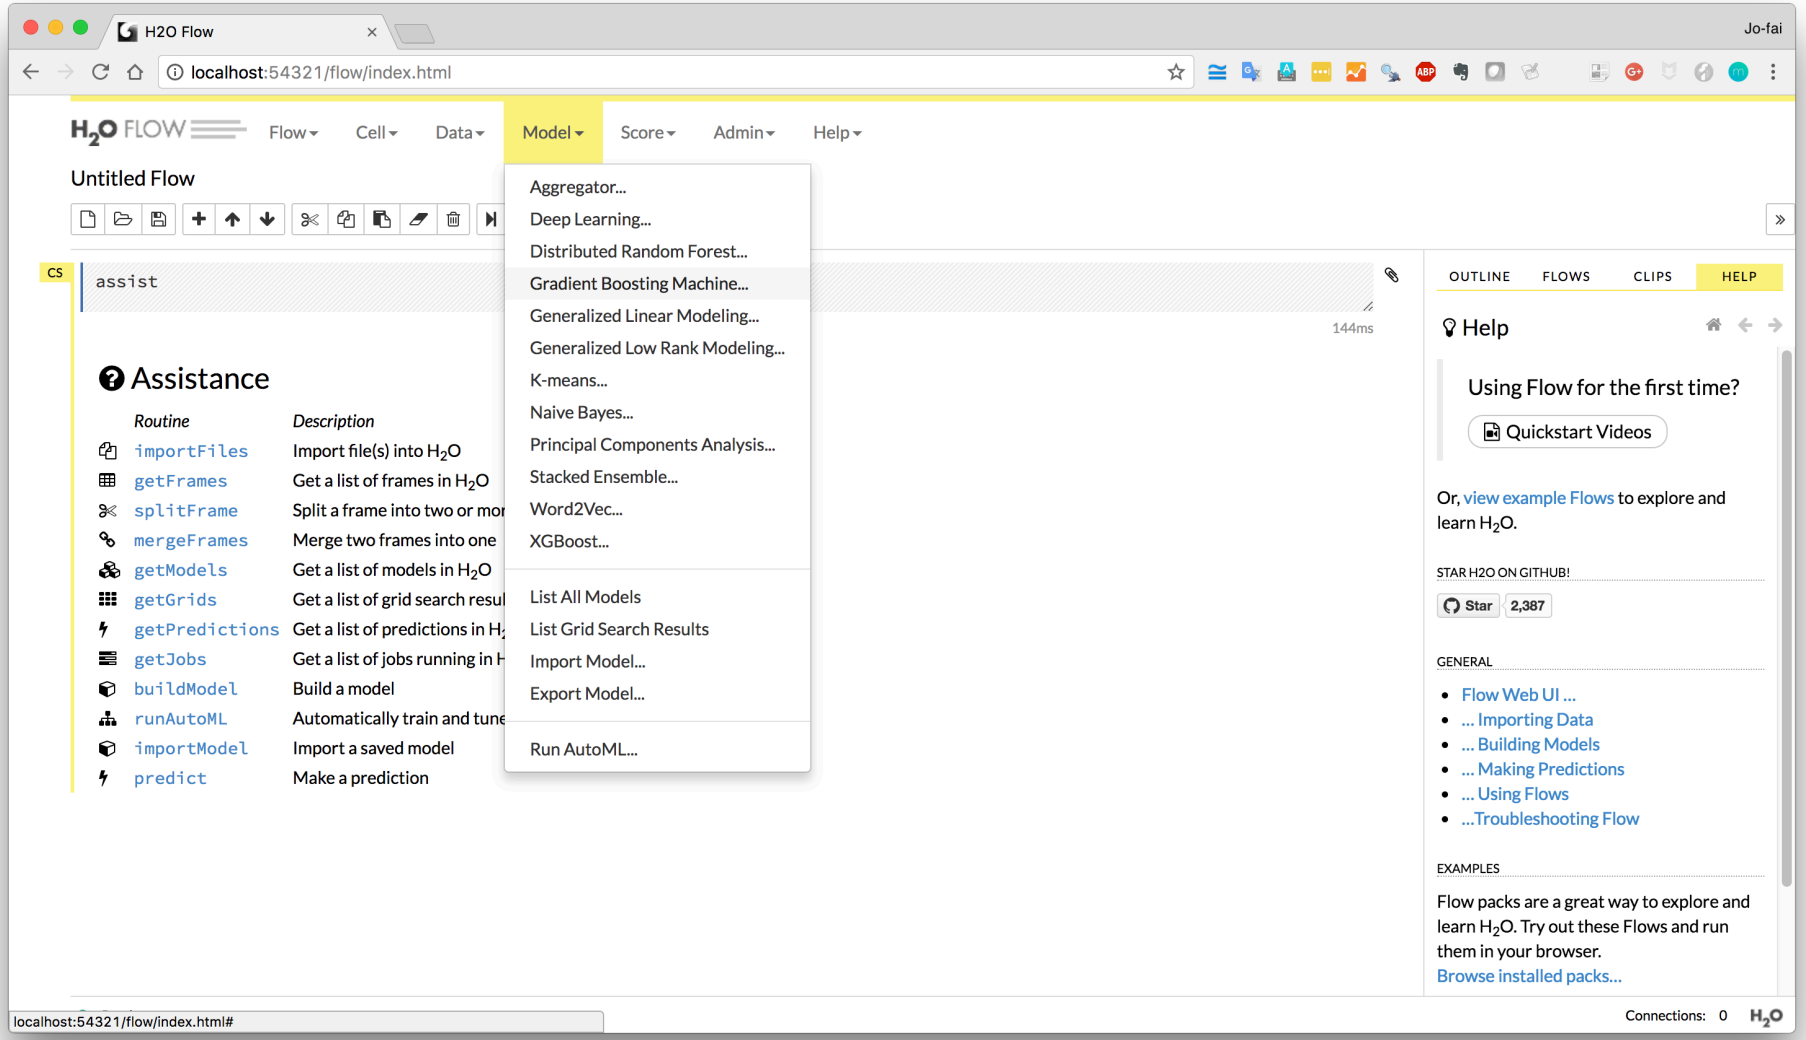

H2O.ai

Using H_2O with R and Python

~/Documents/repo_h2o/sales-engineering - master - RStudio Source Editor

credit_card_example.R

```r
# Credit Card Example

# Datasets:
# https://s3.amazonaws.com/h2o-training/credit_card/credit_card_train.csv
# https://s3.amazonaws.com/h2o-training/credit_card/credit_card_test.csv

# Start and connect to a local H2O cluster
library(h2o)
h2o.init(nthreads = -1)

# Import datasets from s3
df_train = h2o.importFile("https://s3.amazonaws.com/h2o-training/credit_card/credit_card_train.csv")
df_test = h2o.importFile("https://s3.amazonaws.com/h2o-training/credit_card/credit_card_test.csv")

# Look at datasets
summary(df_train)
summary(df_test)

# Define features and target
features = colnames(df_test)
target = "DEFAULT_PAYMENT_NEXT_MONTH"

# Train a GBM model
model_gbm = h2o.gbm(x = features,
                    y = target,
                    training_frame = df_train,
                    seed = 1234)
print(model_gbm)

# Use GBM model for making predictions
yhat_test = h2o.predict(model_gbm, newdata = df_test)
head(yhat_test)

# (Extra) Use H2O's AutoML
aml = h2o.automl(x = features,
                 y = target,
                 training_frame = df_train,
                 max_runtime_secs = 60,
                 seed = 1234)

# Print leaderboard
print(aml@leaderboard)

# Use best model for making predictions
best_model = aml@leader
yhat_test = h2o.predict(best_model, newdata = df_test)
head(yhat_test)
```

49:1 (Top Level) R Script

credit_card_example

localhost:8888/notebooks/credit_card_example.ipynb

jupyter credit_card_example Last Checkpoint: 5 minutes ago (unsaved changes)

File Edit View Insert Cell Kernel Help Python 2

Code CellToolbar

In [2]:
```r
# Start and connect to a local H2O cluster
import h2o
h2o.init(nthreads = -1)
```

```
Checking whether there is an H2O instance running at http://localhost:54321..... not found.
Attempting to start a local H2O server...
  Java Version: java version "1.8.0_72"; Java(TM) SE Runtime Environment (build 1.8.0_72-b15); Java HotSpot(TM) 64-Bit Server VM (build 25.72-b15, mixed mode)
  Starting server from /Users/jofaichow/anaconda/lib/python2.7/site-packages/h2o/backend/bin/h2o.jar
  Ice root: /var/folders/4z/p7yt7_4n4fj1jlyq6g4qhfbw0000gn/T/tmpFdP3Av
  JVM stdout: /var/folders/4z/p7yt7_4n4fj1jlyq6g4qhfbw0000gn/T/tmpFdP3Av/h2o_jofaichow_started_from_python.out
  JVM stderr: /var/folders/4z/p7yt7_4n4fj1jlyq6g4qhfbw0000gn/T/tmpFdP3Av/h2o_jofaichow_started_from_python.err
  Server is running at http://127.0.0.1:54321
Connecting to H2O server at http://127.0.0.1:54321... successful.
```

H2O cluster uptime:	02 secs
H2O cluster version:	3.13.0.3981
H2O cluster version age:	29 days
H2O cluster name:	H2O_from_python_jofaichow_id7iqa
H2O cluster total nodes:	1

In [3]:
```r
# Import datasets from s3
df_train = h2o.import_file("https://s3.amazonaws.com/h2o-training/credit_card/credit_card_train.csv")
df_test = h2o.import_file("https://s3.amazonaws.com/h2o-training/credit_card/credit_card_test.csv")
```

```
Parse progress: |████████████████████████████████████████████████| 100%
Parse progress: |████████████████████████████████████████████████| 100%
```

In [4]:
```r
# Look at datasets
df_train.summary()
df_test.summary()
```

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4
type	int	enum	int	int	int	int	int	int	int
mins	10000.0		0.0	0.0	21.0	-2.0	-2.0	-2.0	-2.0
mean	165471.466667		1.85	1.55578703704	35.4053240741	-0.00523148148148	-0.122361111111	-0.15537037037	-0.21060
maxs	1000000.0		6.0	3.0	79.0	8.0	8.0	8.0	8.0
sigma	128853.314839		0.779559696278	0.522505078476	9.27675421641	1.12668964211	1.20086854503	1.20727030901	1.17217
zeros	0		9	37	0	10563	11284	11309	11905
missing	0	0	0	0	0	0	0	0	0

H2O.ai

High Level Architecture

Export Standalone Models for Production

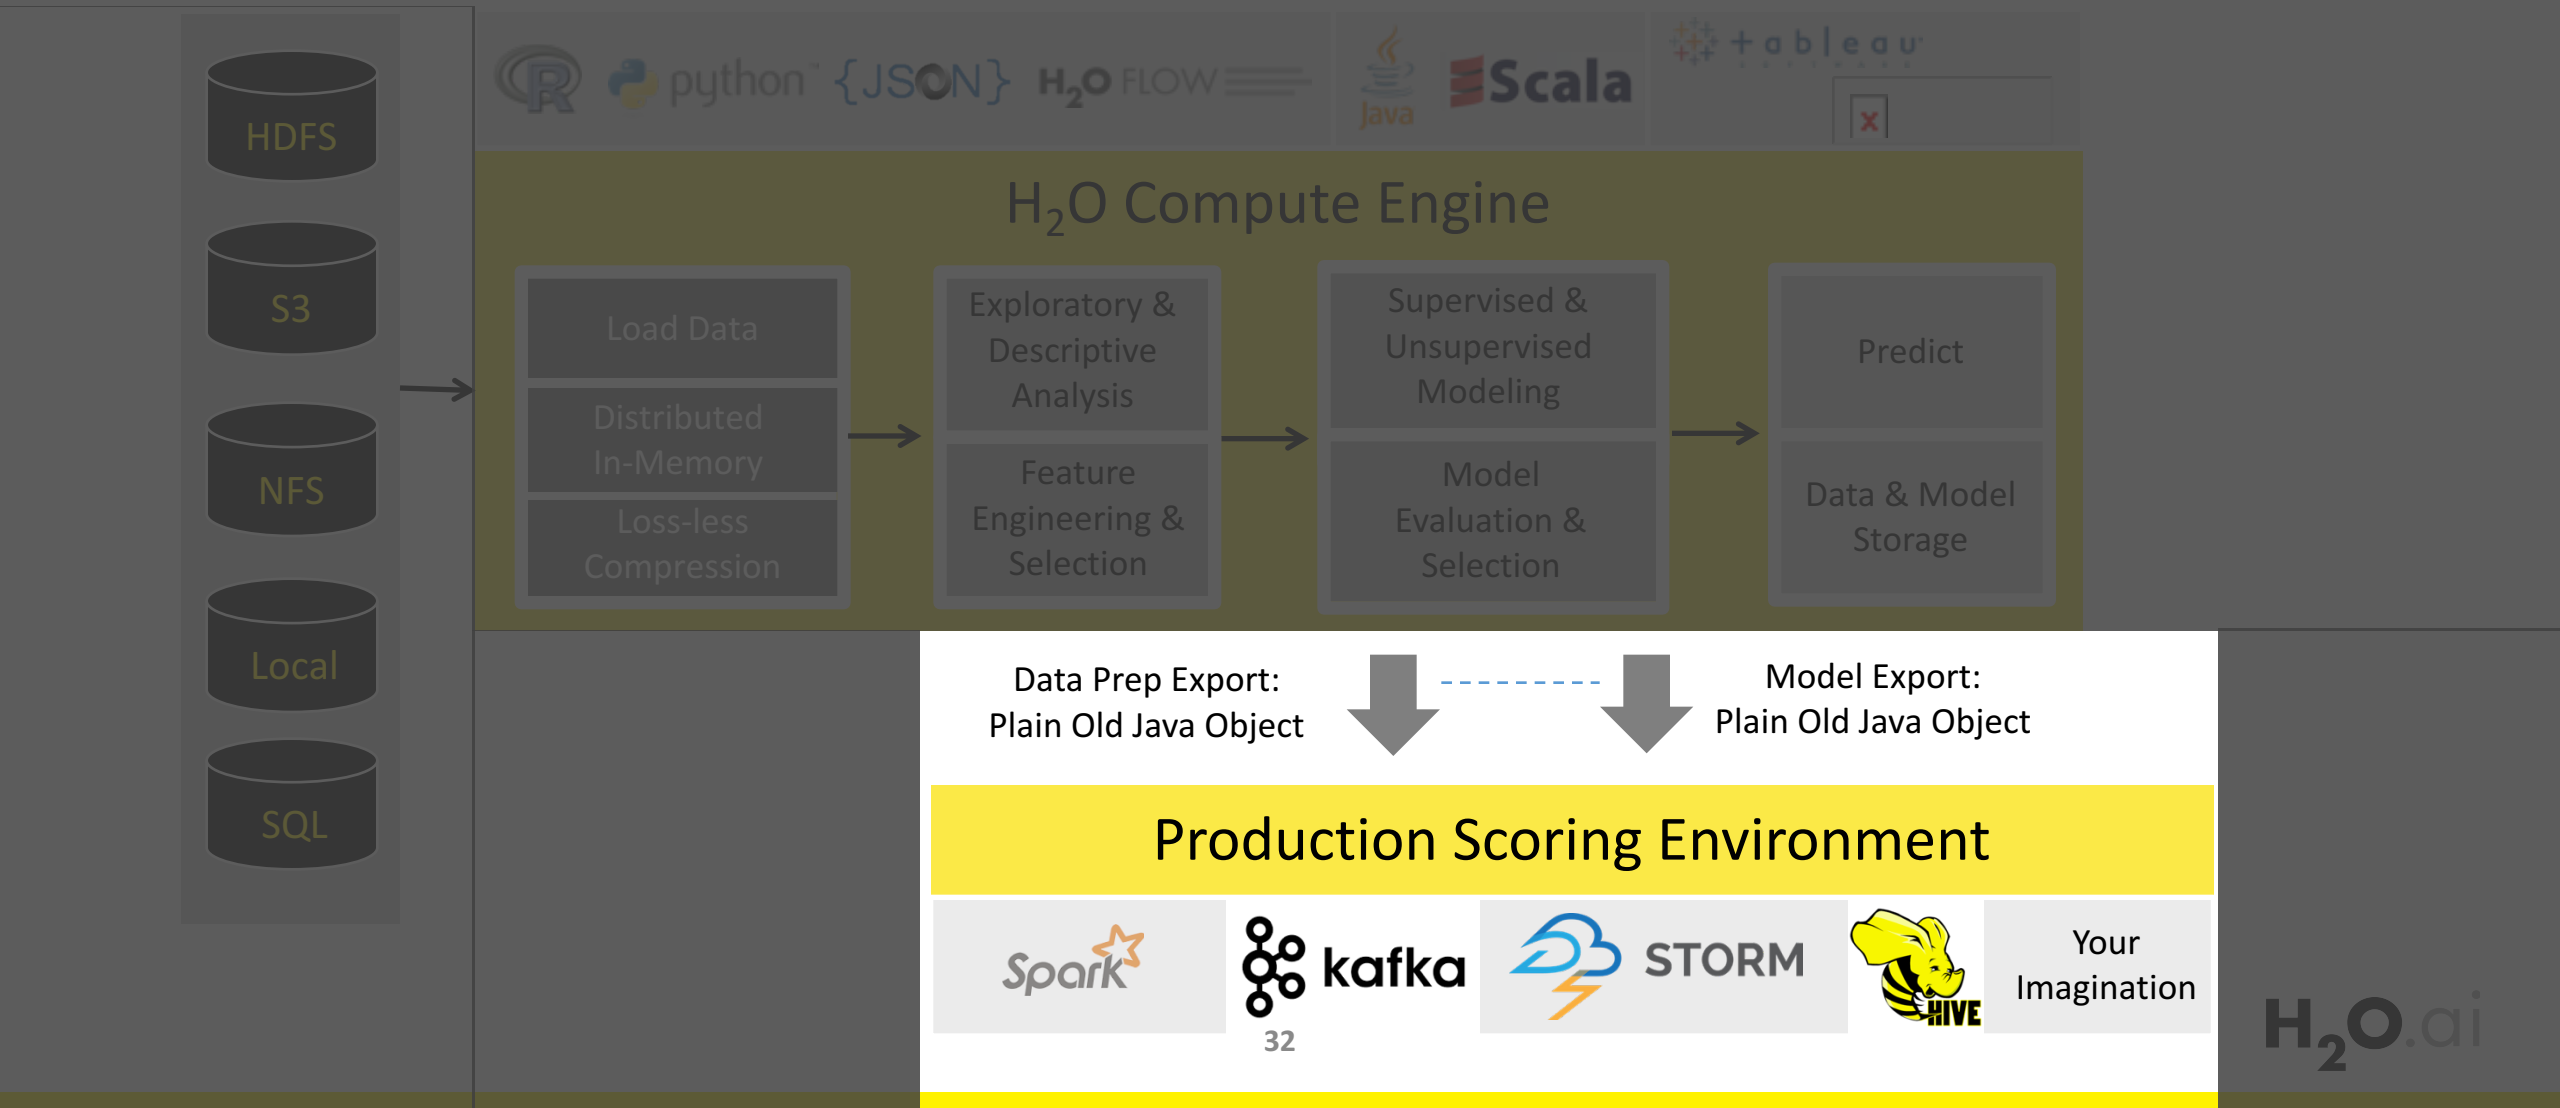

High Level Architecture

This Workshop

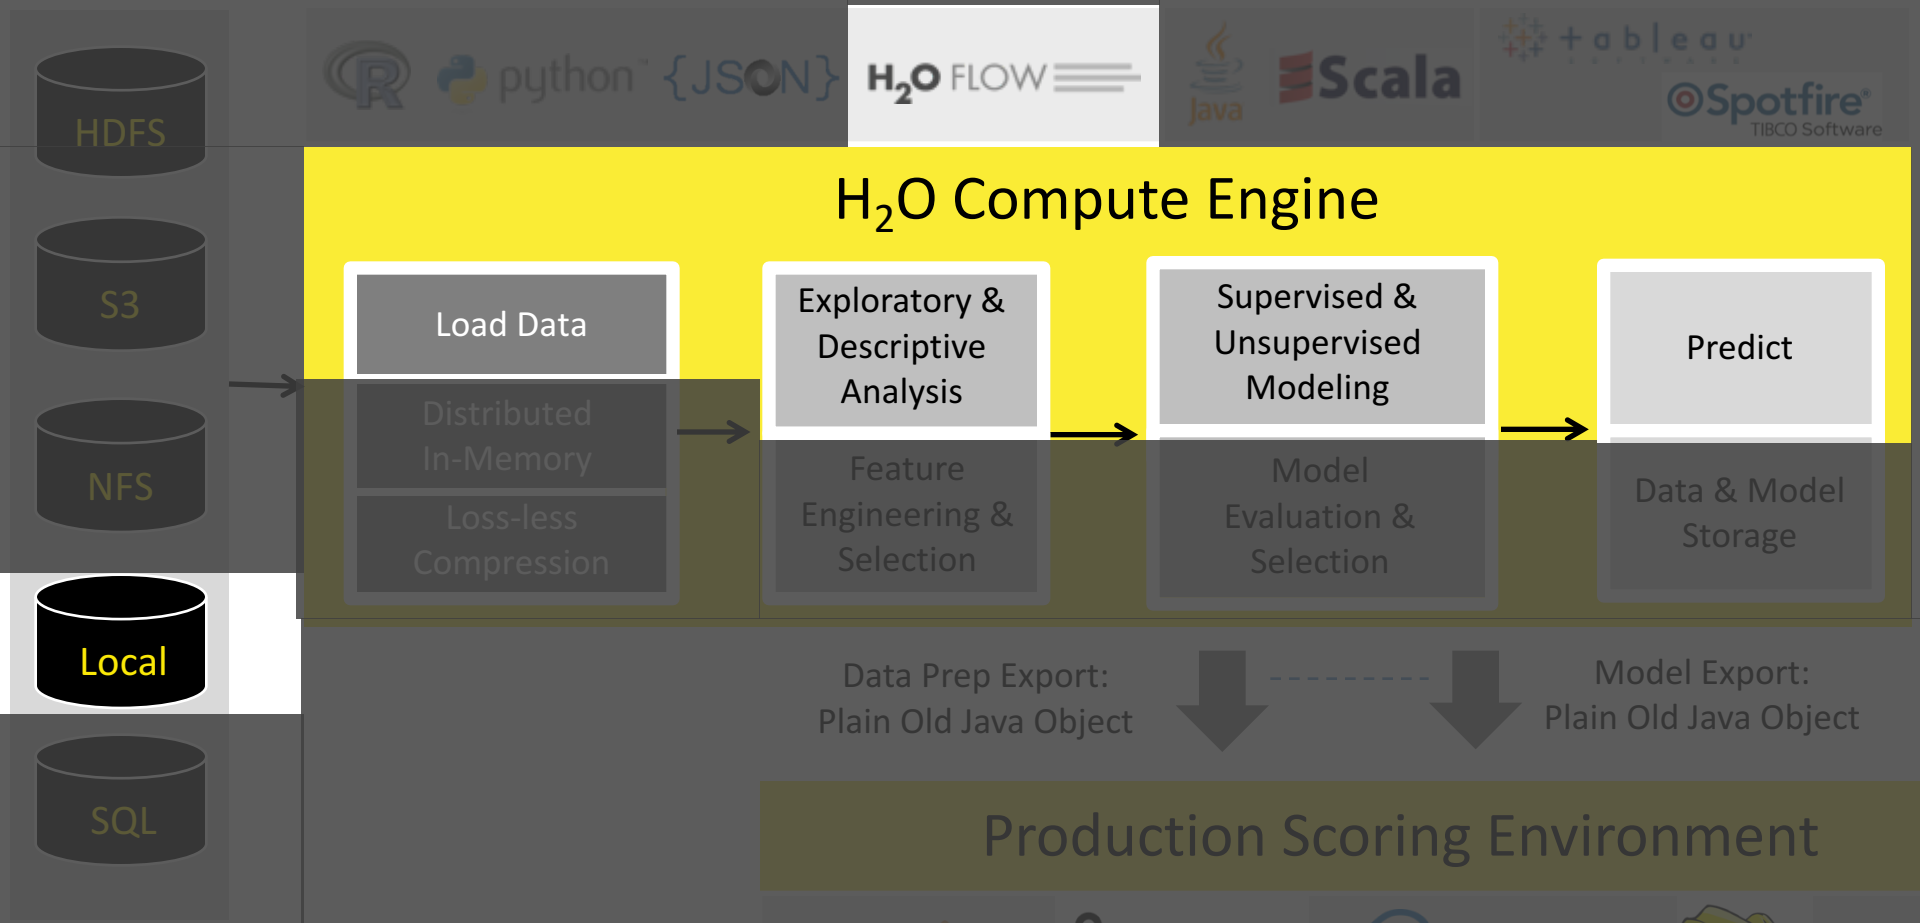

docs.h2o.ai

Getting Started & User Guides

H_2O

What is H2O?
H2O User Guide (Main docs)
H2O Book (O'Reilly)
Recent Changes
Open Source License (Apache V2)

Quick Start Video - Flow Web UI
Quick Start Video - R
Quick Start Video - Python

Download H_2O

Sparkling Water

What is Sparkling Water?
Sparkling Water Booklet
PySparkling Readme 2.0 | 2.1 | 2.2
RSparkling Readme
Open Source License (Apache V2)

Quick Start Video - Scala

Download Sparkling Water

Steam

What is Steam?
Steam User Guide
Recent Changes
Open Source License (AGPL)

Download Steam

Deep Water (preview)

Deep Water Readme
Deep Water Booklet
Deep Water AMI Guide
Deep Water Docker Image
Open Source License (Apache V2)

Launch Deep Water AMI
(choose p2.xlarge)

Q & A

FAQ
Issue Tracking (JIRA)
Stack Overflow
h2ostream Google Group
Gitter
Cross Validated

For Supported Enterprise Customers

Enterprise Support Web | Email

Algorithms

Supervised Learning

Generalized Linear Modeling (GLM)	Tutorial	Booklet	Reference	Tuning
Gradient Boosting Machine (GBM)	Tutorial	Booklet	Reference	Tuning
Deep Learning	Tutorial	Booklet	Reference	Tuning
Distributed Random Forest	Tutorial	Booklet	Reference	Tuning
Naive Bayes	Tutorial	Booklet	Reference	Tuning
Stacked Ensembles	Tutorial	Booklet	Reference	Tuning
XGBoost	Tutorial	Booklet	Reference	Tuning

Unsupervised Learning

Generalized Low Rank Models (GLRM)	Tutorial	Reference
K-Means Clustering	Tutorial	Reference
Principal Components Analysis (PCA)	Tutorial	Reference

Miscellaneous

Word2vec	Tutorial	Reference

TOP

H_2O Tutorials

- Quick Start
- Regression
- Classification
- Clustering

H_2O Quick Start

www.h2o.ai/download

Prerequisite: Java version 7 or 8
(Note: Java version 9 is not yet supported)

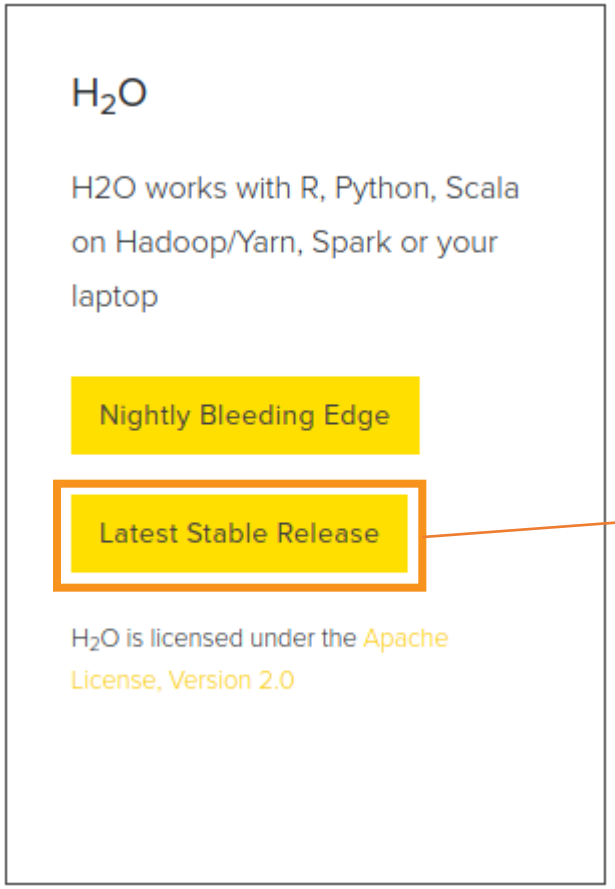

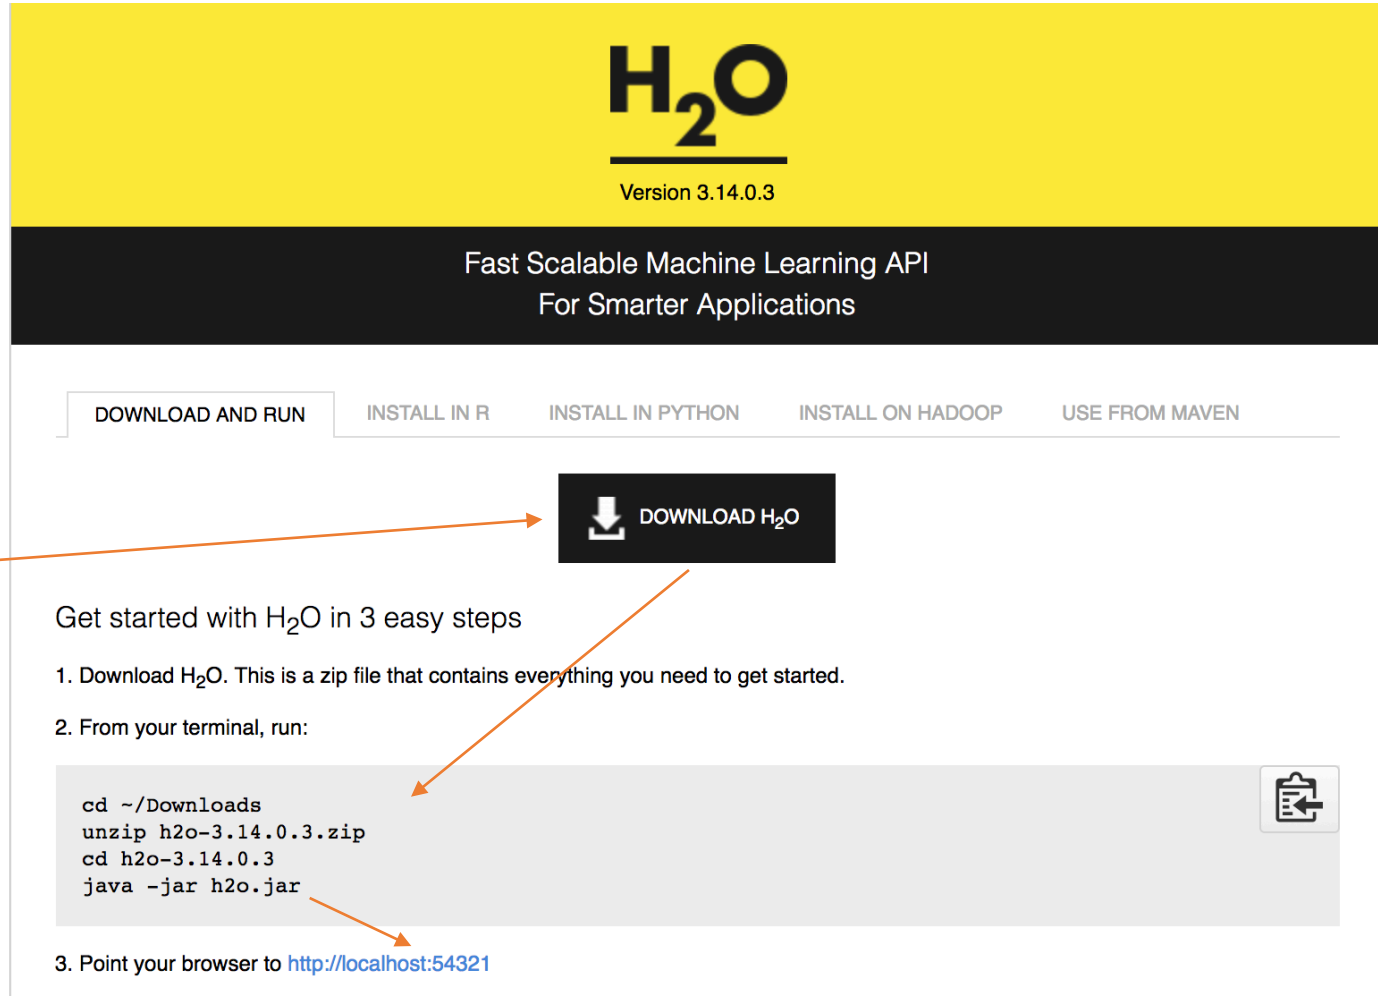

H2O.ai

Install and Start H_2O Flow (Web Interface)

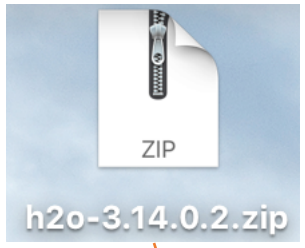

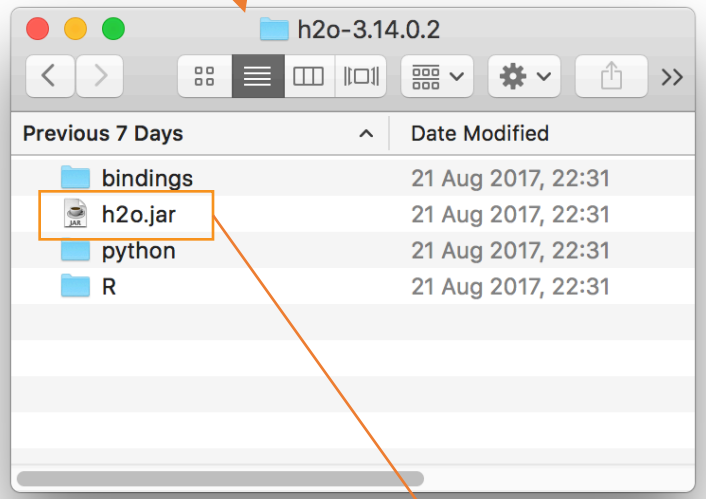

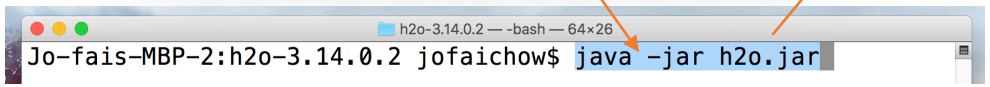

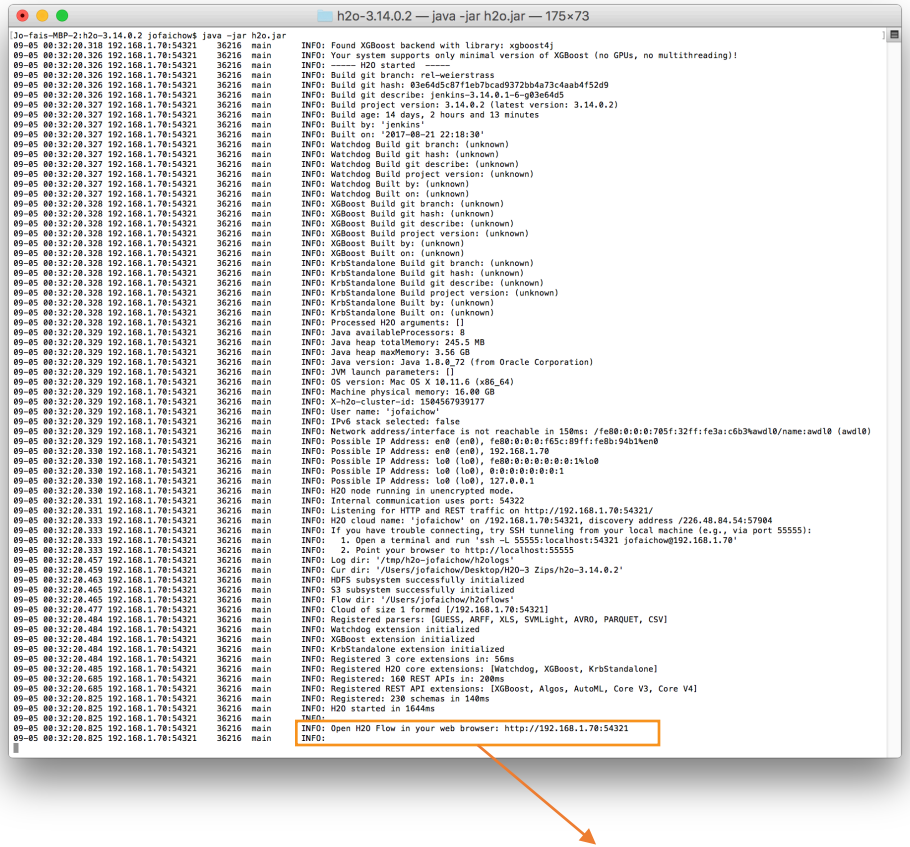

INFO: Open H2O Flow in your web browser: http://192.168.1.70:54321

H2O.ai

H_2O Flow (Web Interface)

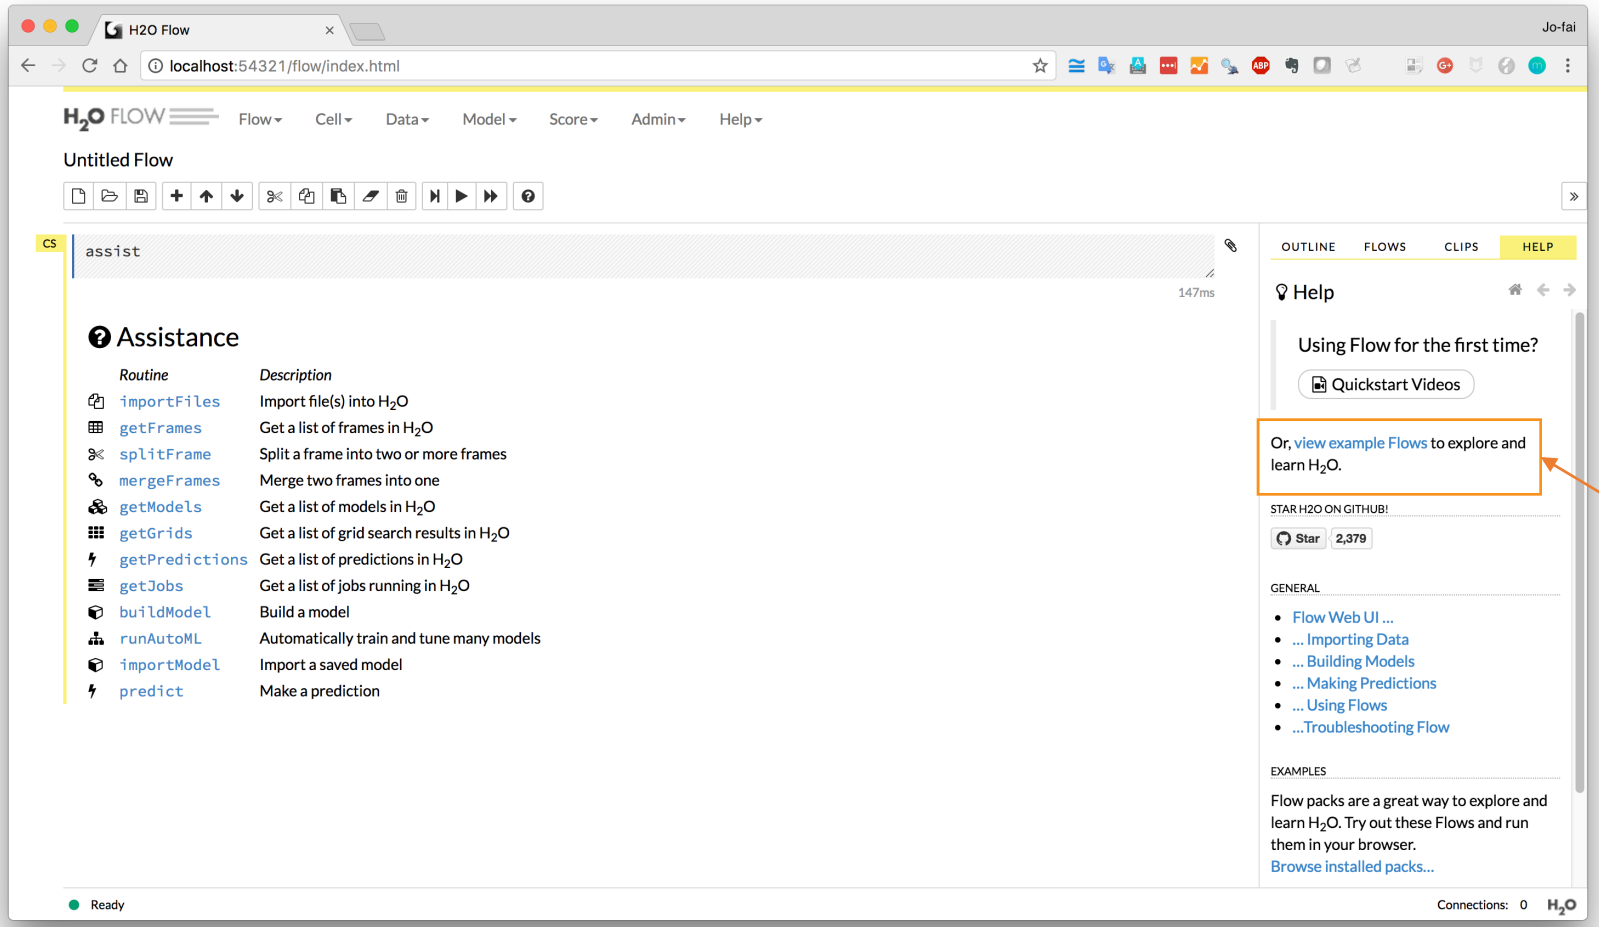

More Examples

H2O.ai

Install and Start H_2O in R / Python

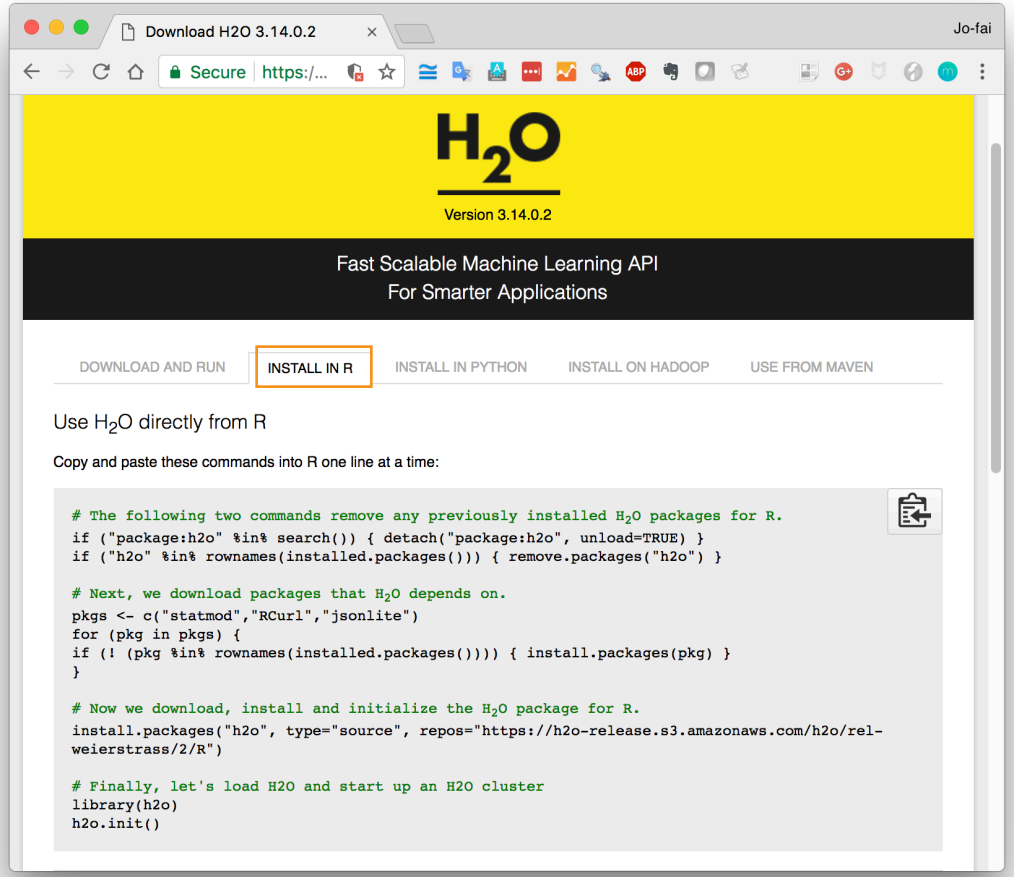

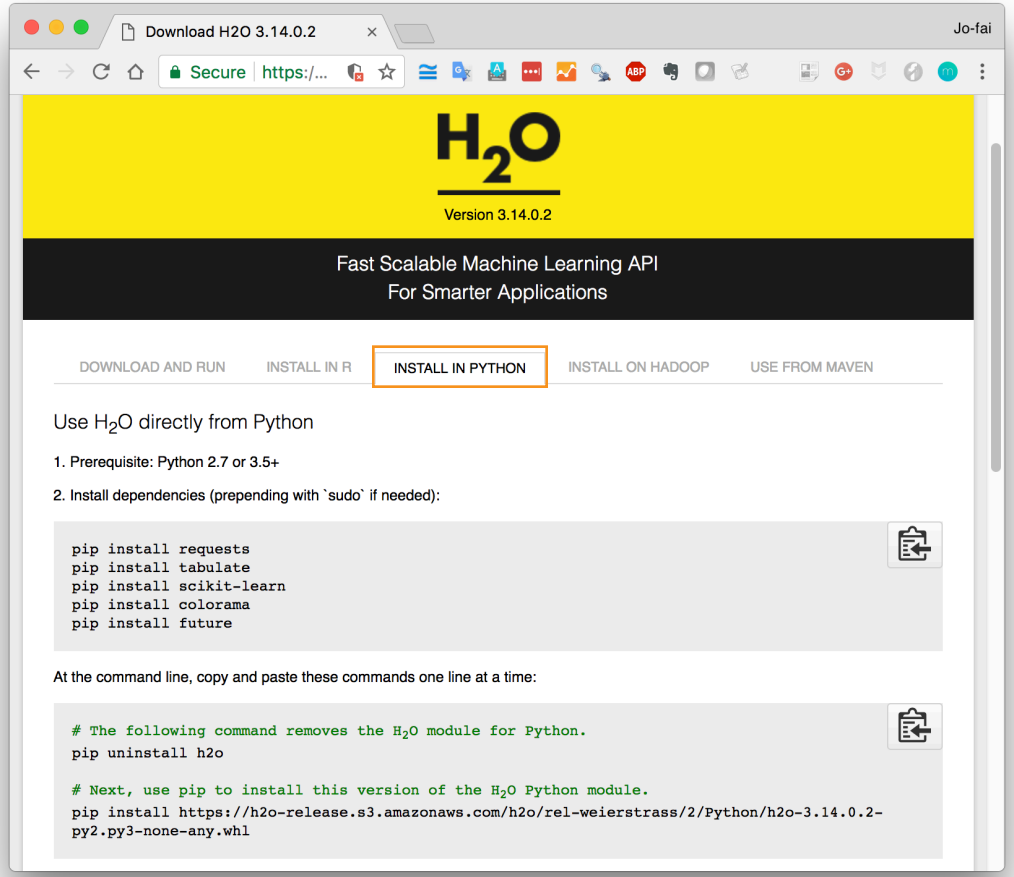

H2O.ai

Tutorial 1 - Regression

- **Data**: Boston Housing (1978)
- **Source**: http://www.cs.toronto.edu/~delve/data/boston/bostonDetail.html

Supervised Learning Example

Machine Learning:
Learn Patterns
from Data

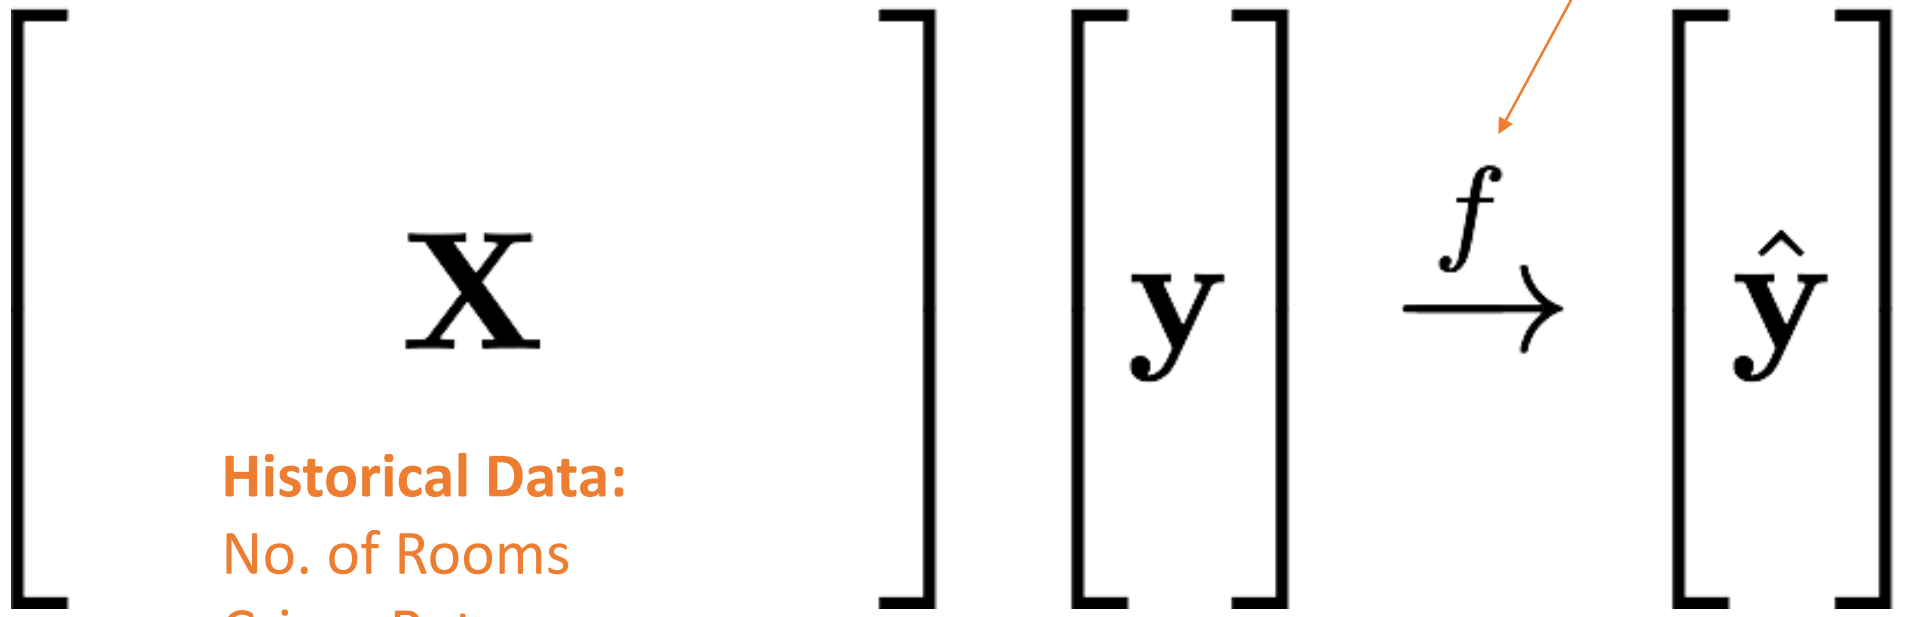

Historical Data:
No. of Rooms
Crime Rate
Pupil-Teacher Ratio
…

Target:
House Value

Predicted Value

H2O.ai

1. CRIM - per capita crime rate by town
2. ZN - proportion of residential land zoned for lots over 25,000 sq.ft.
3. INDUS - proportion of non-retail business acres per town.
4. CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)
5. NOX - nitric oxides concentration (parts per 10 million)
6. RM - average number of rooms per dwelling
7. AGE - proportion of owner-occupied units built prior to 1940
8. DIS - weighted distances to five Boston employment centres
9. RAD - index of accessibility to radial highways
10. TAX - full-value property-tax rate per $10,000
11. PTRATIO - pupil-teacher ratio by town
12. B - 1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town
13. LSTAT - % lower status of the population
14. MEDV - Median value of owner-occupied homes in $1000's

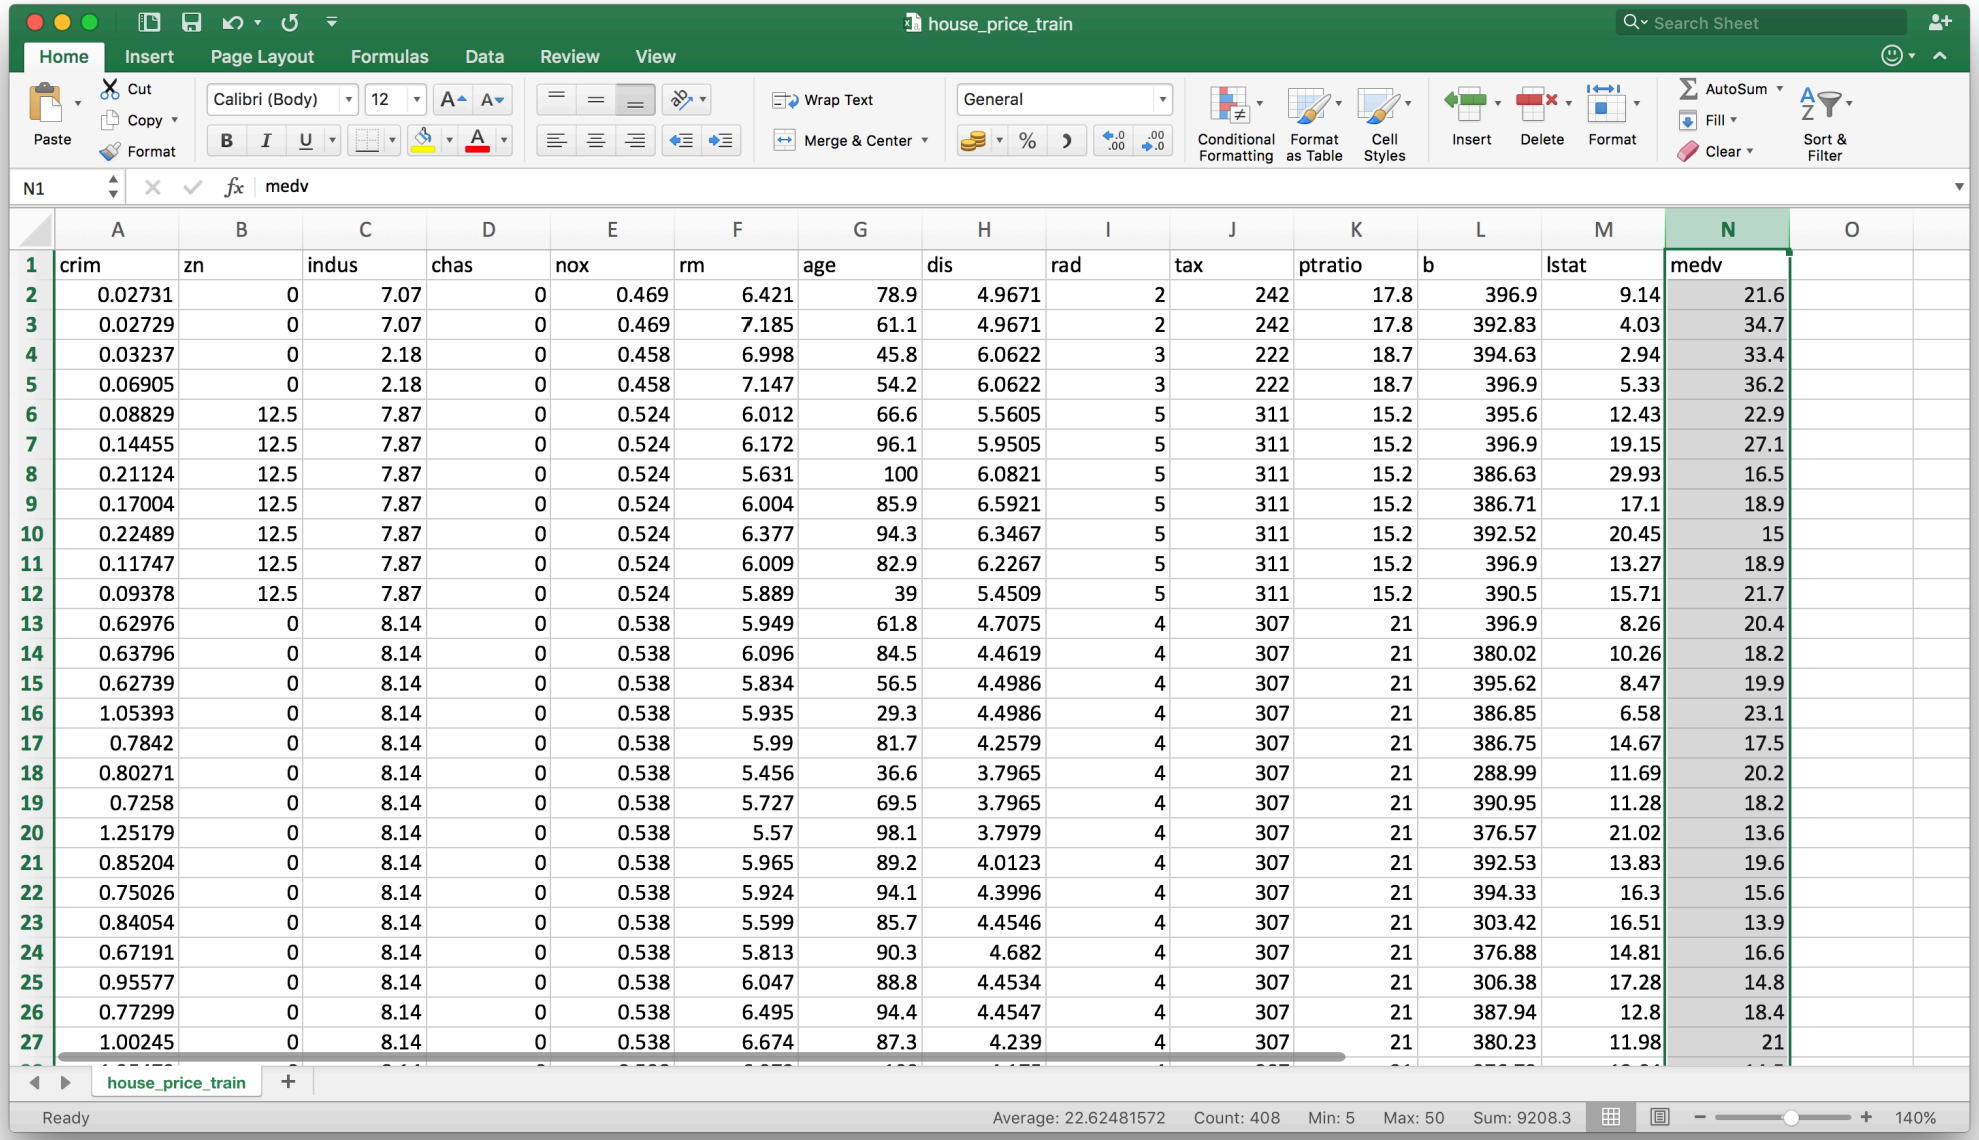

house_price_train

Search Sheet

Home Insert Page Layout Formulas Data Review View

Paste Cut Copy Format Calibri (Body) 12 Wrap Text General Conditional Formatting Format as Table Cell Styles Insert Delete Format AutoSum Fill Clear Sort & Filter

N1 medv

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv	
2	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.9	9.14	21.6	
3	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7	
4	0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4	
5	0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.9	5.33	36.2	
6	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.6	12.43	22.9	
7	0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311	15.2	396.9	19.15	27.1	
8	0.21124	12.5	7.87	0	0.524	5.631	100	6.0821	5	311	15.2	386.63	29.93	16.5	
9	0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311	15.2	386.71	17.1	18.9	
10	0.22489	12.5	7.87	0	0.524	6.377	94.3	6.3467	5	311	15.2	392.52	20.45	15	
11	0.11747	12.5	7.87	0	0.524	6.009	82.9	6.2267	5	311	15.2	396.9	13.27	18.9	
12	0.09378	12.5	7.87	0	0.524	5.889	39	5.4509	5	311	15.2	390.5	15.71	21.7	
13	0.62976	0	8.14	0	0.538	5.949	61.8	4.7075	4	307	21	396.9	8.26	20.4	
14	0.63796	0	8.14	0	0.538	6.096	84.5	4.4619	4	307	21	380.02	10.26	18.2	
15	0.62739	0	8.14	0	0.538	5.834	56.5	4.4986	4	307	21	395.62	8.47	19.9	
16	1.05393	0	8.14	0	0.538	5.935	29.3	4.4986	4	307	21	386.85	6.58	23.1	
17	0.7842	0	8.14	0	0.538	5.99	81.7	4.2579	4	307	21	386.75	14.67	17.5	
18	0.80271	0	8.14	0	0.538	5.456	36.6	3.7965	4	307	21	288.99	11.69	20.2	
19	0.7258	0	8.14	0	0.538	5.727	69.5	3.7965	4	307	21	390.95	11.28	18.2	
20	1.25179	0	8.14	0	0.538	5.57	98.1	3.7979	4	307	21	376.57	21.02	13.6	
21	0.85204	0	8.14	0	0.538	5.965	89.2	4.0123	4	307	21	392.53	13.83	19.6	
22	0.75026	0	8.14	0	0.538	5.924	94.1	4.3996	4	307	21	394.33	16.3	15.6	
23	0.84054	0	8.14	0	0.538	5.599	85.7	4.4546	4	307	21	303.42	16.51	13.9	
24	0.67191	0	8.14	0	0.538	5.813	90.3	4.682	4	307	21	376.88	14.81	16.6	
25	0.95577	0	8.14	0	0.538	6.047	88.8	4.4534	4	307	21	306.38	17.28	14.8	
26	0.77299	0	8.14	0	0.538	6.495	94.4	4.4547	4	307	21	387.94	12.8	18.4	
27	1.00245	0	8.14	0	0.538	6.674	87.3	4.239	4	307	21	380.23	11.98	21	

house_price_train

Ready Average: 22.62481572 Count: 408 Min: 5 Max: 50 Sum: 9208.3 140%

Supervised Learning – You Already Have Target Data

Regression:
How much will a customers spend?

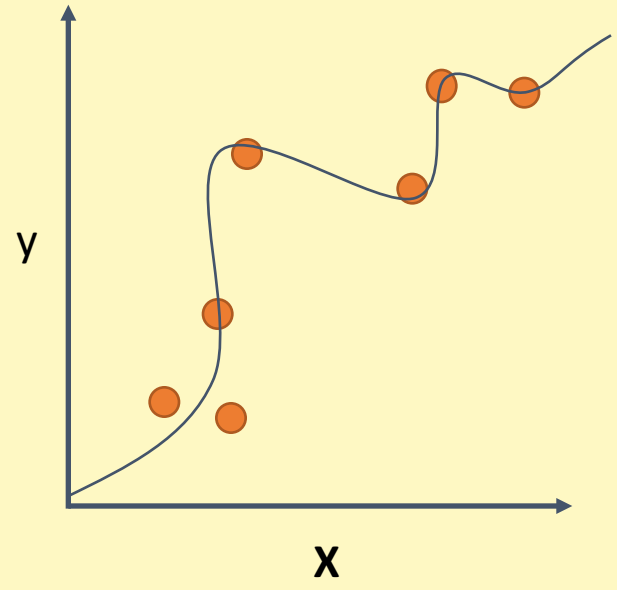

H_2O algos:
Penalized Linear Models
Random Forest
Gradient Boosting
Neural Networks
Stacked Ensembles

Classification:
Will a customer make a purchase? Yes or No

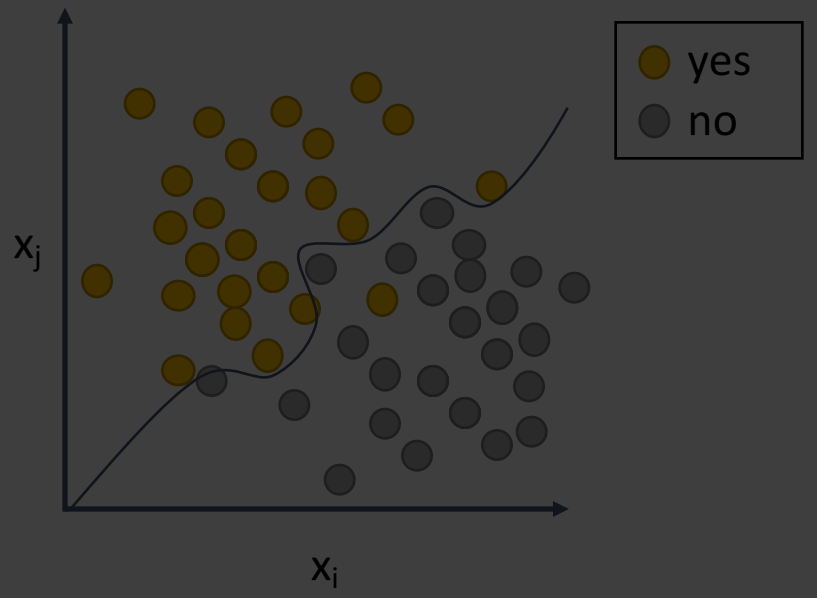

H_2O algos:
Penalized Linear Models
Naïve Bayes
Random Forest
Gradient Boosting
Neural Networks
Stacked Ensembles

H_2O.ai

Regression – Key Steps

- Import Data (CSV Files)
  - ./data/regression/ …
  - Have a quick look
- Train a Random Forest Model
  - Look at variable importance
- Make Predictions
  - Compare with ground truth
- Build Partial Dependence Plots
  - Explain variable

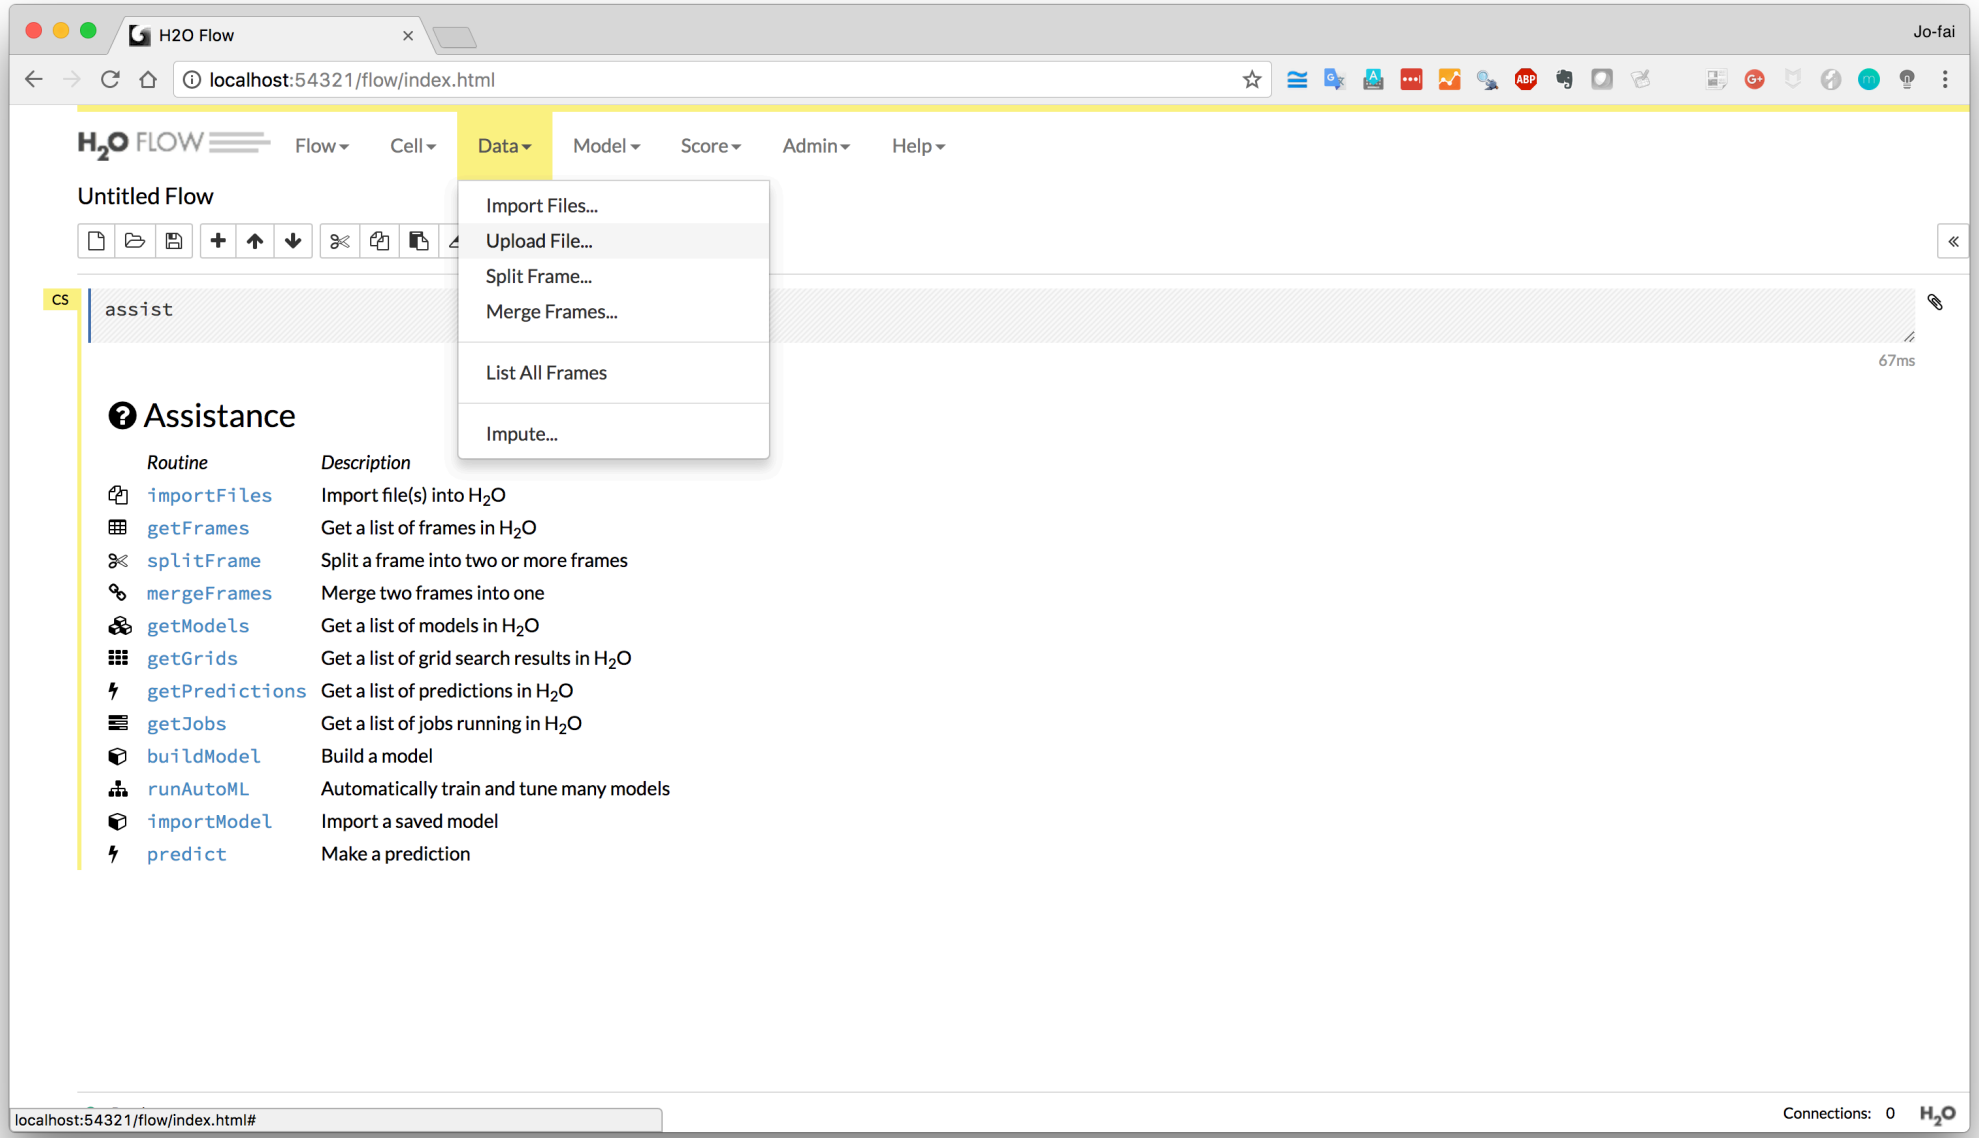
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Import Files...
Upload File...
Split Frame...
Merge Frames...
List All Frames
Impute...
CS
assist
67ms
Assistance
Routine
Description
importFiles
Import file(s) into H2O
getFrames
Get a list of frames in H2O
splitFrame
Split a frame into two or more frames
mergeFrames
Merge two frames into one
getModels
Get a list of models in H2O
getGrids
Get a list of grid search results in H2O
getPredictions
Get a list of predictions in H2O
getJobs
Get a list of jobs running in H2O
buildModel
Build a model
runAutoML
Automatically train and tune many models
importModel
Import a saved model
predict
Make a prediction
localhost:54321/flow/index.html#
Connections: 0
H2O

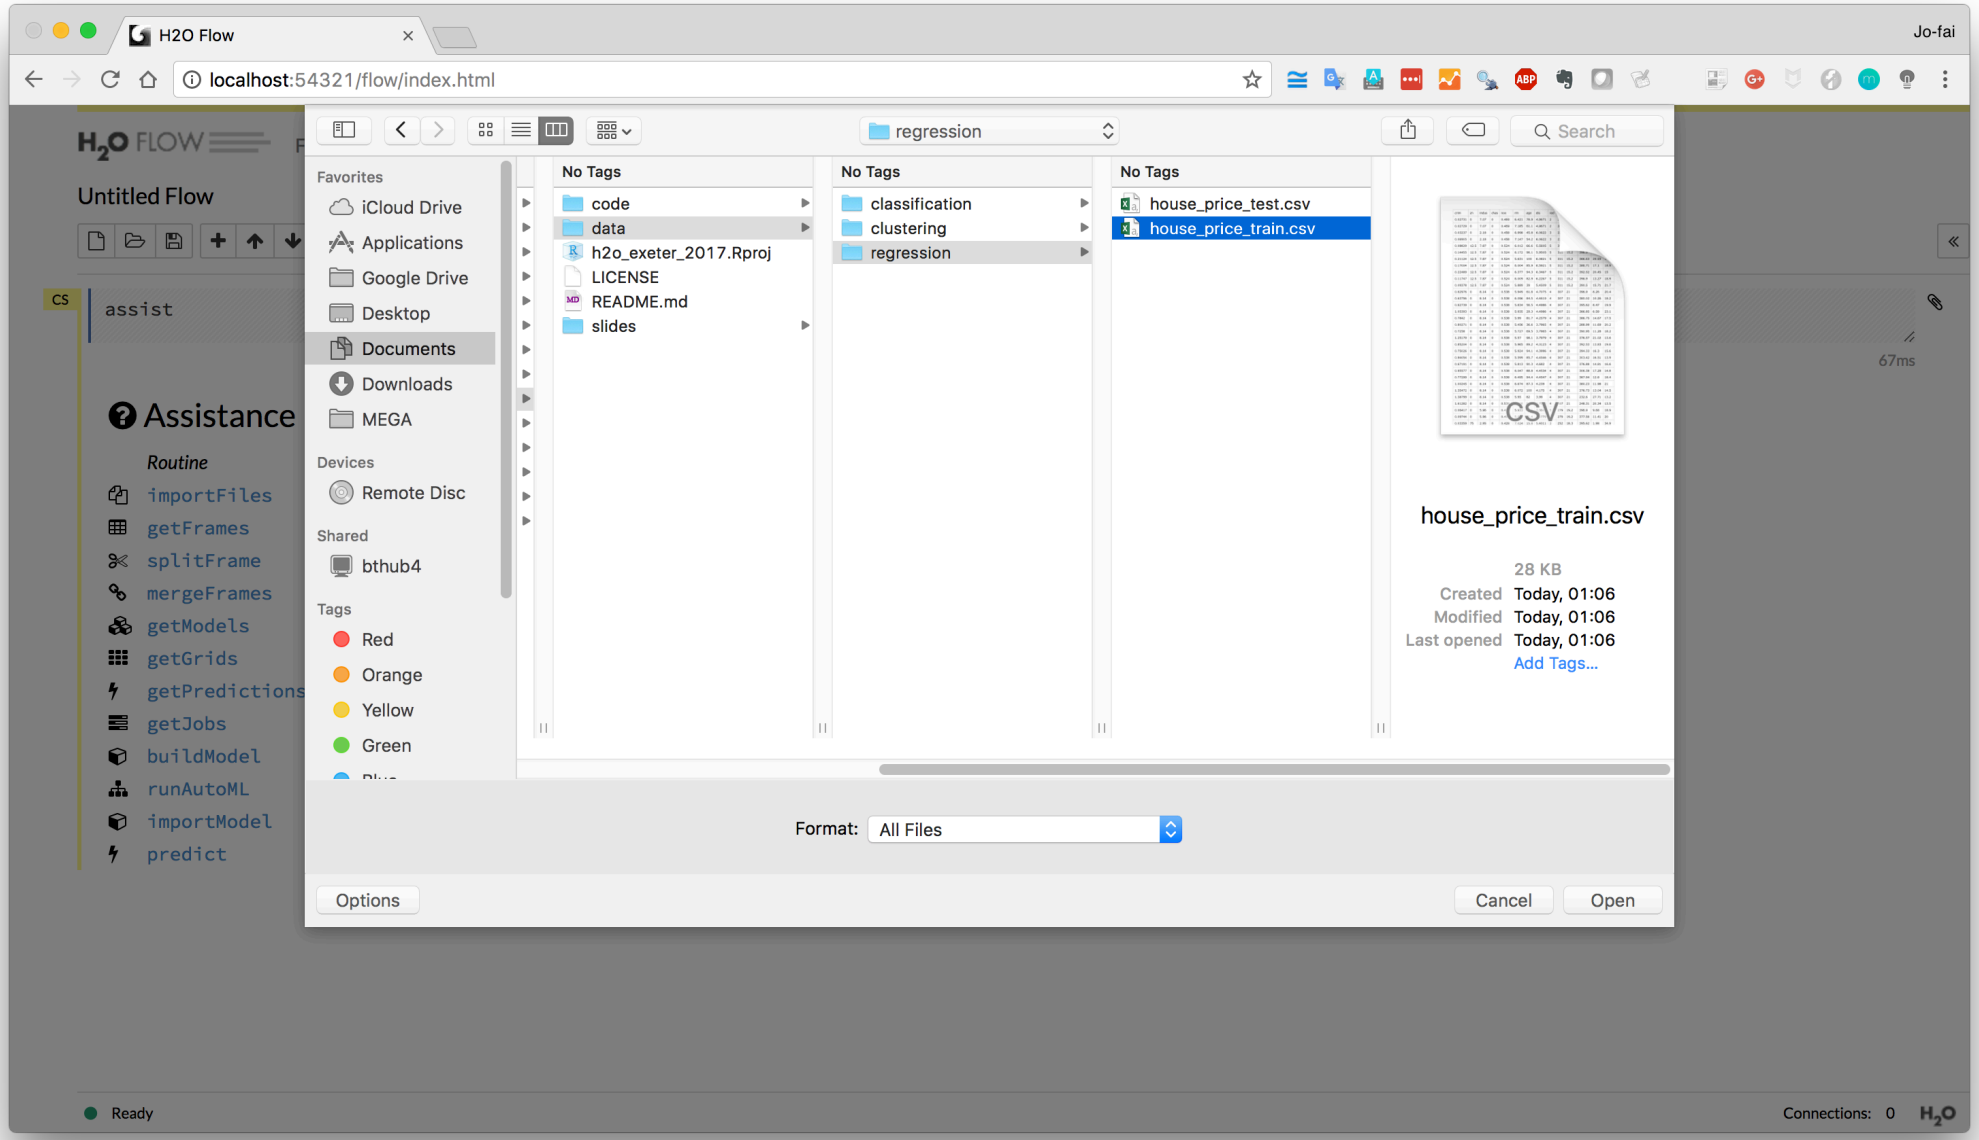
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Untitled Flow
assist
Assistance
Routine
importFiles
getFrames
splitFrame
mergeFrames
getModels
getGrids
getPredictions
getJobs
buildModel
runAutoML
importModel
predict
67ms
Ready
Connections: 0
regression
Search
Favorites
iCloud Drive
Applications
Google Drive
Desktop
Documents
Downloads
MEGA
Devices
Remote Disc
Shared
bthub4
Tags
Red
Orange
Yellow
Green
No Tags
code
data
h2o_exeter_2017.Rproj
LICENSE
README.md
slides
No Tags
classification
clustering
regression
No Tags
house_price_test.csv
house_price_train.csv
CSV
house_price_train.csv
28 KB
Created Today, 01:06
Modified Today, 01:06
Last opened Today, 01:06
Add Tags...
Format: All Files
Options
Cancel
Open

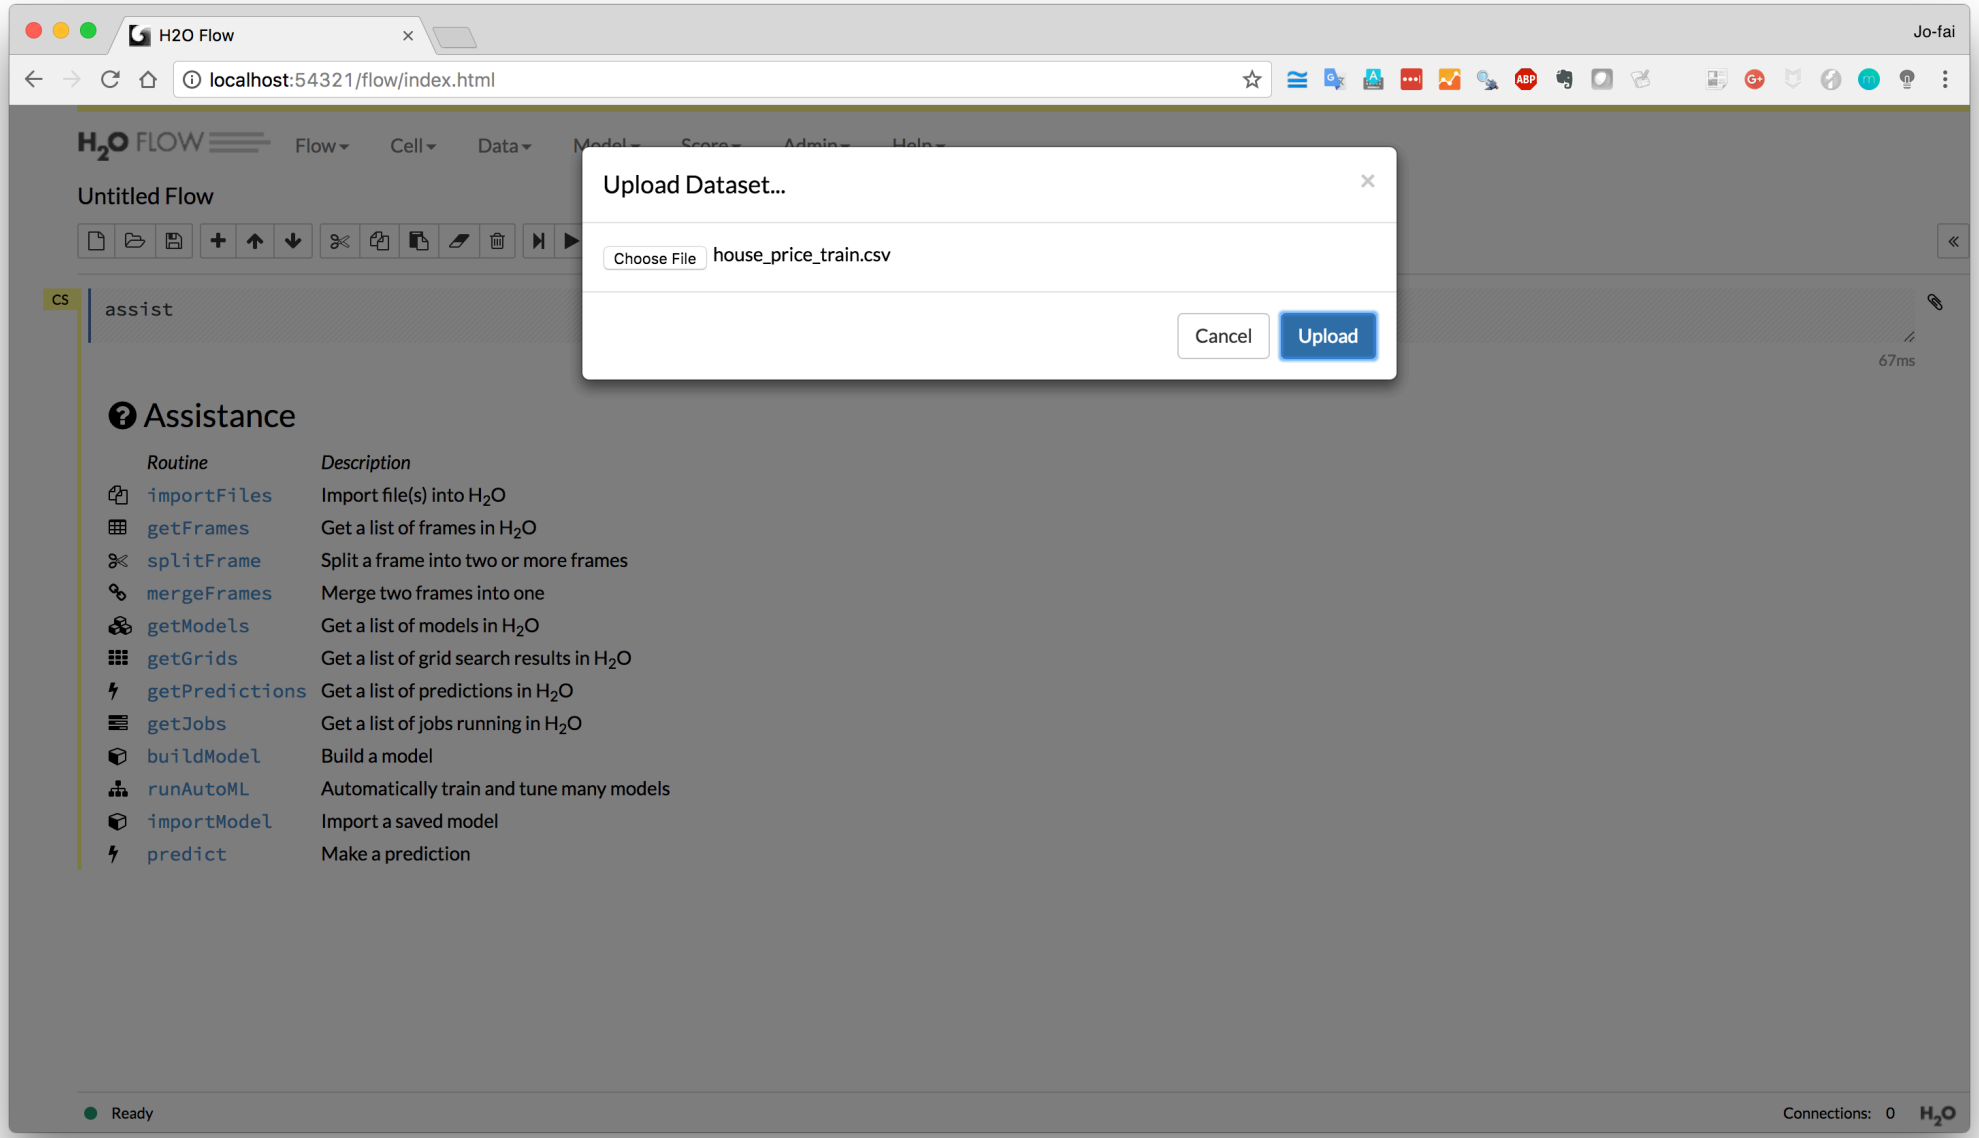
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Upload Dataset...
Choose File
house_price_train.csv
Cancel
Upload
CS
assist
67ms
Assistance
Routine
Description
importFiles
Import file(s) into H2O
getFrames
Get a list of frames in H2O
splitFrame
Split a frame into two or more frames
mergeFrames
Merge two frames into one
getModels
Get a list of models in H2O
getGrids
Get a list of grid search results in H2O
getPredictions
Get a list of predictions in H2O
getJobs
Get a list of jobs running in H2O
buildModel
Build a model
runAutoML
Automatically train and tune many models
importModel
Import a saved model
predict
Make a prediction
Ready
Connections: 0
H2O

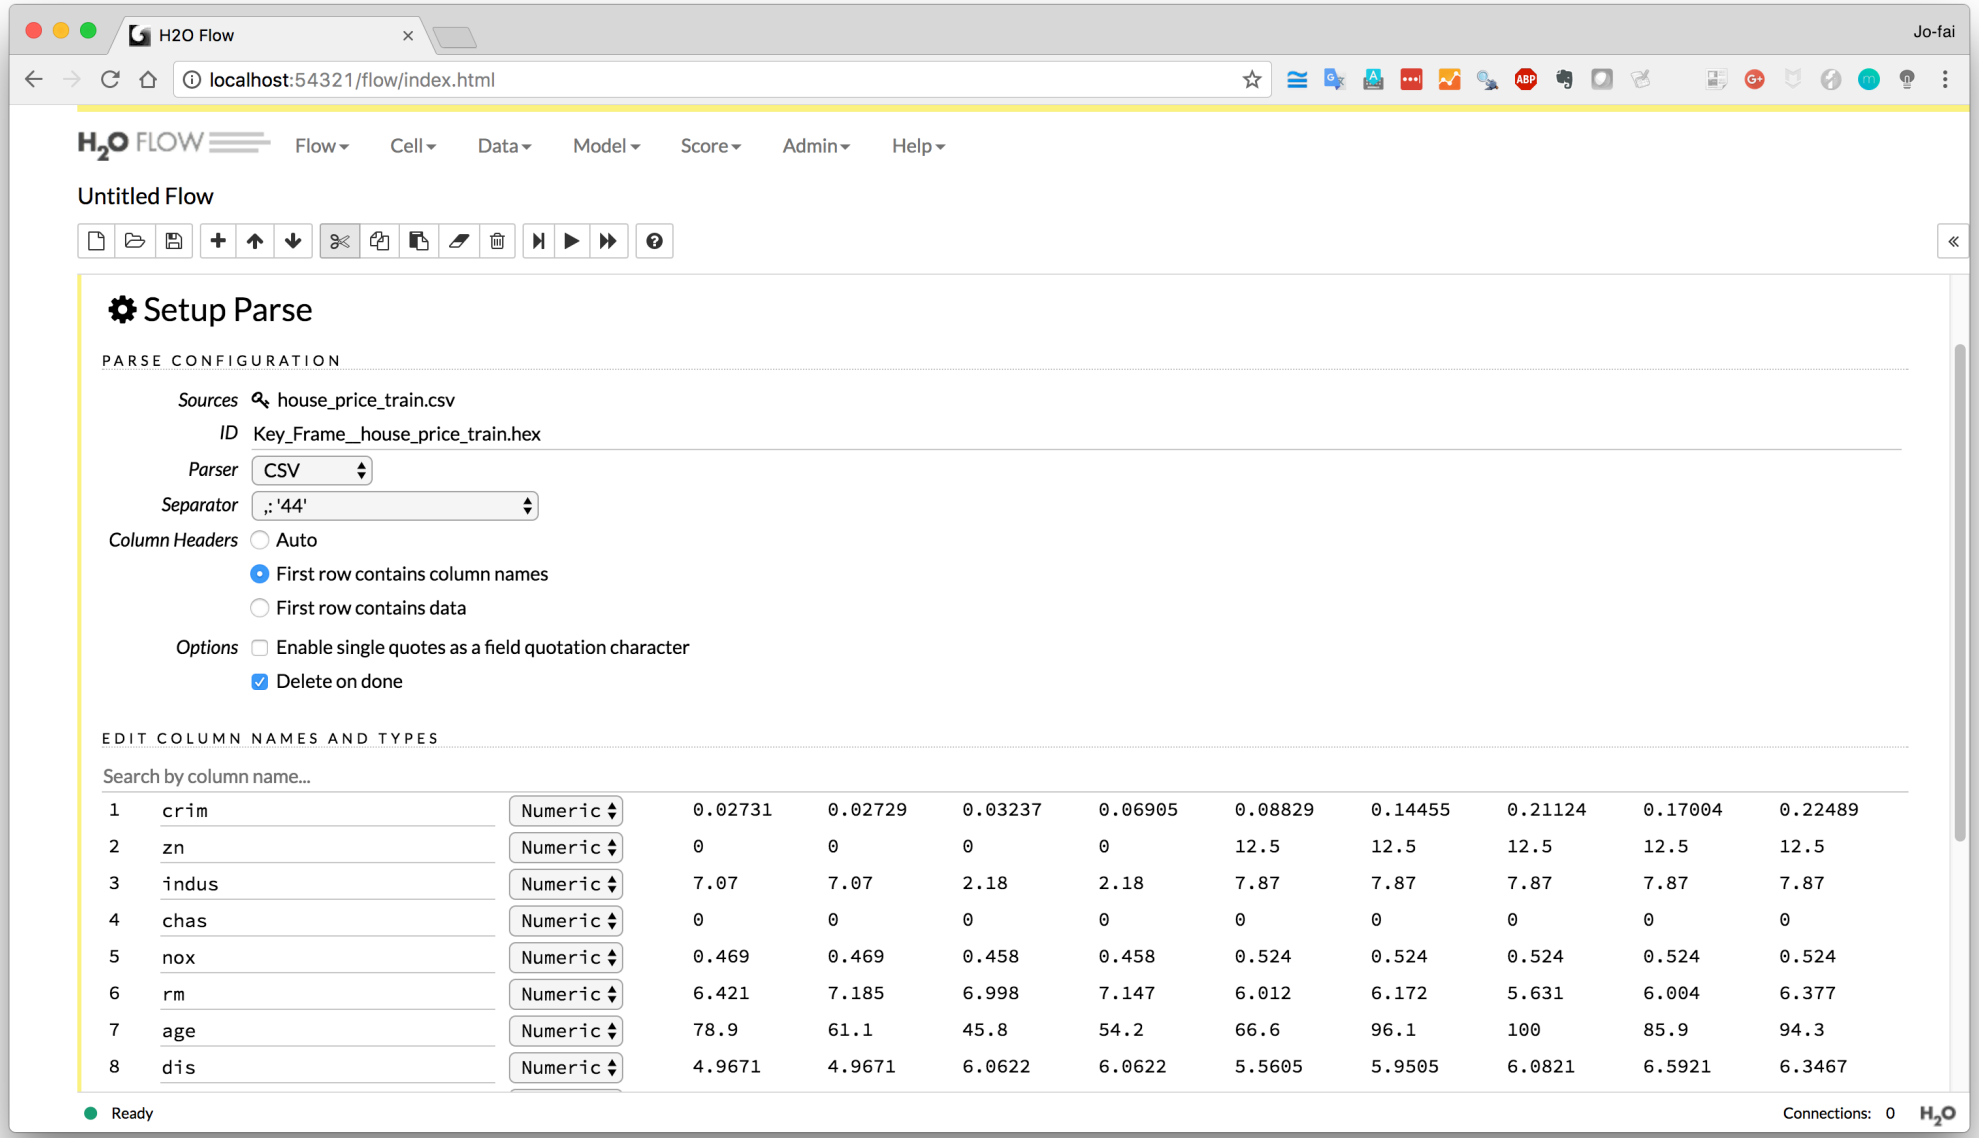
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Setup Parse
PARSE CONFIGURATION
Sources house_price_train.csv
ID Key_Frame__house_price_train.hex
Parser CSV
Separator ,: '44'
Column Headers Auto
First row contains column names
First row contains data
Options Enable single quotes as a field quotation character
Delete on done
EDIT COLUMN NAMES AND TYPES
Search by column name...
1 crim Numeric 0.02731 0.02729 0.03237 0.06905 0.08829 0.14455 0.21124 0.17004 0.22489
2 zn Numeric 0 0 0 0 12.5 12.5 12.5 12.5 12.5
3 indus Numeric 7.07 7.07 2.18 2.18 7.87 7.87 7.87 7.87 7.87
4 chas Numeric 0 0 0 0 0 0 0 0 0
5 nox Numeric 0.469 0.469 0.458 0.458 0.524 0.524 0.524 0.524 0.524
6 rm Numeric 6.421 7.185 6.998 7.147 6.012 6.172 5.631 6.004 6.377
7 age Numeric 78.9 61.1 45.8 54.2 66.6 96.1 100 85.9 94.3
8 dis Numeric 4.9671 4.9671 6.0622 6.0622 5.5605 5.9505 6.0821 6.5921 6.3467
Ready
Connections: 0
H2O

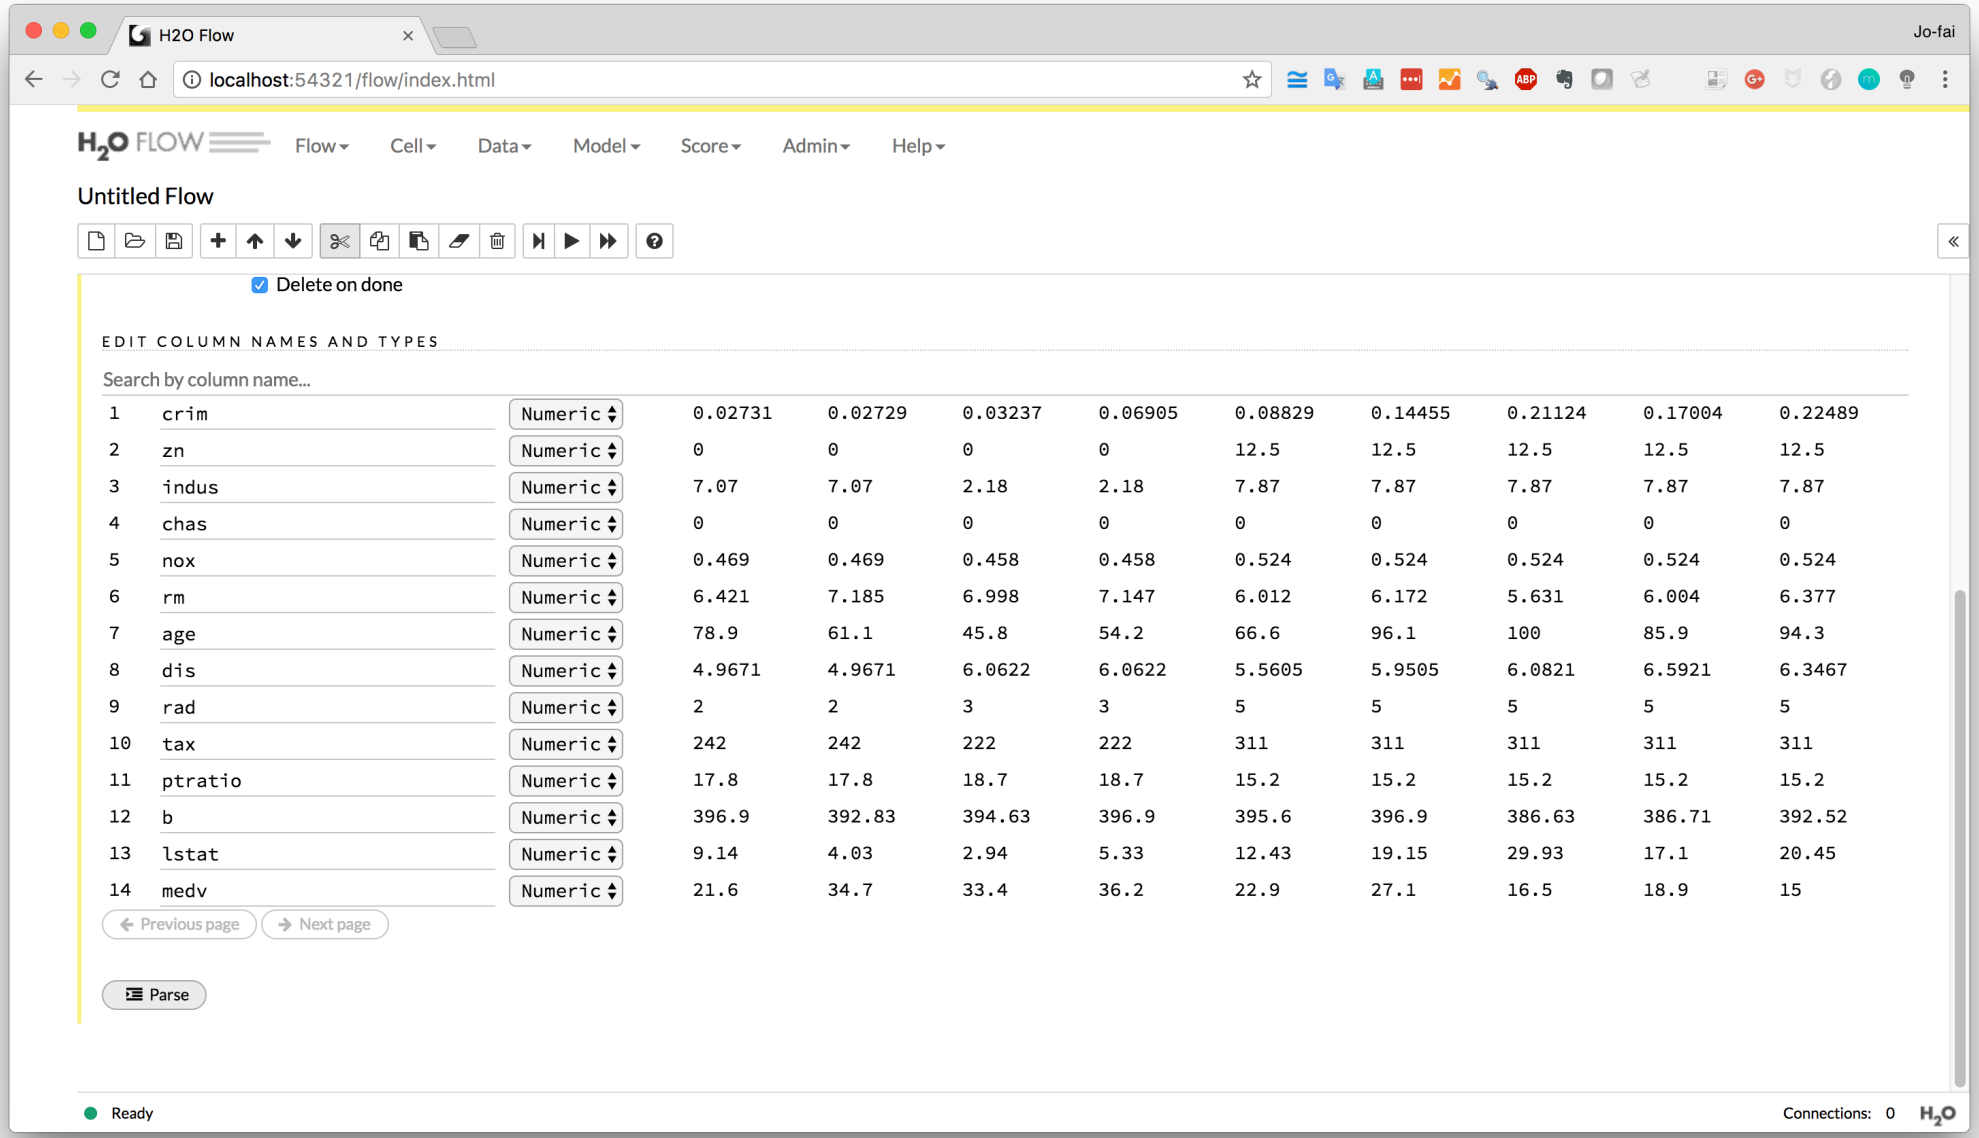
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Delete on done
EDIT COLUMN NAMES AND TYPES
Search by column name...
1 crim Numeric 0.02731 0.02729 0.03237 0.06905 0.08829 0.14455 0.21124 0.17004 0.22489
2 zn Numeric 0 0 0 0 12.5 12.5 12.5 12.5 12.5
3 indus Numeric 7.07 7.07 2.18 2.18 7.87 7.87 7.87 7.87 7.87
4 chas Numeric 0 0 0 0 0 0 0 0 0
5 nox Numeric 0.469 0.469 0.458 0.458 0.524 0.524 0.524 0.524 0.524
6 rm Numeric 6.421 7.185 6.998 7.147 6.012 6.172 5.631 6.004 6.377
7 age Numeric 78.9 61.1 45.8 54.2 66.6 96.1 100 85.9 94.3
8 dis Numeric 4.9671 4.9671 6.0622 6.0622 5.5605 5.9505 6.0821 6.5921 6.3467
9 rad Numeric 2 2 3 3 5 5 5 5 5
10 tax Numeric 242 242 222 222 311 311 311 311 311
11 ptratio Numeric 17.8 17.8 18.7 18.7 15.2 15.2 15.2 15.2 15.2
12 b Numeric 396.9 392.83 394.63 396.9 395.6 396.9 386.63 386.71 392.52
13 lstat Numeric 9.14 4.03 2.94 5.33 12.43 19.15 29.93 17.1 20.45
14 medv Numeric 21.6 34.7 33.4 36.2 22.9 27.1 16.5 18.9 15
Previous page
Next page
Parse
Ready
Connections: 0
H2O

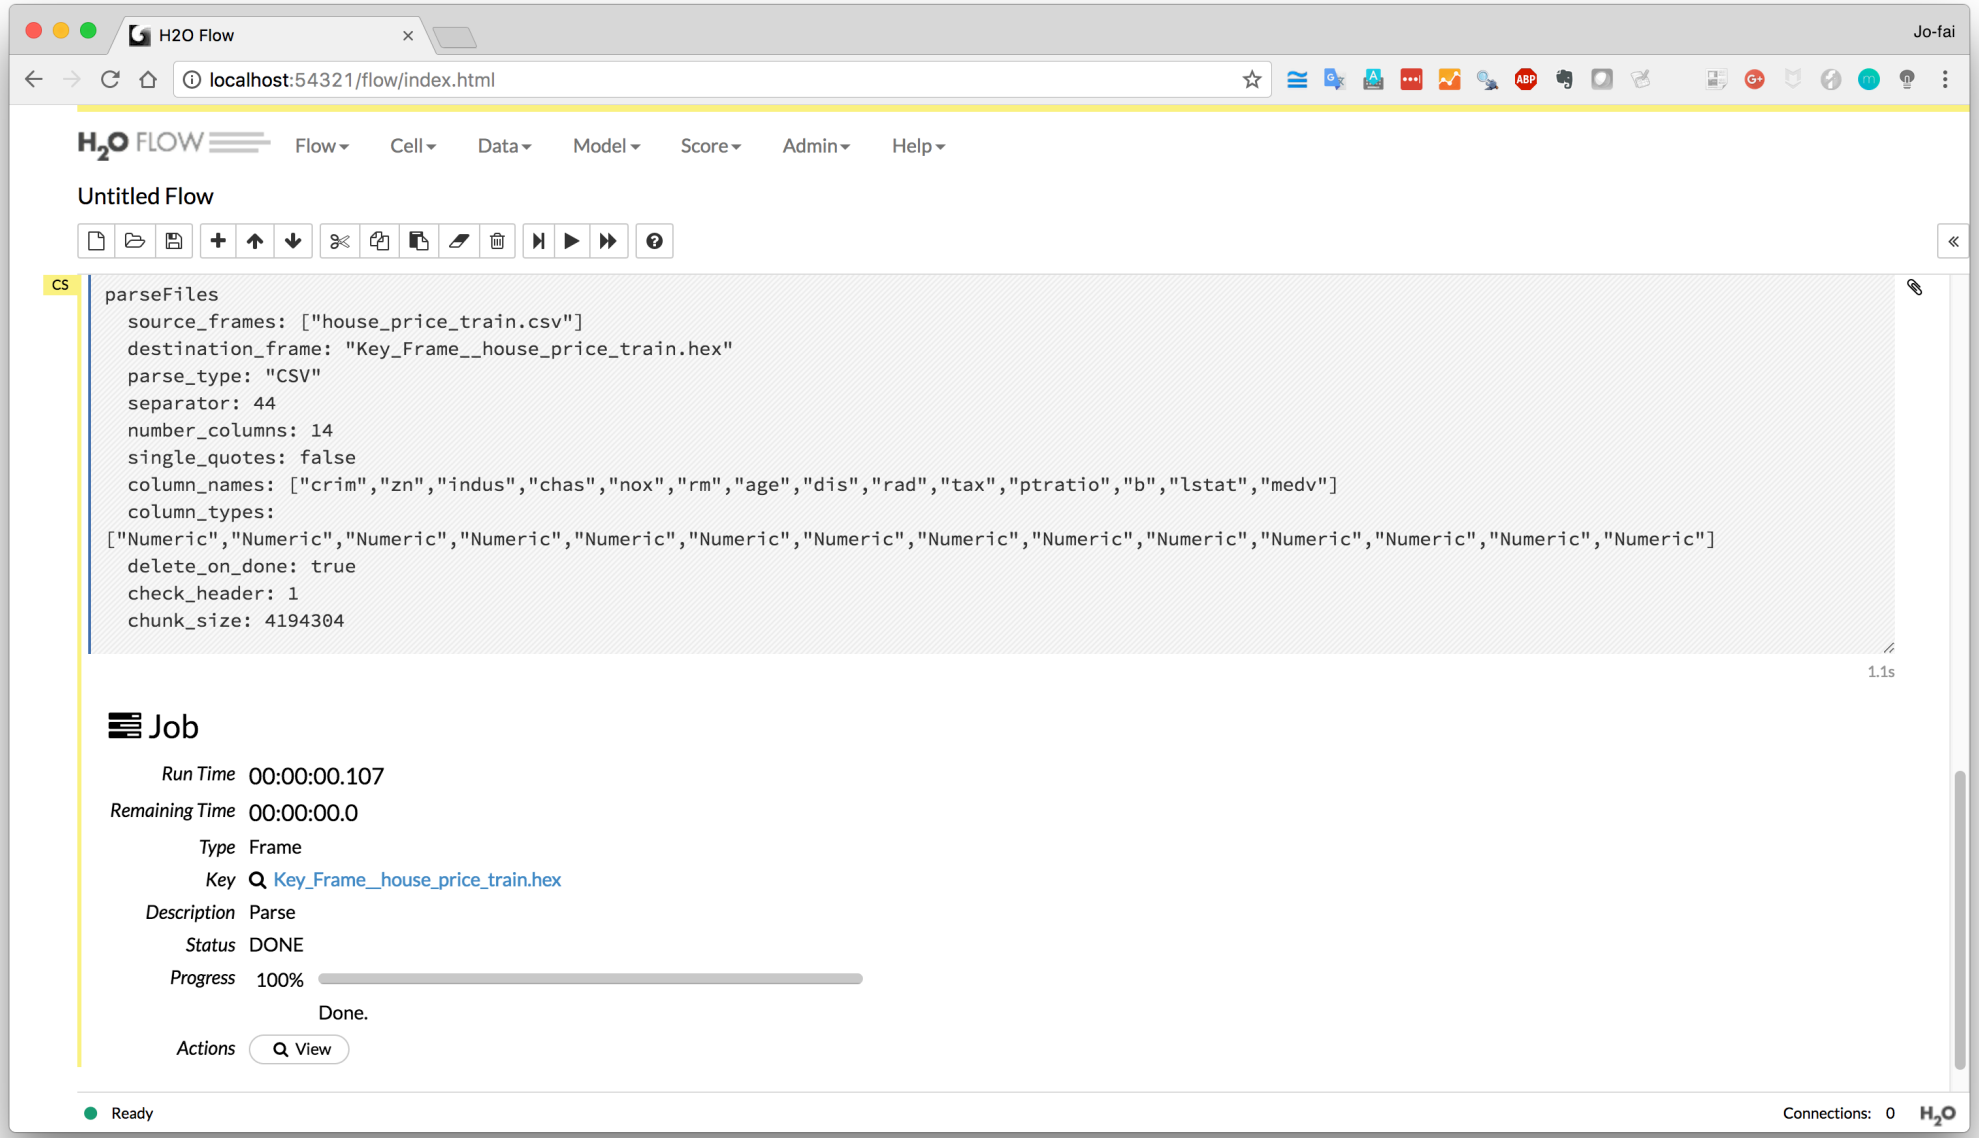
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
CS
parseFiles
source_frames: ["house_price_train.csv"]
destination_frame: "Key_Frame__house_price_train.hex"
parse_type: "CSV"
separator: 44
number_columns: 14
single_quotes: false
column_names: ["crim","zn","indus","chas","nox","rm","age","dis","rad","tax","ptratio","b","lstat","medv"]
column_types:
["Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric","Numeric"]
delete_on_done: true
check_header: 1
chunk_size: 4194304
1.1s
Job
Run Time 00:00:00.107
Remaining Time 00:00:00.0
Type Frame
Key Key_Frame__house_price_train.hex
Description Parse
Status DONE
Progress 100%
Done.
Actions View
Ready
Connections: 0
H2O

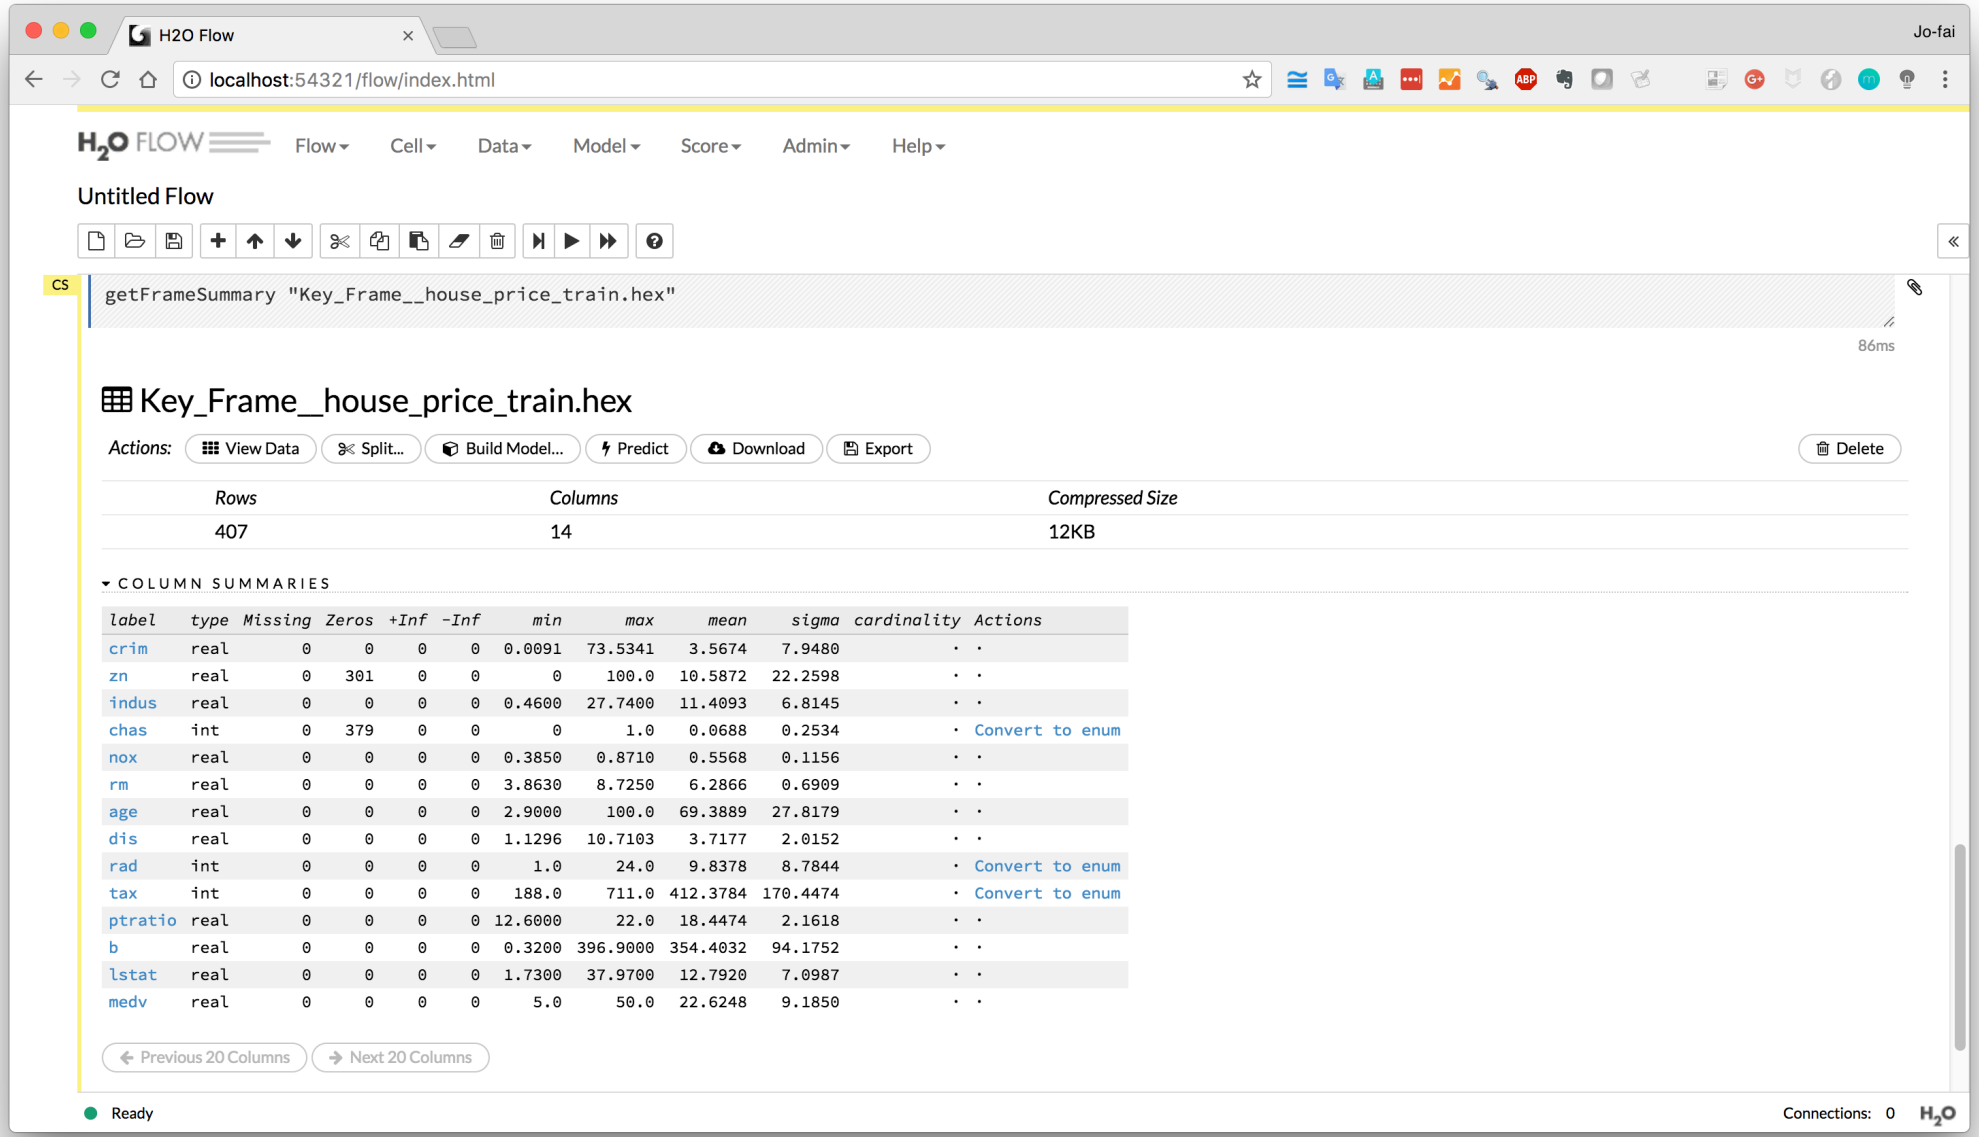
H2O Flow
Jo-fai
localhost:54321/flow/index.html
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
CS
getFrameSummary "Key_Frame__house_price_train.hex"
86ms
Key_Frame__house_price_train.hex
Actions:
View Data
Split...
Build Model...
Predict
Download
Export
Delete
Rows
Columns
Compressed Size
407
14
12KB
COLUMN SUMMARIES
label type Missing Zeros +Inf -Inf min max mean sigma cardinality Actions
crim real 0 0 0 0 0.0091 73.5341 3.5674 7.9480 · ·
zn real 0 301 0 0 0 100.0 10.5872 22.2598 · ·
indus real 0 0 0 0 0.4600 27.7400 11.4093 6.8145 · ·
chas int 0 379 0 0 0 1.0 0.0688 0.2534 · Convert to enum
nox real 0 0 0 0 0.3850 0.8710 0.5568 0.1156 · ·
rm real 0 0 0 0 3.8630 8.7250 6.2866 0.6909 · ·
age real 0 0 0 0 2.9000 100.0 69.3889 27.8179 · ·
dis real 0 0 0 0 1.1296 10.7103 3.7177 2.0152 · ·
rad int 0 0 0 0 1.0 24.0 9.8378 8.7844 · Convert to enum
tax int 0 0 0 0 188.0 711.0 412.3784 170.4474 · Convert to enum
ptratio real 0 0 0 0 12.6000 22.0 18.4474 2.1618 · ·
b real 0 0 0 0 0.3200 396.9000 354.4032 94.1752 · ·
lstat real 0 0 0 0 1.7300 37.9700 12.7920 7.0987 · ·
medv real 0 0 0 0 5.0 50.0 22.6248 9.1850 · ·
Previous 20 Columns
Next 20 Columns
Ready
Connections: 0
H2O

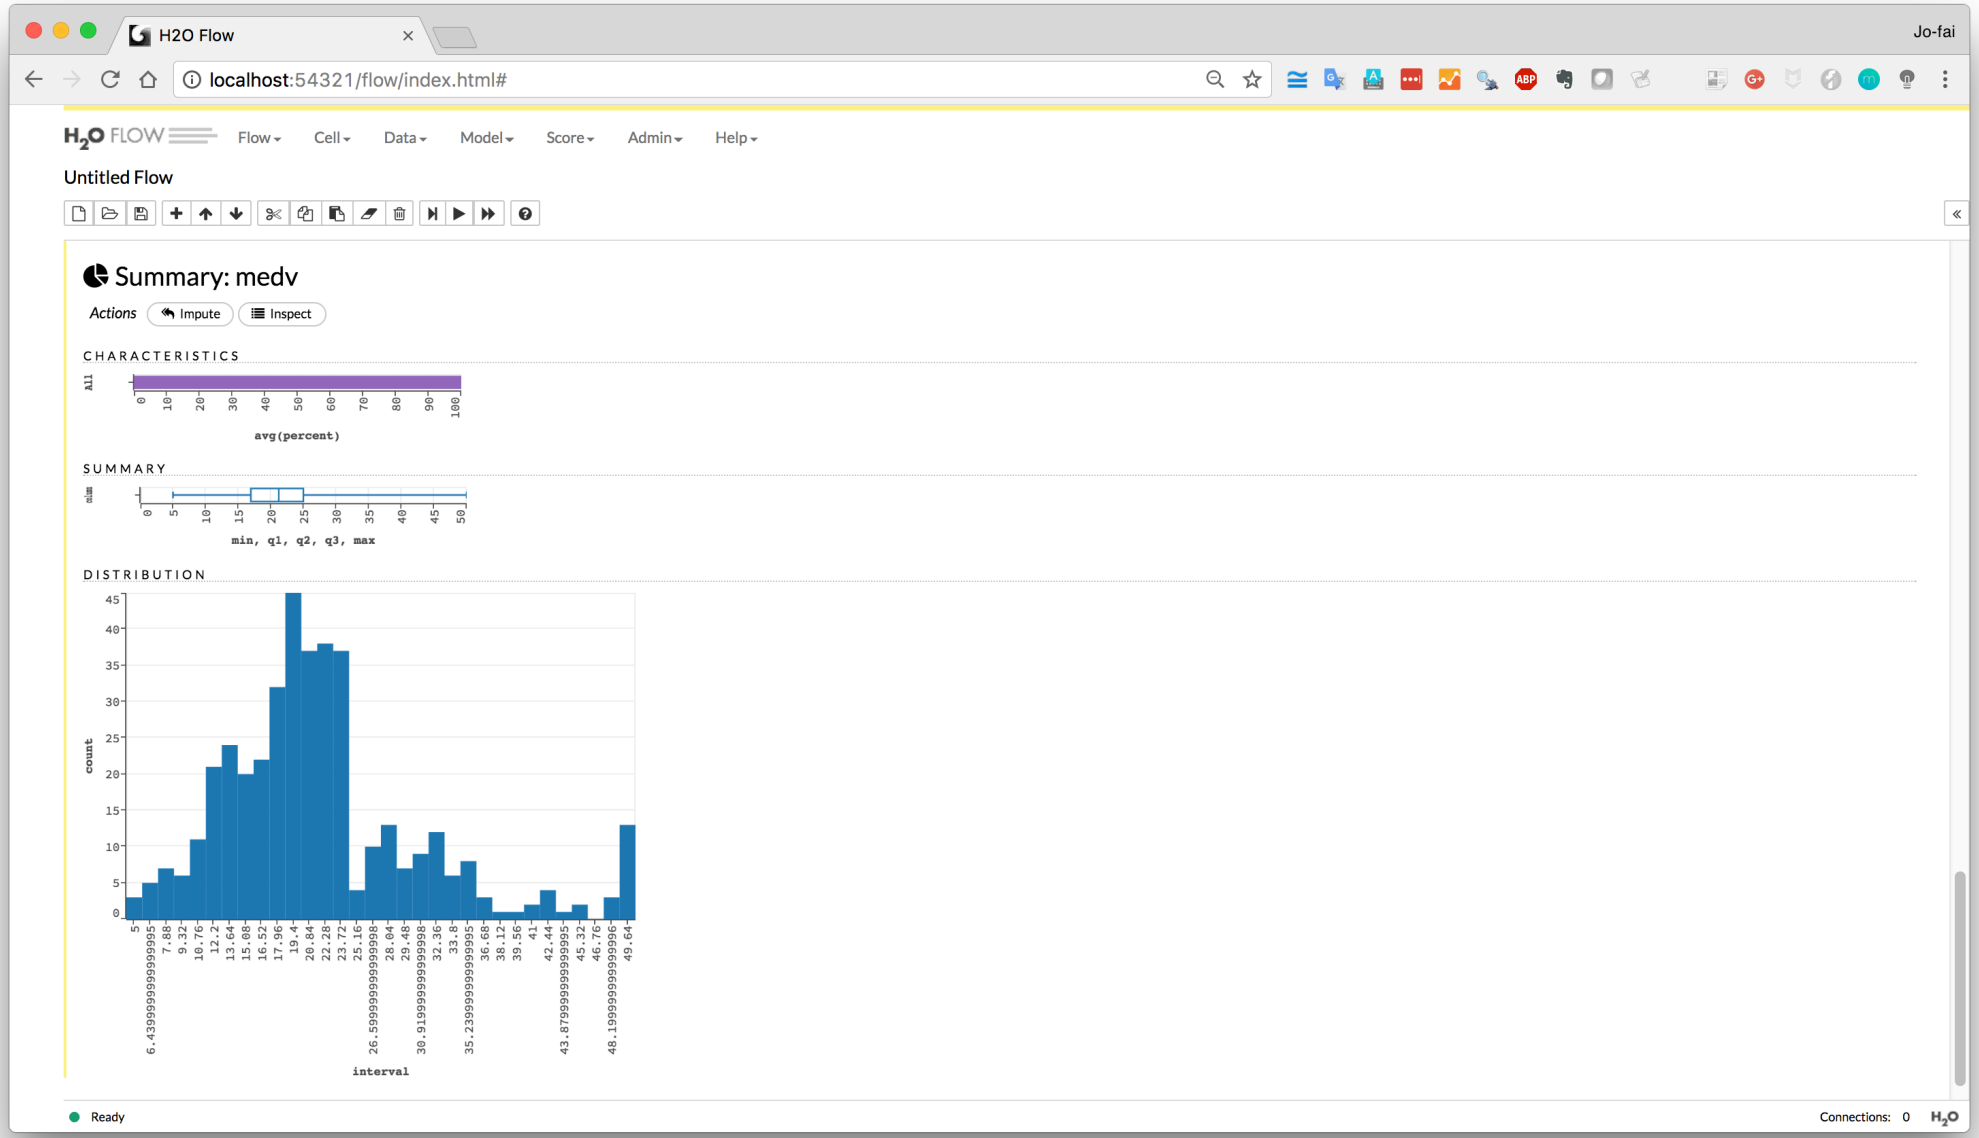
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Summary: medv
Actions
Impute
Inspect
CHARACTERISTICS
All
0 10 20 30 40 50 60 70 80 90 100
avg(percent)
SUMMARY
min
0 5 10 15 20 25 30 35 40 45 50
min, q1, q2, q3, max
DISTRIBUTION
count
0 5 10 15 20 25 30 35 40 45
6.439999999999995 7.88 9.32 10.76 12.2 13.64 15.08 16.52 17.96 19.4 20.84 22.28 23.72 25.16 26.599999999999998 28.04 29.48 30.919999999999998 32.36 33.8 35.239999999999995 36.68 38.12 39.56 41 42.44 43.879999999999995 45.32 46.76 48.199999999999996 49.64
interval
Ready
Connections: 0
H2O

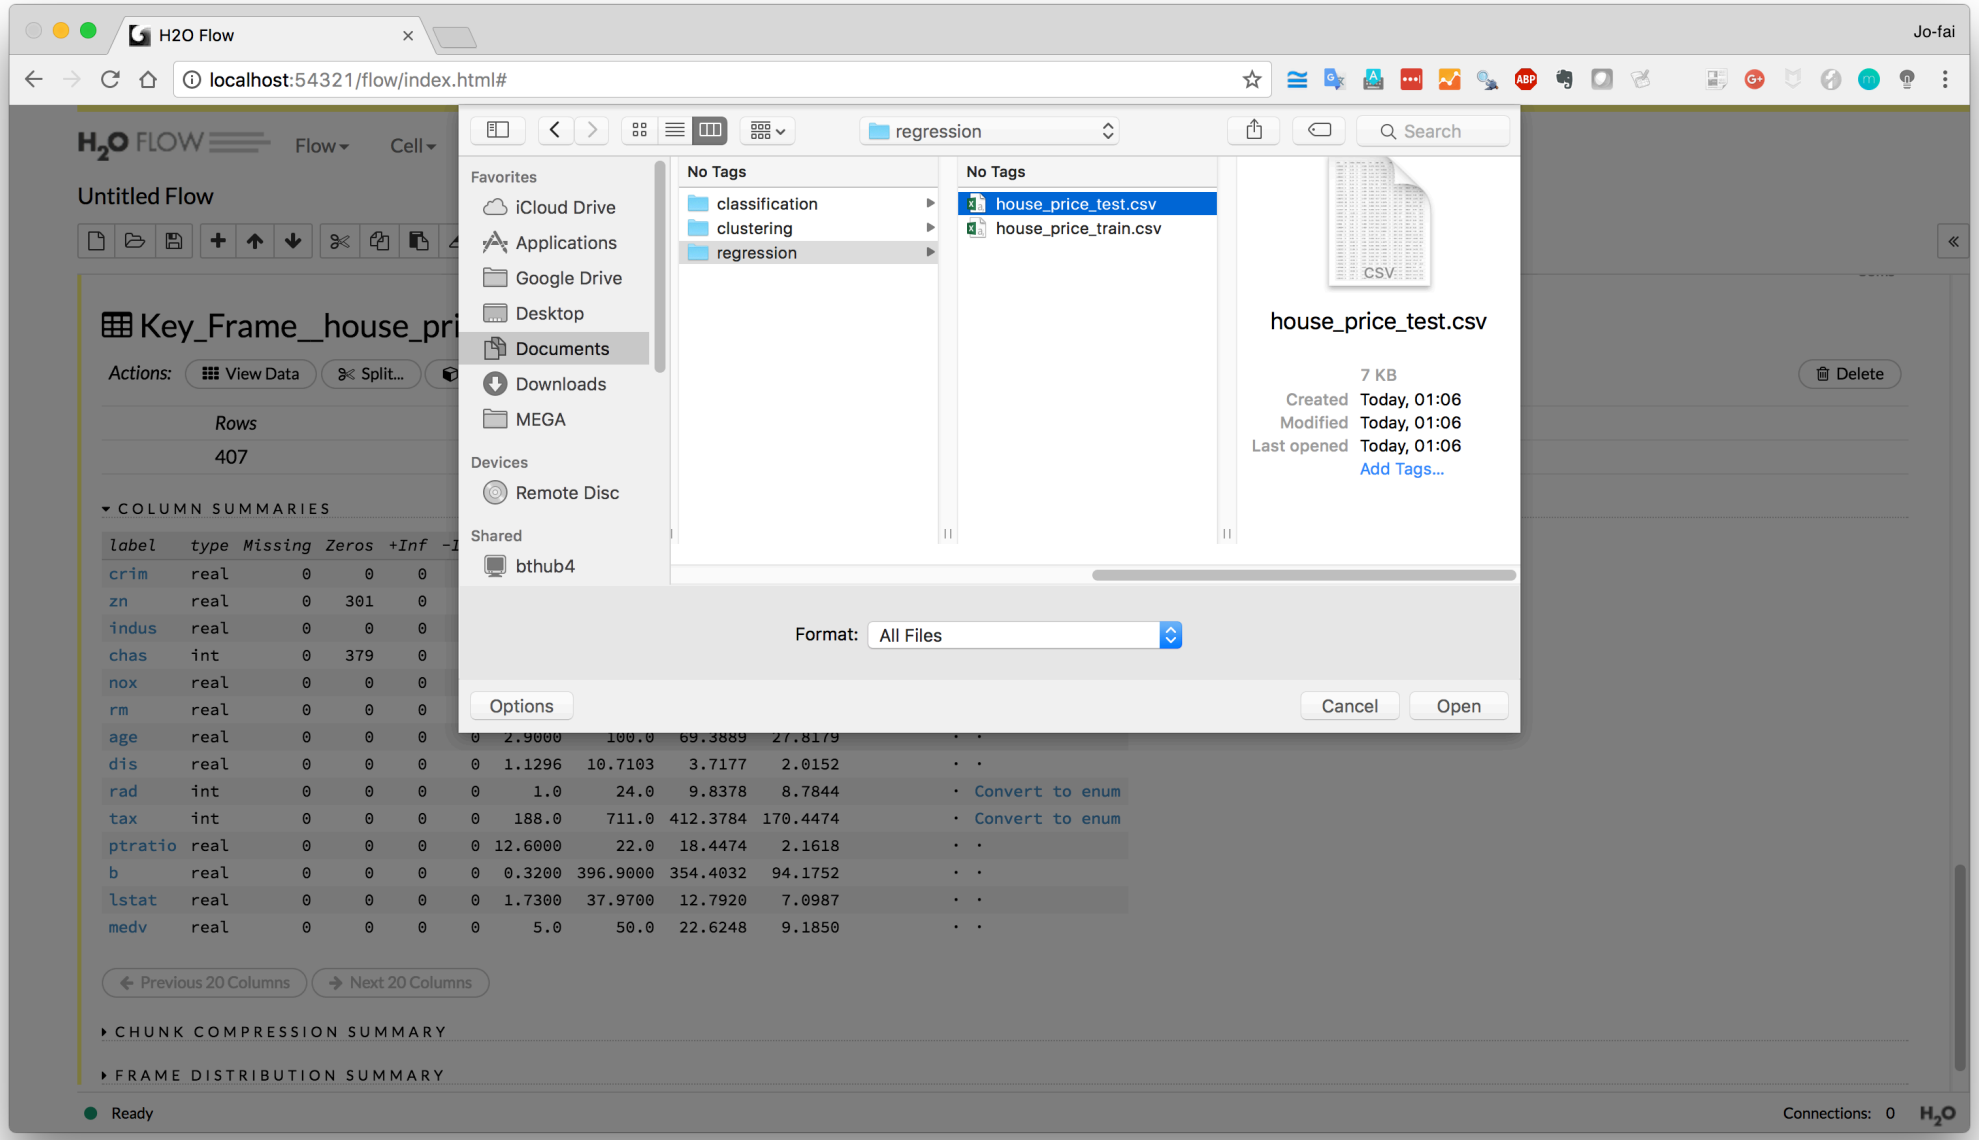
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Untitled Flow
Key_Frame__house_pri
Actions:
View Data
Split...
Delete
Rows
407
COLUMN SUMMARIES
label type Missing Zeros +Inf
crim real 0 0 0
zn real 0 301 0
indus real 0 0 0
chas int 0 379 0
nox real 0 0 0
rm real 0 0 0
age real 0 0 0 0 2.9000 100.0 69.3889 27.8179
dis real 0 0 0 0 1.1296 10.7103 3.7177 2.0152
rad int 0 0 0 0 1.0 24.0 9.8378 8.7844 Convert to enum
tax int 0 0 0 0 188.0 711.0 412.3784 170.4474 Convert to enum
ptratio real 0 0 0 0 12.6000 22.0 18.4474 2.1618
b real 0 0 0 0 0.3200 396.9000 354.4032 94.1752
lstat real 0 0 0 0 1.7300 37.9700 12.7920 7.0987
medv real 0 0 0 0 5.0 50.0 22.6248 9.1850
Previous 20 Columns
Next 20 Columns
CHUNK COMPRESSION SUMMARY
FRAME DISTRIBUTION SUMMARY
Ready
Connections: 0
regression
Search
Favorites
iCloud Drive
Applications
Google Drive
Desktop
Documents
Downloads
MEGA
Devices
Remote Disc
Shared
bthub4
No Tags
classification
clustering
regression
No Tags
house_price_test.csv
house_price_train.csv
CSV
house_price_test.csv
7 KB
Created Today, 01:06
Modified Today, 01:06
Last opened Today, 01:06
Add Tags...
Format: All Files
Options
Cancel
Open

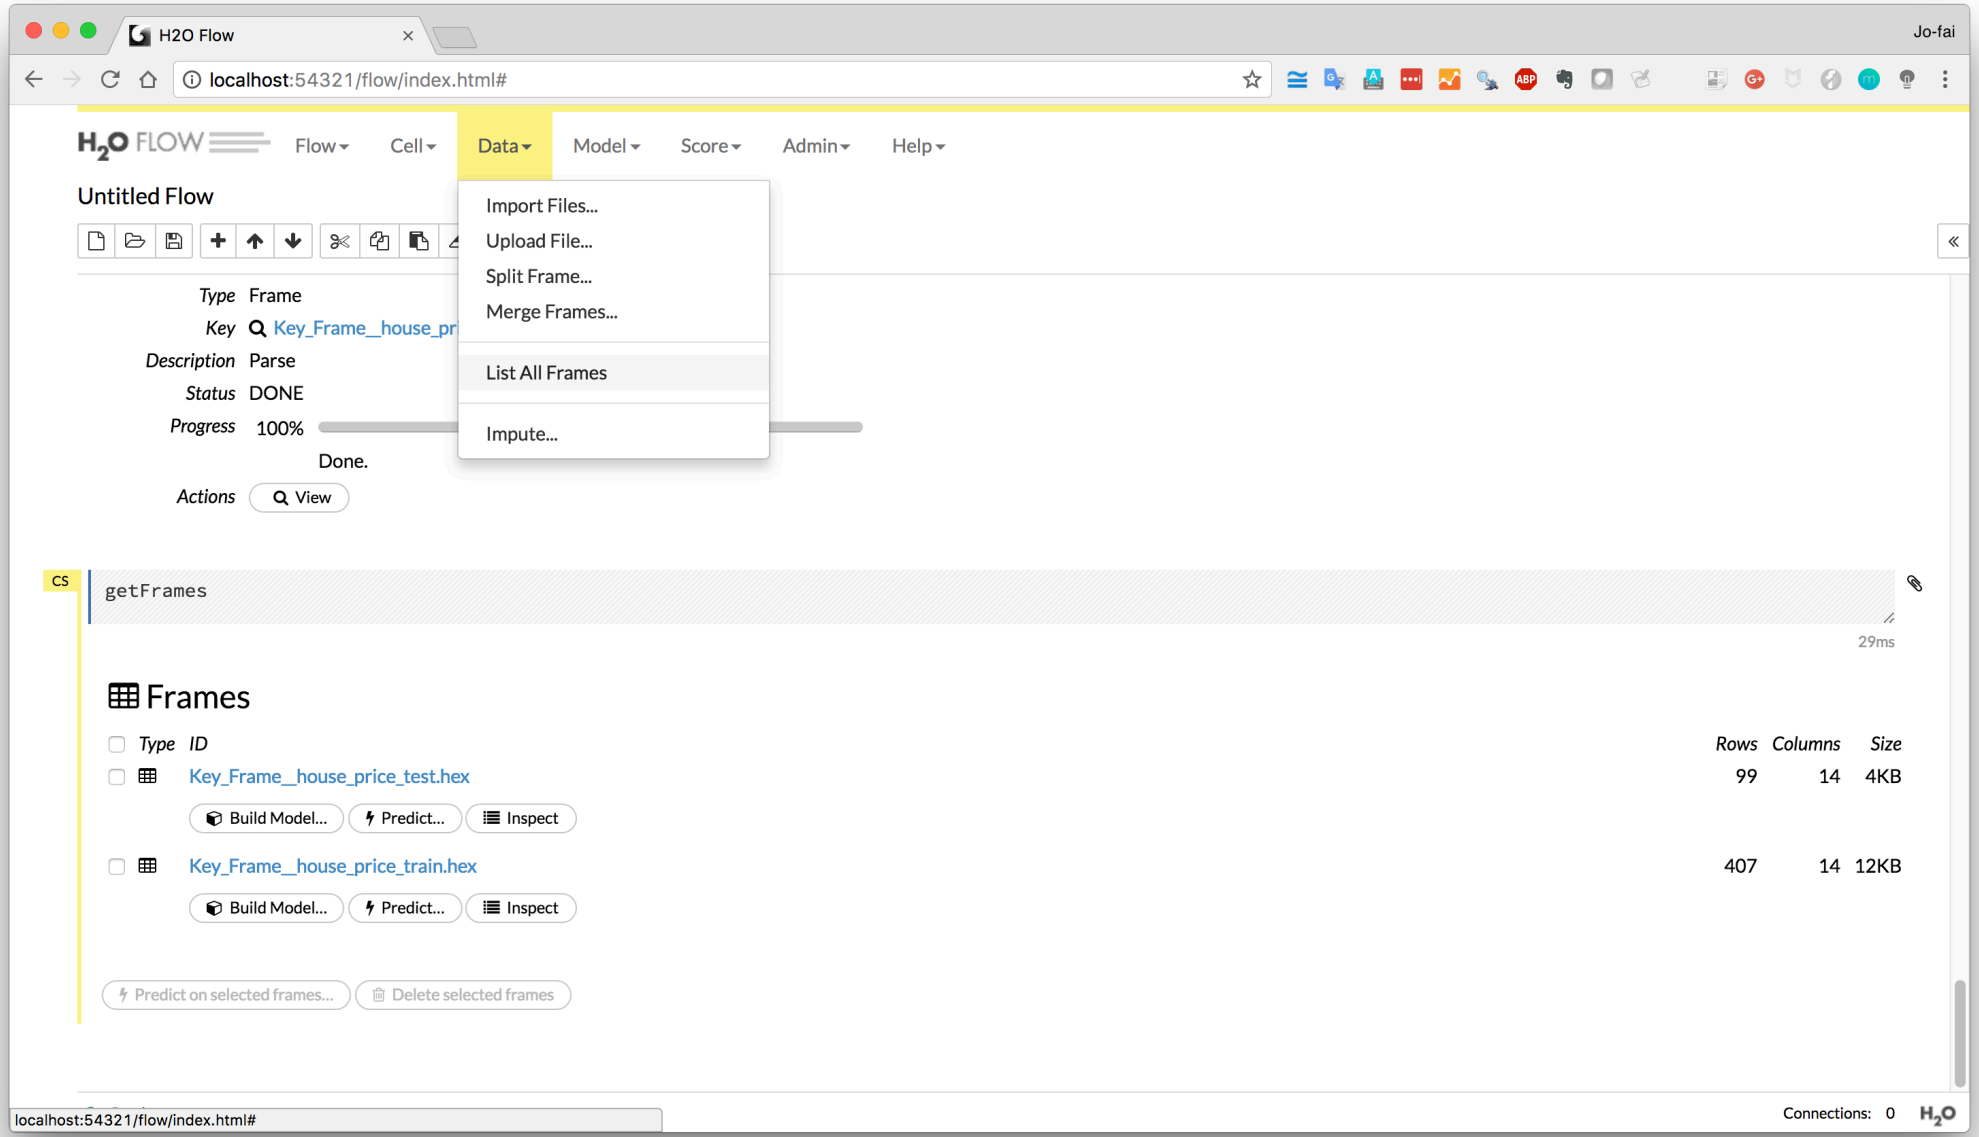
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Import Files...
Upload File...
Split Frame...
Merge Frames...
List All Frames
Impute...
Type Frame
Key Key_Frame__house_pr
Description Parse
Status DONE
Progress 100%
Done.
Actions View
CS
getFrames
29ms
Frames
Type ID
Rows Columns Size
Key_Frame__house_price_test.hex
99 14 4KB
Build Model... Predict... Inspect
Key_Frame__house_price_train.hex
407 14 12KB
Build Model... Predict... Inspect
Predict on selected frames... Delete selected frames
Connections: 0
localhost:54321/flow/index.html#

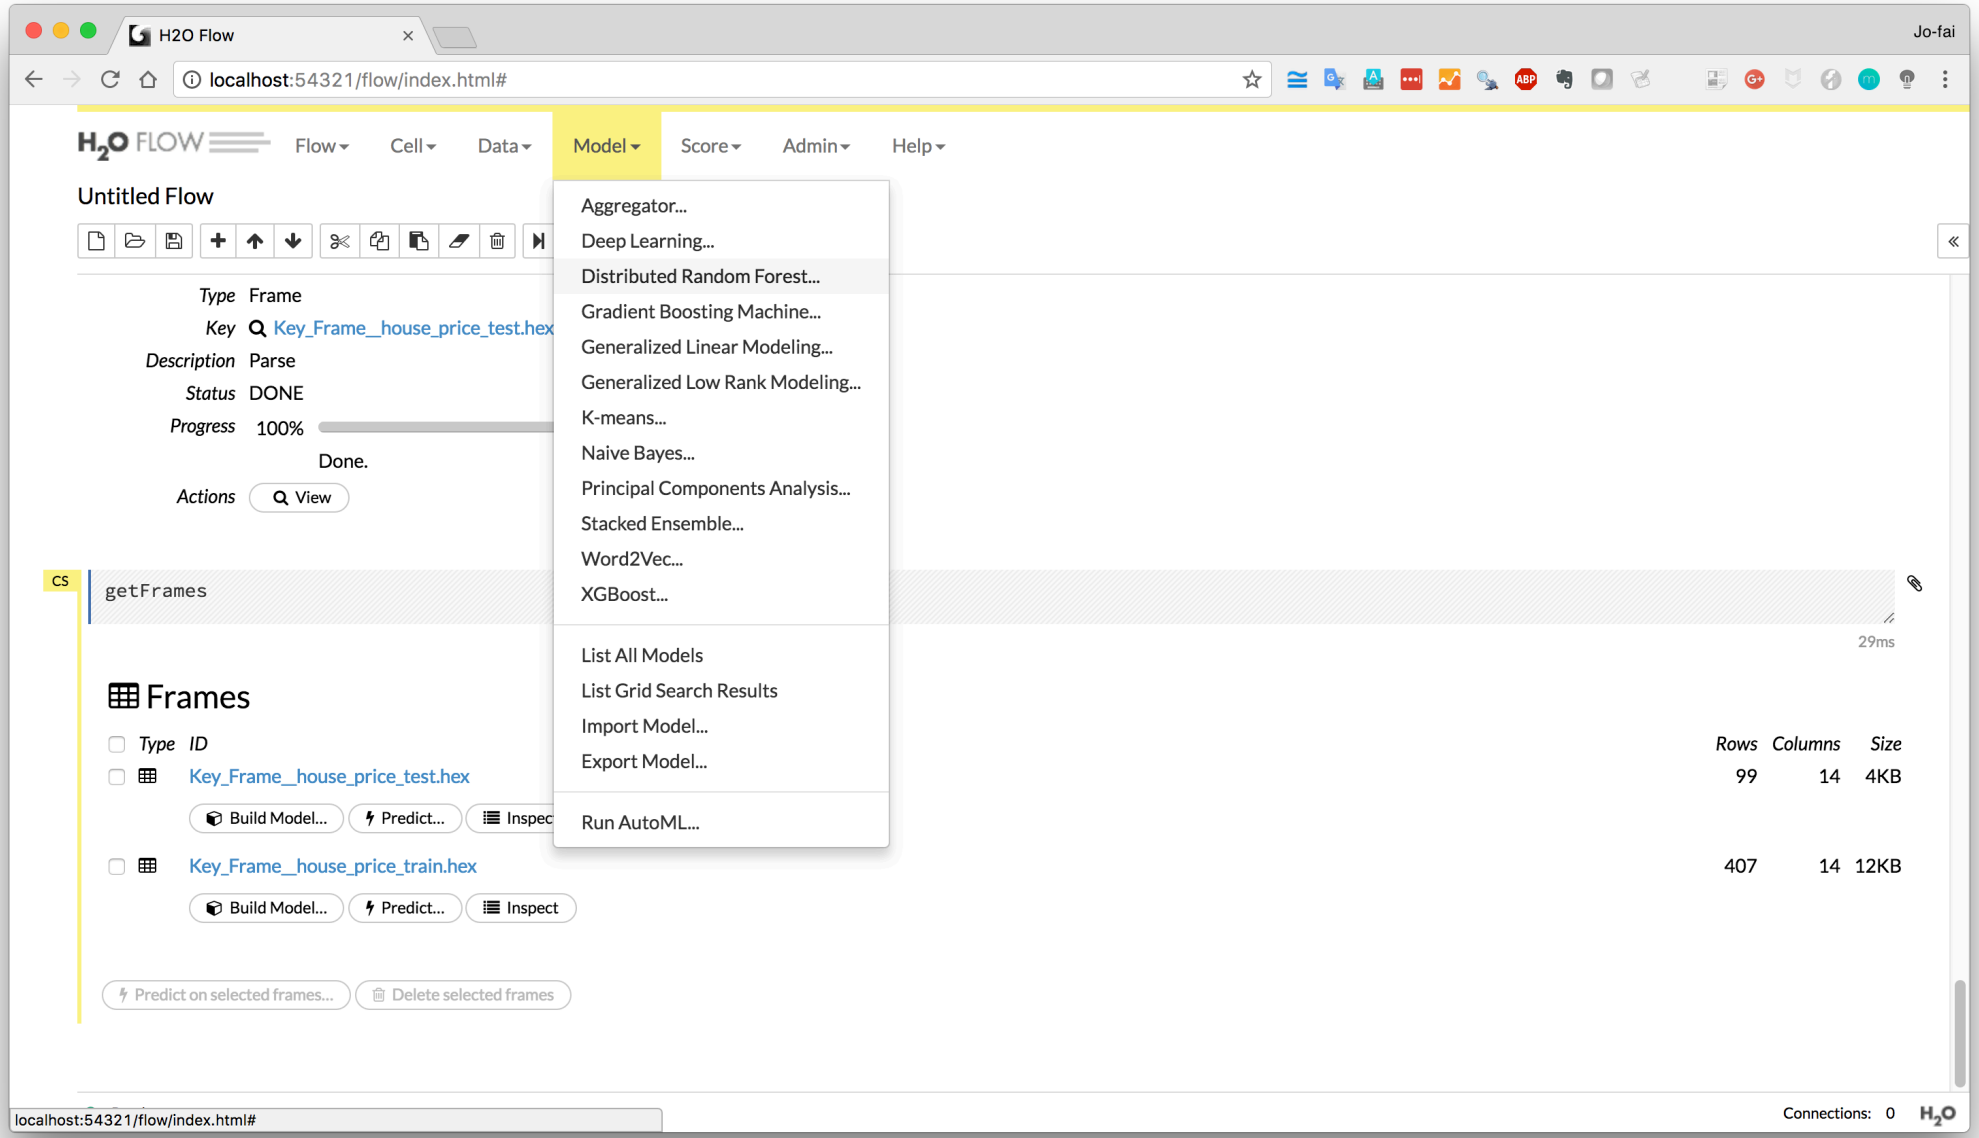
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Type Frame
Key Key_Frame__house_price_test.hex
Description Parse
Status DONE
Progress 100%
Done.
Actions View
Aggregator...
Deep Learning...
Distributed Random Forest...
Gradient Boosting Machine...
Generalized Linear Modeling...
Generalized Low Rank Modeling...
K-means...
Naive Bayes...
Principal Components Analysis...
Stacked Ensemble...
Word2Vec...
XGBoost...
List All Models
List Grid Search Results
Import Model...
Export Model...
Run AutoML...
CS
getFrames
29ms
Frames
Type ID
Rows Columns Size
Key_Frame__house_price_test.hex
99 14 4KB
Build Model...
Predict...
Inspect
Key_Frame__house_price_train.hex
407 14 12KB
Build Model...
Predict...
Inspect
Predict on selected frames...
Delete selected frames
Connections: 0
localhost:54321/flow/index.html#

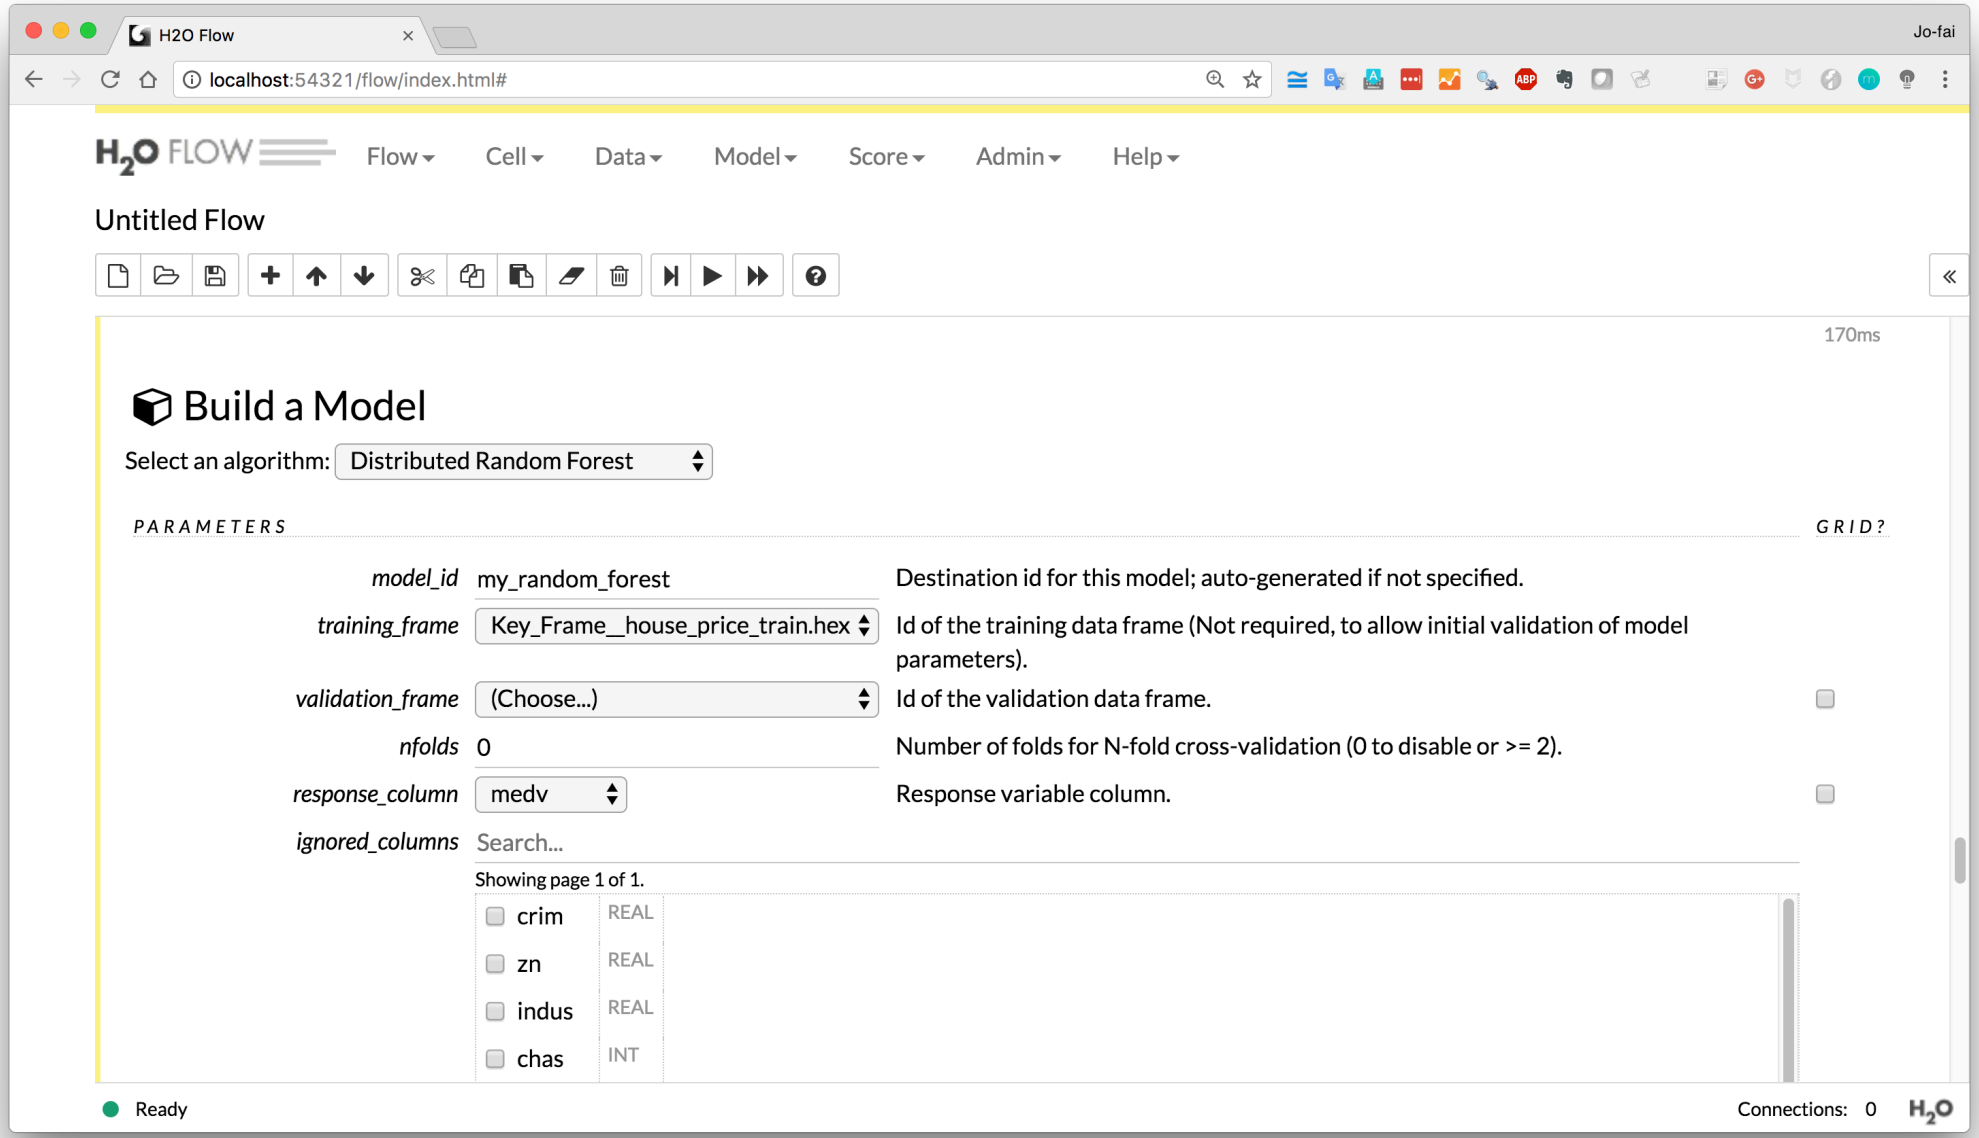
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
170ms
Build a Model
Select an algorithm: Distributed Random Forest
PARAMETERS
GRID?
model_id
my_random_forest
Destination id for this model; auto-generated if not specified.
training_frame
Key_Frame__house_price_train.hex
Id of the training data frame (Not required, to allow initial validation of model parameters).
validation_frame
(Choose...)
Id of the validation data frame.
nfolds
0
Number of folds for N-fold cross-validation (0 to disable or >= 2).
response_column
medv
Response variable column.
ignored_columns
Search...
Showing page 1 of 1.
crim REAL
zn REAL
indus REAL
chas INT
Ready
Connections: 0

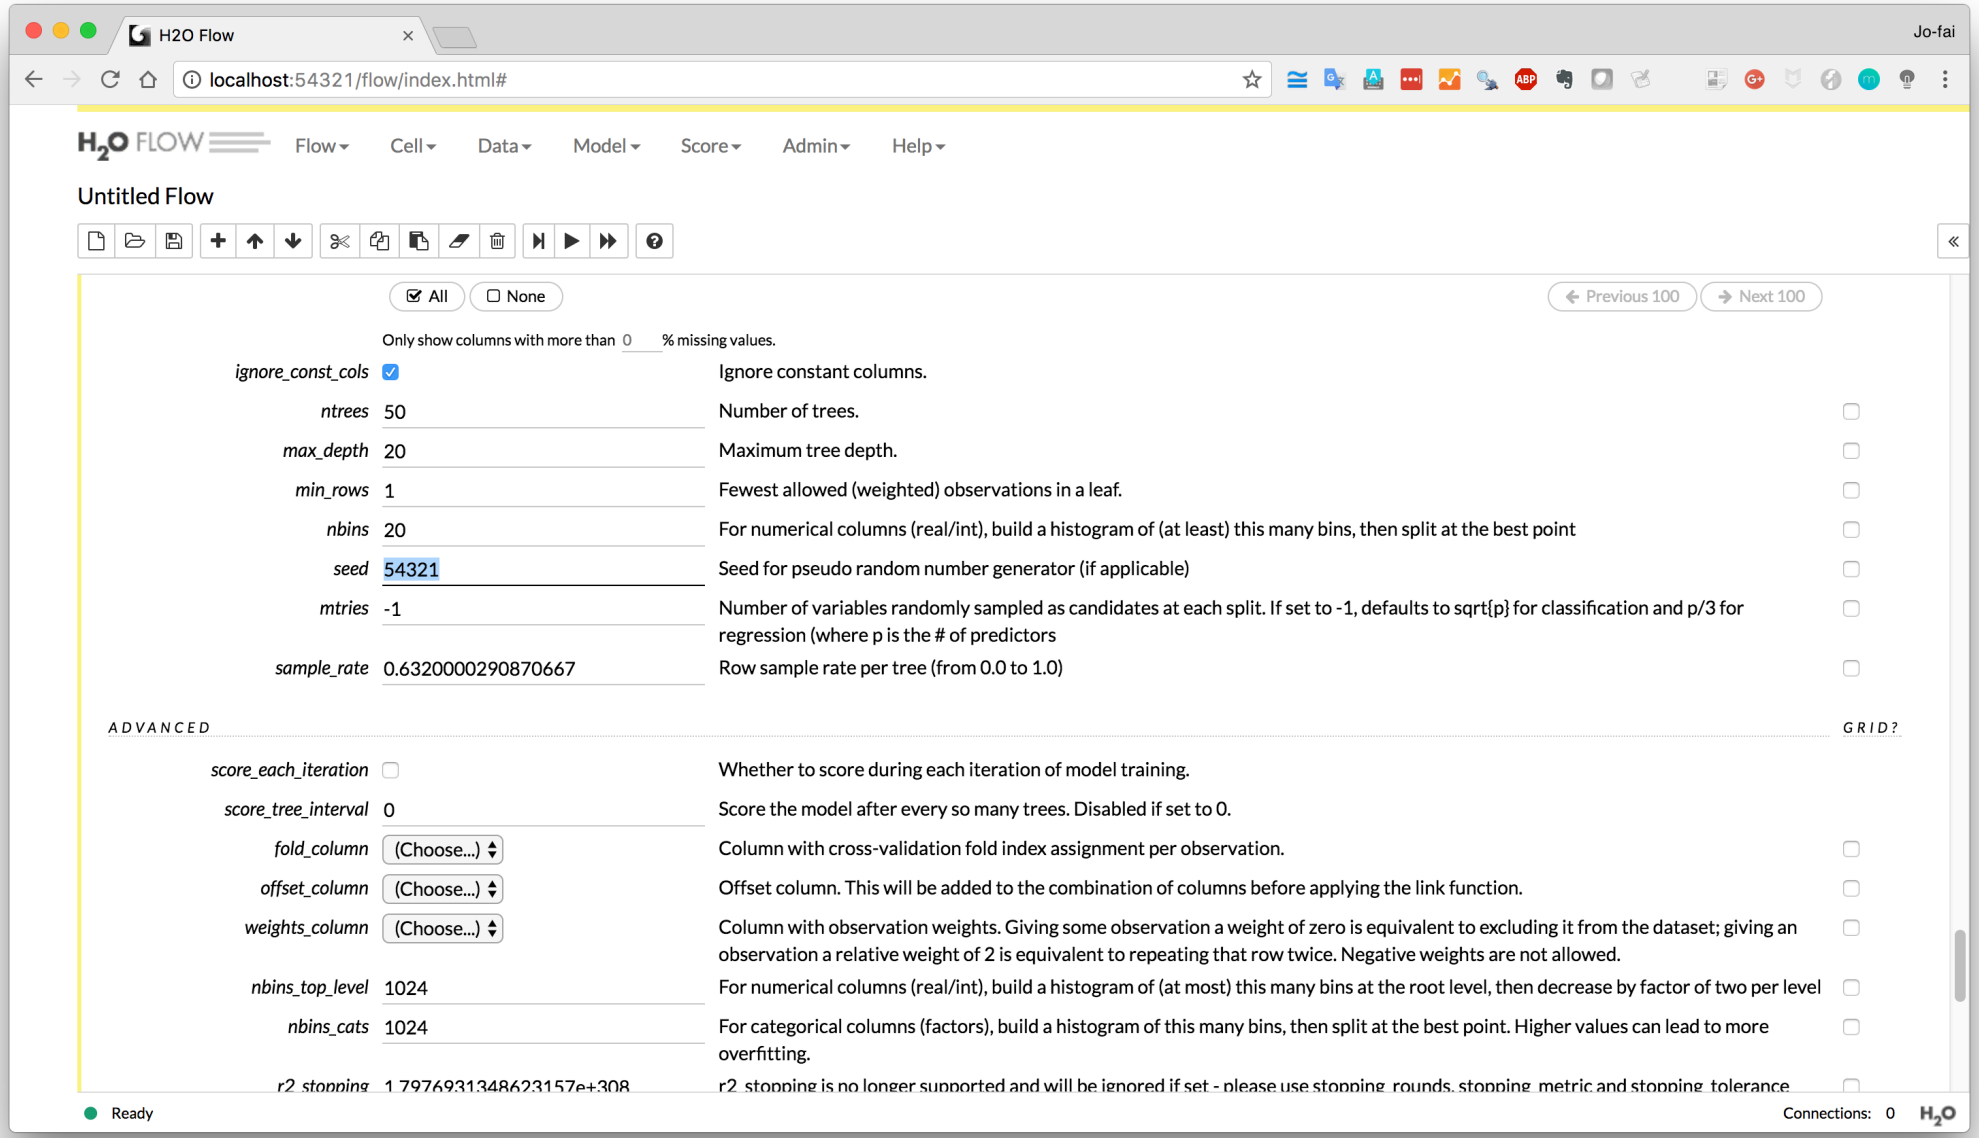
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
All
None
Previous 100
Next 100
Only show columns with more than 0 % missing values.
ignore_const_cols
Ignore constant columns.
ntrees
50
Number of trees.
max_depth
20
Maximum tree depth.
min_rows
1
Fewest allowed (weighted) observations in a leaf.
nbins
20
For numerical columns (real/int), build a histogram of (at least) this many bins, then split at the best point
seed
54321
Seed for pseudo random number generator (if applicable)
mtries
-1
Number of variables randomly sampled as candidates at each split. If set to -1, defaults to sqrt{p} for classification and p/3 for regression (where p is the # of predictors
sample_rate
0.6320000290870667
Row sample rate per tree (from 0.0 to 1.0)
ADVANCED
GRID?
score_each_iteration
Whether to score during each iteration of model training.
score_tree_interval
0
Score the model after every so many trees. Disabled if set to 0.
fold_column
(Choose...)
Column with cross-validation fold index assignment per observation.
offset_column
(Choose...)
Offset column. This will be added to the combination of columns before applying the link function.
weights_column
(Choose...)
Column with observation weights. Giving some observation a weight of zero is equivalent to excluding it from the dataset; giving an observation a relative weight of 2 is equivalent to repeating that row twice. Negative weights are not allowed.
nbins_top_level
1024
For numerical columns (real/int), build a histogram of (at most) this many bins at the root level, then decrease by factor of two per level
nbins_cats
1024
For categorical columns (factors), build a histogram of this many bins, then split at the best point. Higher values can lead to more overfitting.
r2_stopping
1.7976931348623157e+308
r2_stopping is no longer supported and will be ignored if set - please use stopping_rounds, stopping_metric and stopping_tolerance
Ready
Connections: 0
H2O

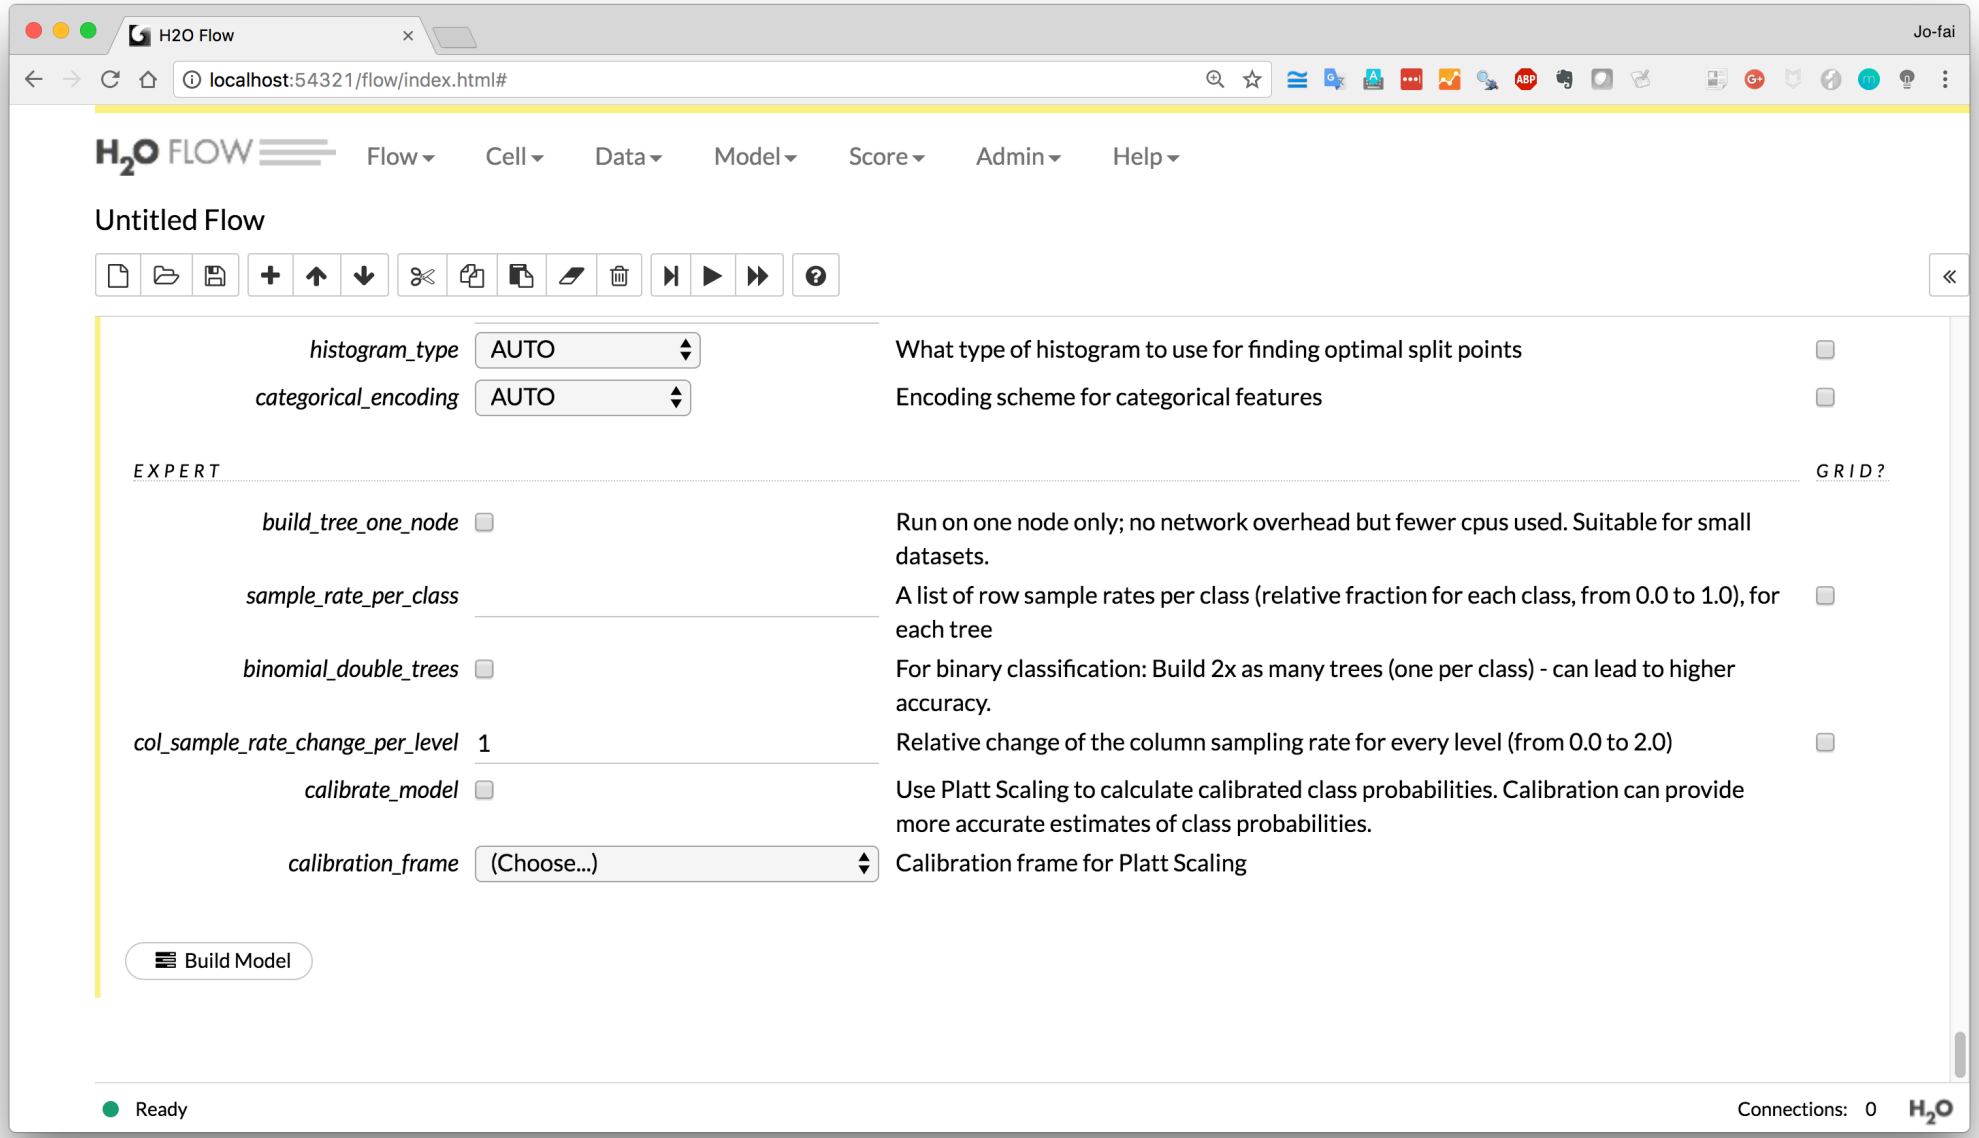
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
histogram_type
AUTO
What type of histogram to use for finding optimal split points
categorical_encoding
AUTO
Encoding scheme for categorical features
EXPERT
GRID?
build_tree_one_node
Run on one node only; no network overhead but fewer cpus used. Suitable for small datasets.
sample_rate_per_class
A list of row sample rates per class (relative fraction for each class, from 0.0 to 1.0), for each tree
binomial_double_trees
For binary classification: Build 2x as many trees (one per class) - can lead to higher accuracy.
col_sample_rate_change_per_level
1
Relative change of the column sampling rate for every level (from 0.0 to 2.0)
calibrate_model
Use Platt Scaling to calculate calibrated class probabilities. Calibration can provide more accurate estimates of class probabilities.
calibration_frame
(Choose...)
Calibration frame for Platt Scaling
Build Model
Ready
Connections: 0
H2O

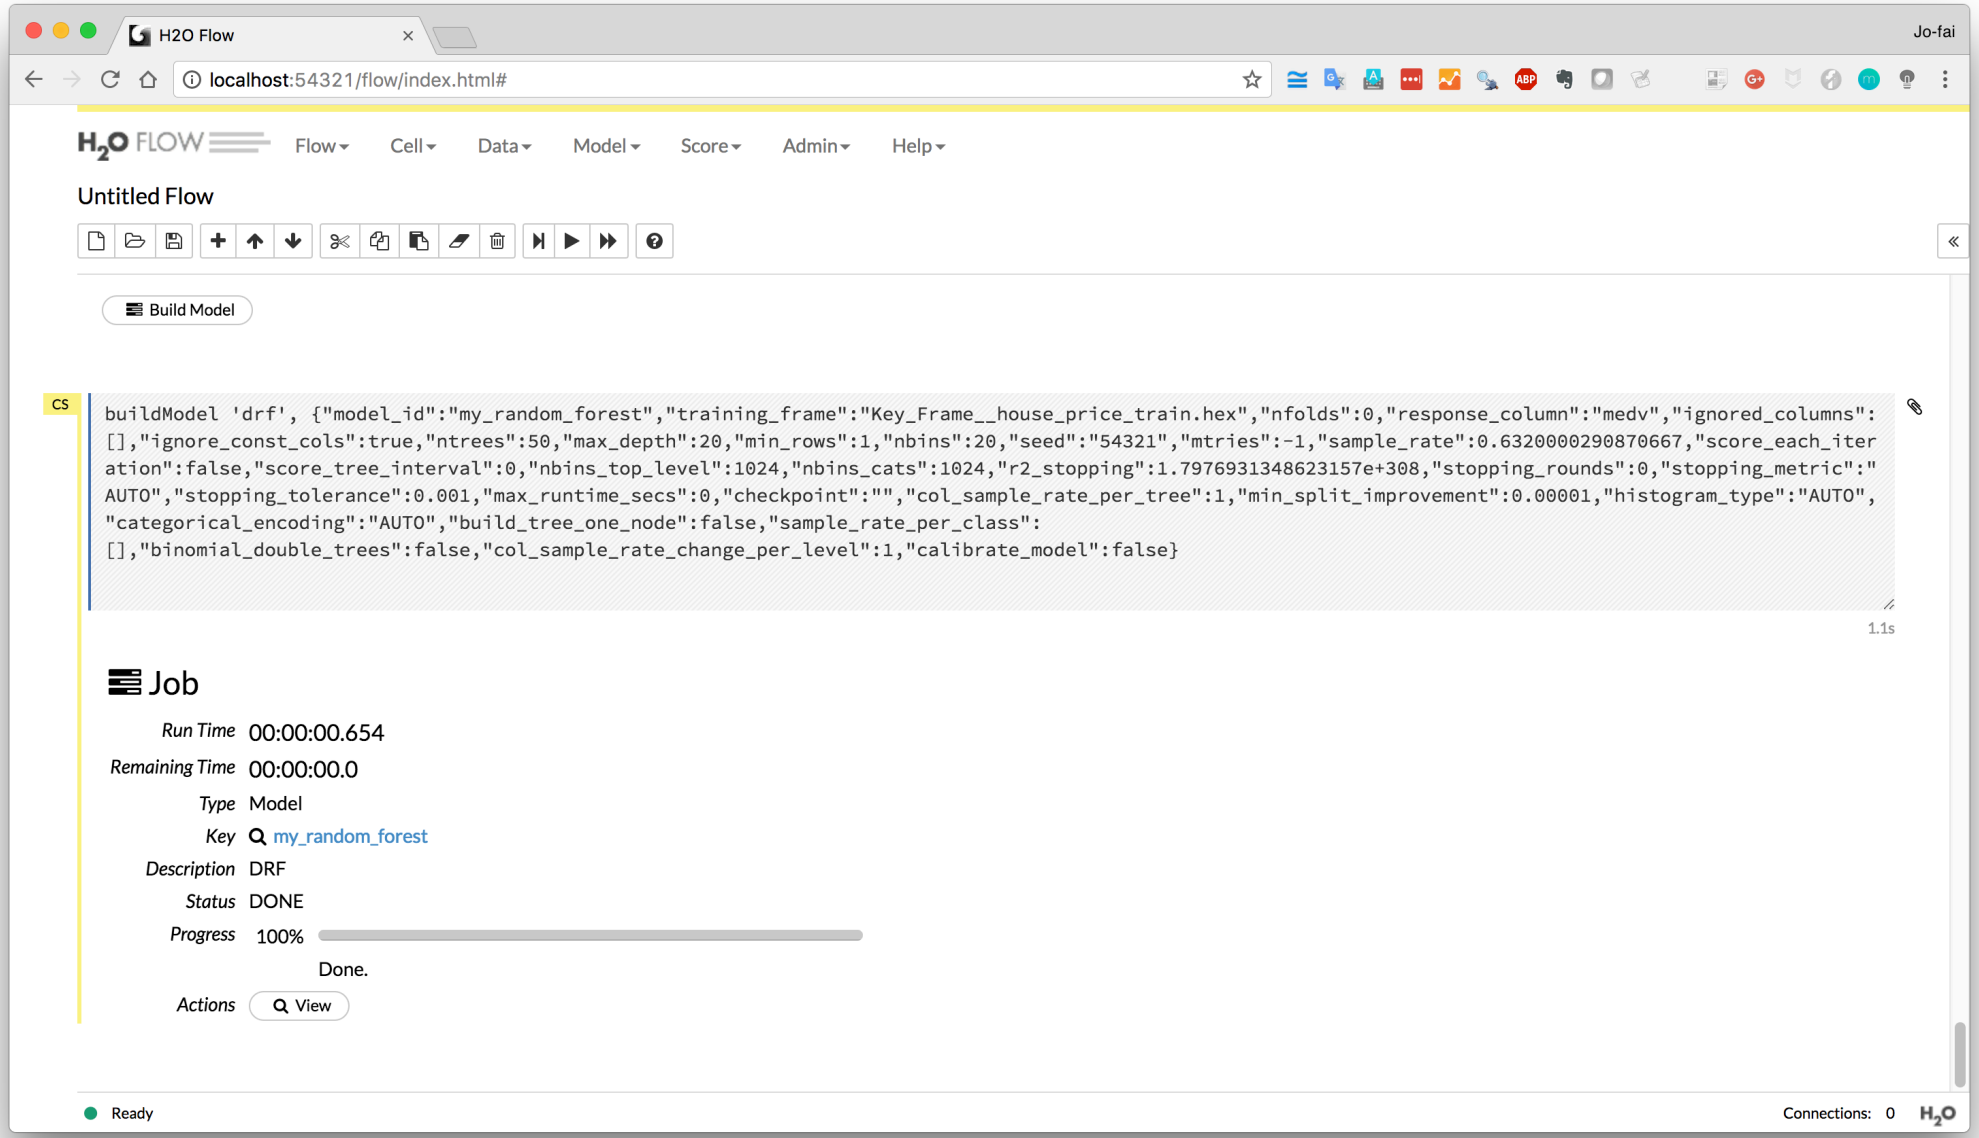
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Build Model
CS
buildModel 'drf', {"model_id":"my_random_forest","training_frame":"Key_Frame__house_price_train.hex","nfolds":0,"response_column":"medv","ignored_columns":[],"ignore_const_cols":true,"ntrees":50,"max_depth":20,"min_rows":1,"nbins":20,"seed":"54321","mtries":-1,"sample_rate":0.6320000290870667,"score_each_iteration":false,"score_tree_interval":0,"nbins_top_level":1024,"nbins_cats":1024,"r2_stopping":1.7976931348623157e+308,"stopping_rounds":0,"stopping_metric":"AUTO","stopping_tolerance":0.001,"max_runtime_secs":0,"checkpoint":"","col_sample_rate_per_tree":1,"min_split_improvement":0.00001,"histogram_type":"AUTO","categorical_encoding":"AUTO","build_tree_one_node":false,"sample_rate_per_class":[],"binomial_double_trees":false,"col_sample_rate_change_per_level":1,"calibrate_model":false}
1.1s
Job
Run Time 00:00:00.654
Remaining Time 00:00:00.0
Type Model
Key my_random_forest
Description DRF
Status DONE
Progress 100%
Done.
Actions View
Ready
Connections: 0
H2O

H2O Flow

Jo-fai

localhost:54321/flow/index.html#

H2O FLOW Flow Cell Data Model Score Admin Help

Untitled Flow

Model

Model ID: my_random_forest

Algorithm: Distributed Random Forest

Actions: Refresh | Predict... | Download POJO | Download Model Deployment Package (MOJO) | Export | Inspect | Delete | Download Gen Model

▸ MODEL PARAMETERS

▾ SCORING HISTORY - DEVIANCE

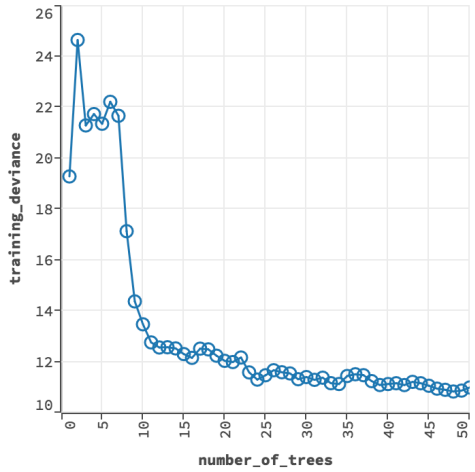

▾ VARIABLE IMPORTANCES

Ready

Connections: 0

H2O Flow

Jo-fai

localhost:54321/flow/index.html#

H2O FLOW Flow Cell Data Model Score Admin Help

Untitled Flow

VARIABLE IMPORTANCES

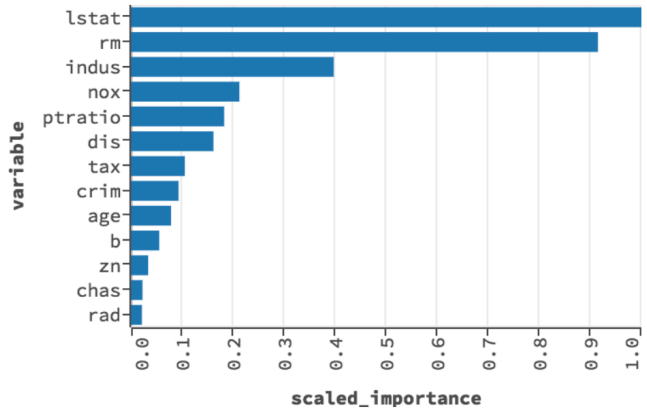

OUTPUT

OUTPUT - MODEL SUMMARY

OUTPUT - SCORING HISTORY

OUTPUT - TRAINING_METRICS

Ready

Connections: 0

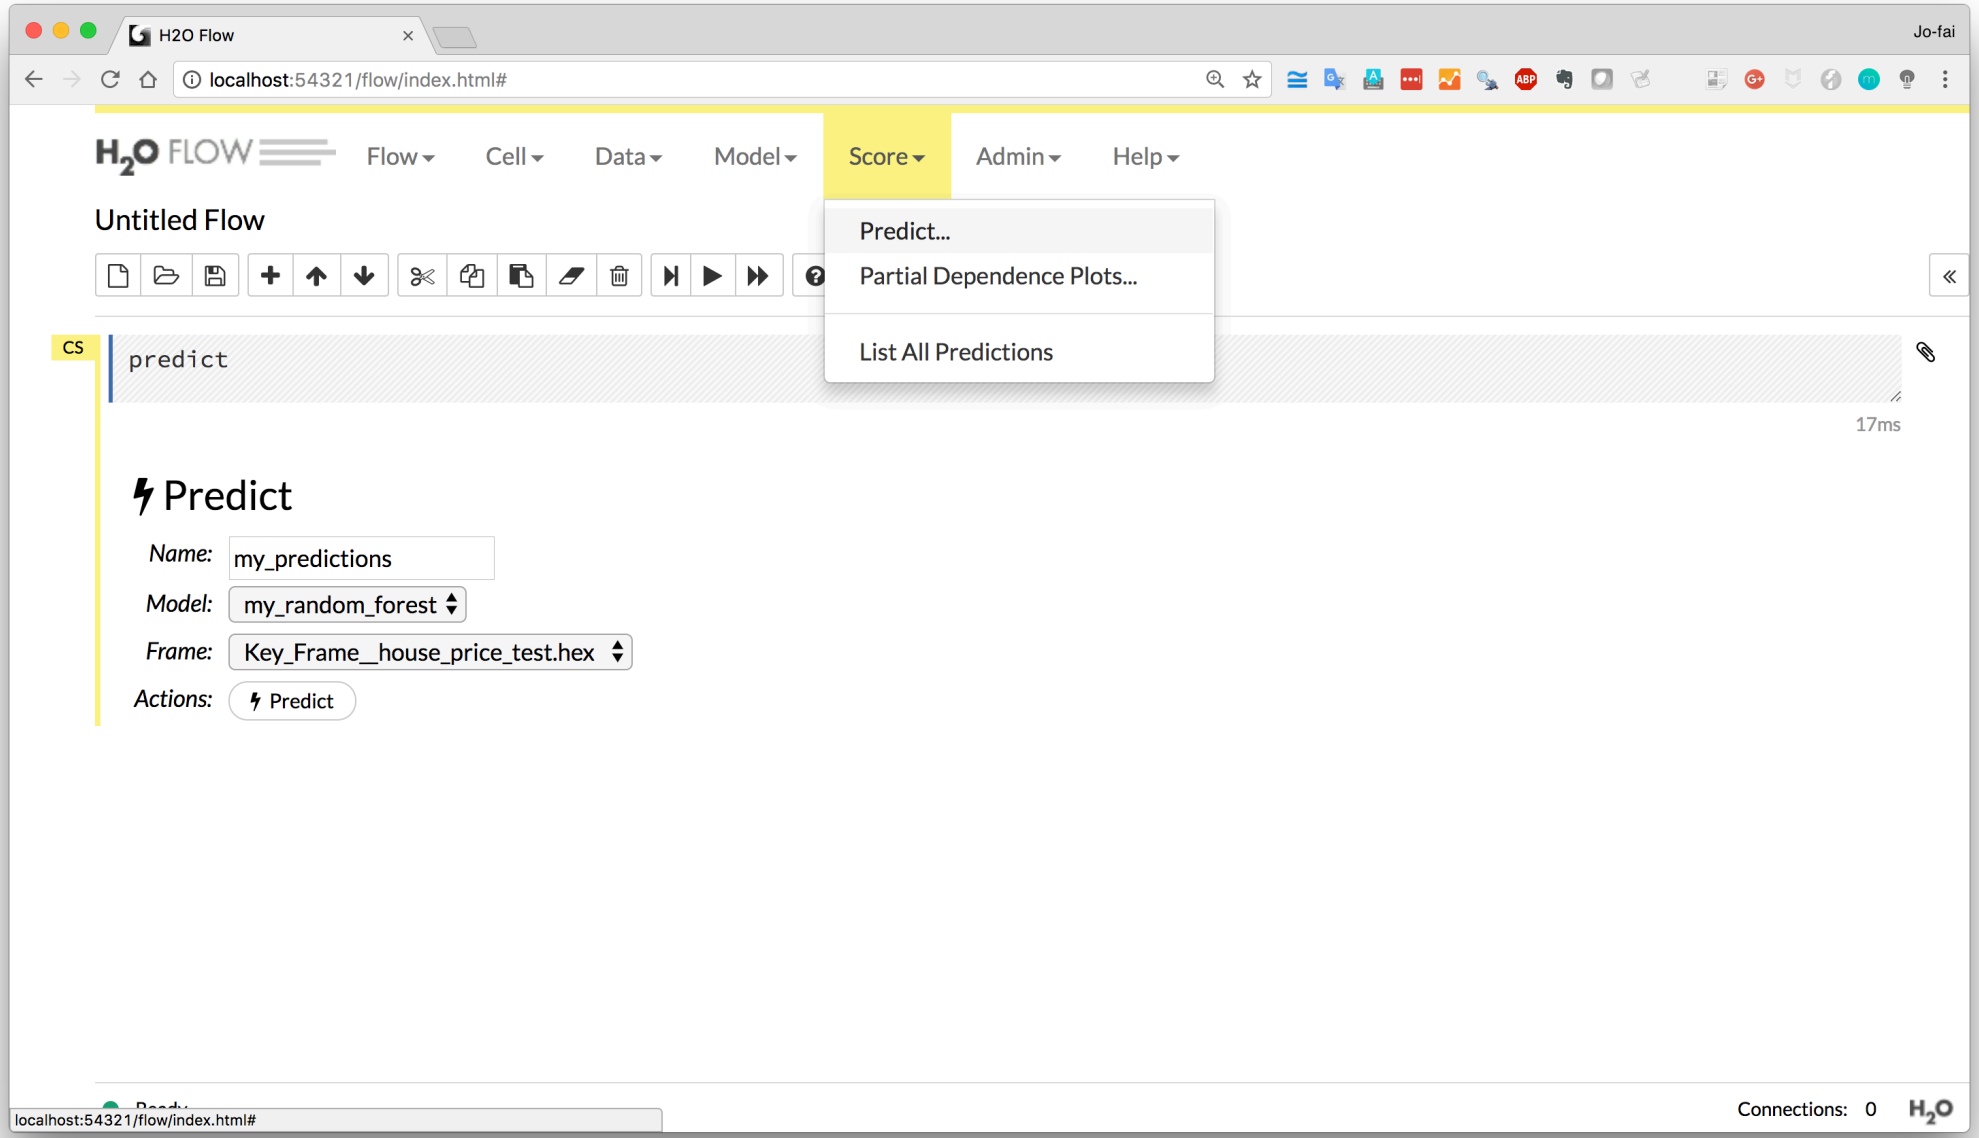
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
Predict...
Partial Dependence Plots...
List All Predictions
CS
predict
17ms
Predict
Name:
my_predictions
Model:
my_random_forest
Frame:
Key_Frame__house_price_test.hex
Actions:
Predict
Connections: 0
localhost:54321/flow/index.html#

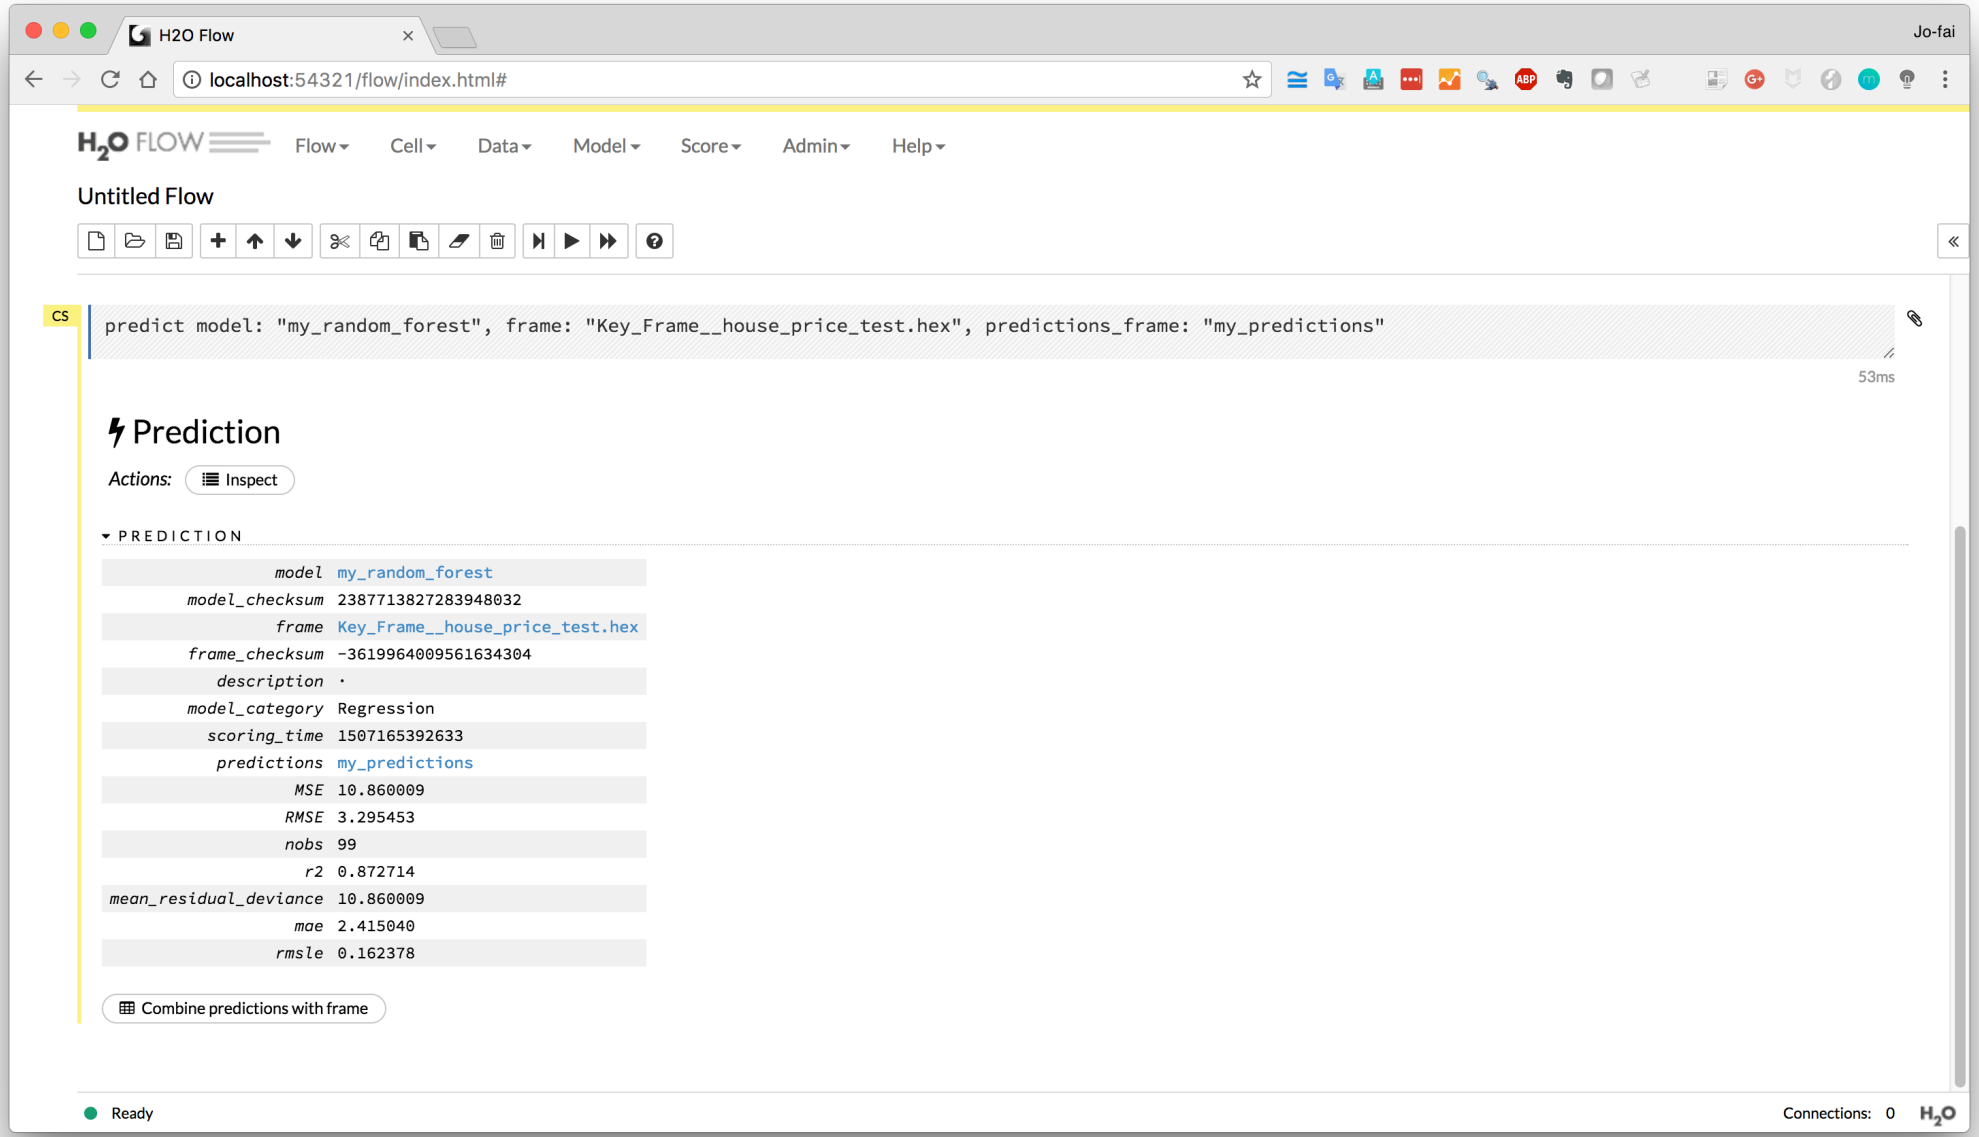
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
CS
predict model: "my_random_forest", frame: "Key_Frame__house_price_test.hex", predictions_frame: "my_predictions"
53ms
Prediction
Actions:
Inspect
PREDICTION
model my_random_forest
model_checksum 2387713827283948032
frame Key_Frame__house_price_test.hex
frame_checksum -3619964009561634304
description ·
model_category Regression
scoring_time 1507165392633
predictions my_predictions
MSE 10.860009
RMSE 3.295453
nobs 99
r2 0.872714
mean_residual_deviance 10.860009
mae 2.415040
rmsle 0.162378
Combine predictions with frame
Ready
Connections: 0

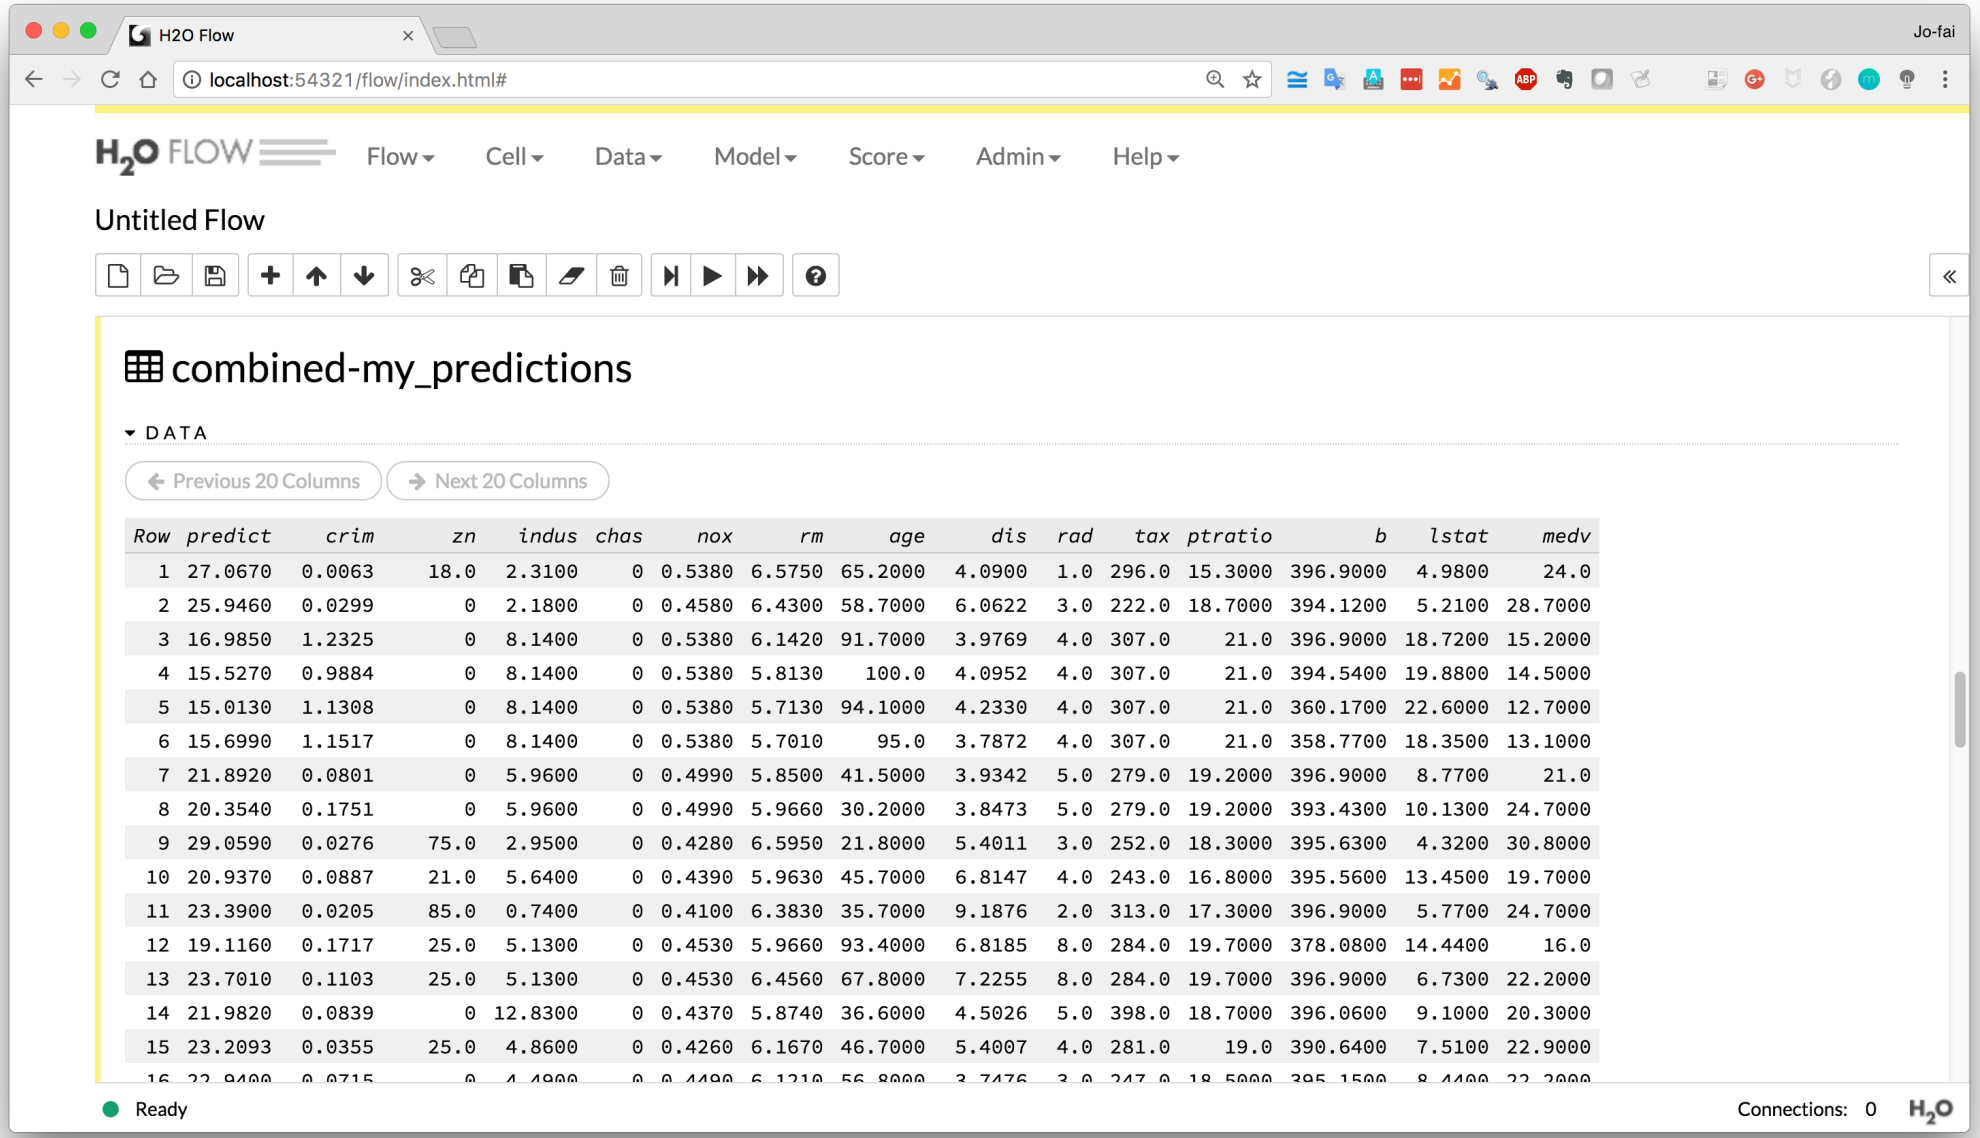

Row	predict	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
1	27.0670	0.0063	18.0	2.3100	0	0.5380	6.5750	65.2000	4.0900	1.0	296.0	15.3000	396.9000	4.9800	24.0
2	25.9460	0.0299	0	2.1800	0	0.4580	6.4300	58.7000	6.0622	3.0	222.0	18.7000	394.1200	5.2100	28.7000
3	16.9850	1.2325	0	8.1400	0	0.5380	6.1420	91.7000	3.9769	4.0	307.0	21.0	396.9000	18.7200	15.2000
4	15.5270	0.9884	0	8.1400	0	0.5380	5.8130	100.0	4.0952	4.0	307.0	21.0	394.5400	19.8800	14.5000
5	15.0130	1.1308	0	8.1400	0	0.5380	5.7130	94.1000	4.2330	4.0	307.0	21.0	360.1700	22.6000	12.7000
6	15.6990	1.1517	0	8.1400	0	0.5380	5.7010	95.0	3.7872	4.0	307.0	21.0	358.7700	18.3500	13.1000
7	21.8920	0.0801	0	5.9600	0	0.4990	5.8500	41.5000	3.9342	5.0	279.0	19.2000	396.9000	8.7700	21.0
8	20.3540	0.1751	0	5.9600	0	0.4990	5.9660	30.2000	3.8473	5.0	279.0	19.2000	393.4300	10.1300	24.7000
9	29.0590	0.0276	75.0	2.9500	0	0.4280	6.5950	21.8000	5.4011	3.0	252.0	18.3000	395.6300	4.3200	30.8000
10	20.9370	0.0887	21.0	5.6400	0	0.4390	5.9630	45.7000	6.8147	4.0	243.0	16.8000	395.5600	13.4500	19.7000
11	23.3900	0.0205	85.0	0.7400	0	0.4100	6.3830	35.7000	9.1876	2.0	313.0	17.3000	396.9000	5.7700	24.7000
12	19.1160	0.1717	25.0	5.1300	0	0.4530	5.9660	93.4000	6.8185	8.0	284.0	19.7000	378.0800	14.4400	16.0
13	23.7010	0.1103	25.0	5.1300	0	0.4530	6.4560	67.8000	7.2255	8.0	284.0	19.7000	396.9000	6.7300	22.2000
14	21.9820	0.0839	0	12.8300	0	0.4370	5.8740	36.6000	4.5026	5.0	398.0	18.7000	396.0600	9.1000	20.3000
15	23.2093	0.0355	25.0	4.8600	0	0.4260	6.1670	46.7000	5.4007	4.0	281.0	19.0	390.6400	7.5100	22.9000
16	22.9400	0.0715	0	4.4900	0	0.4490	6.1210	56.8000	3.7476	3.0	247.0	18.5000	395.1500	8.4400	22.2000

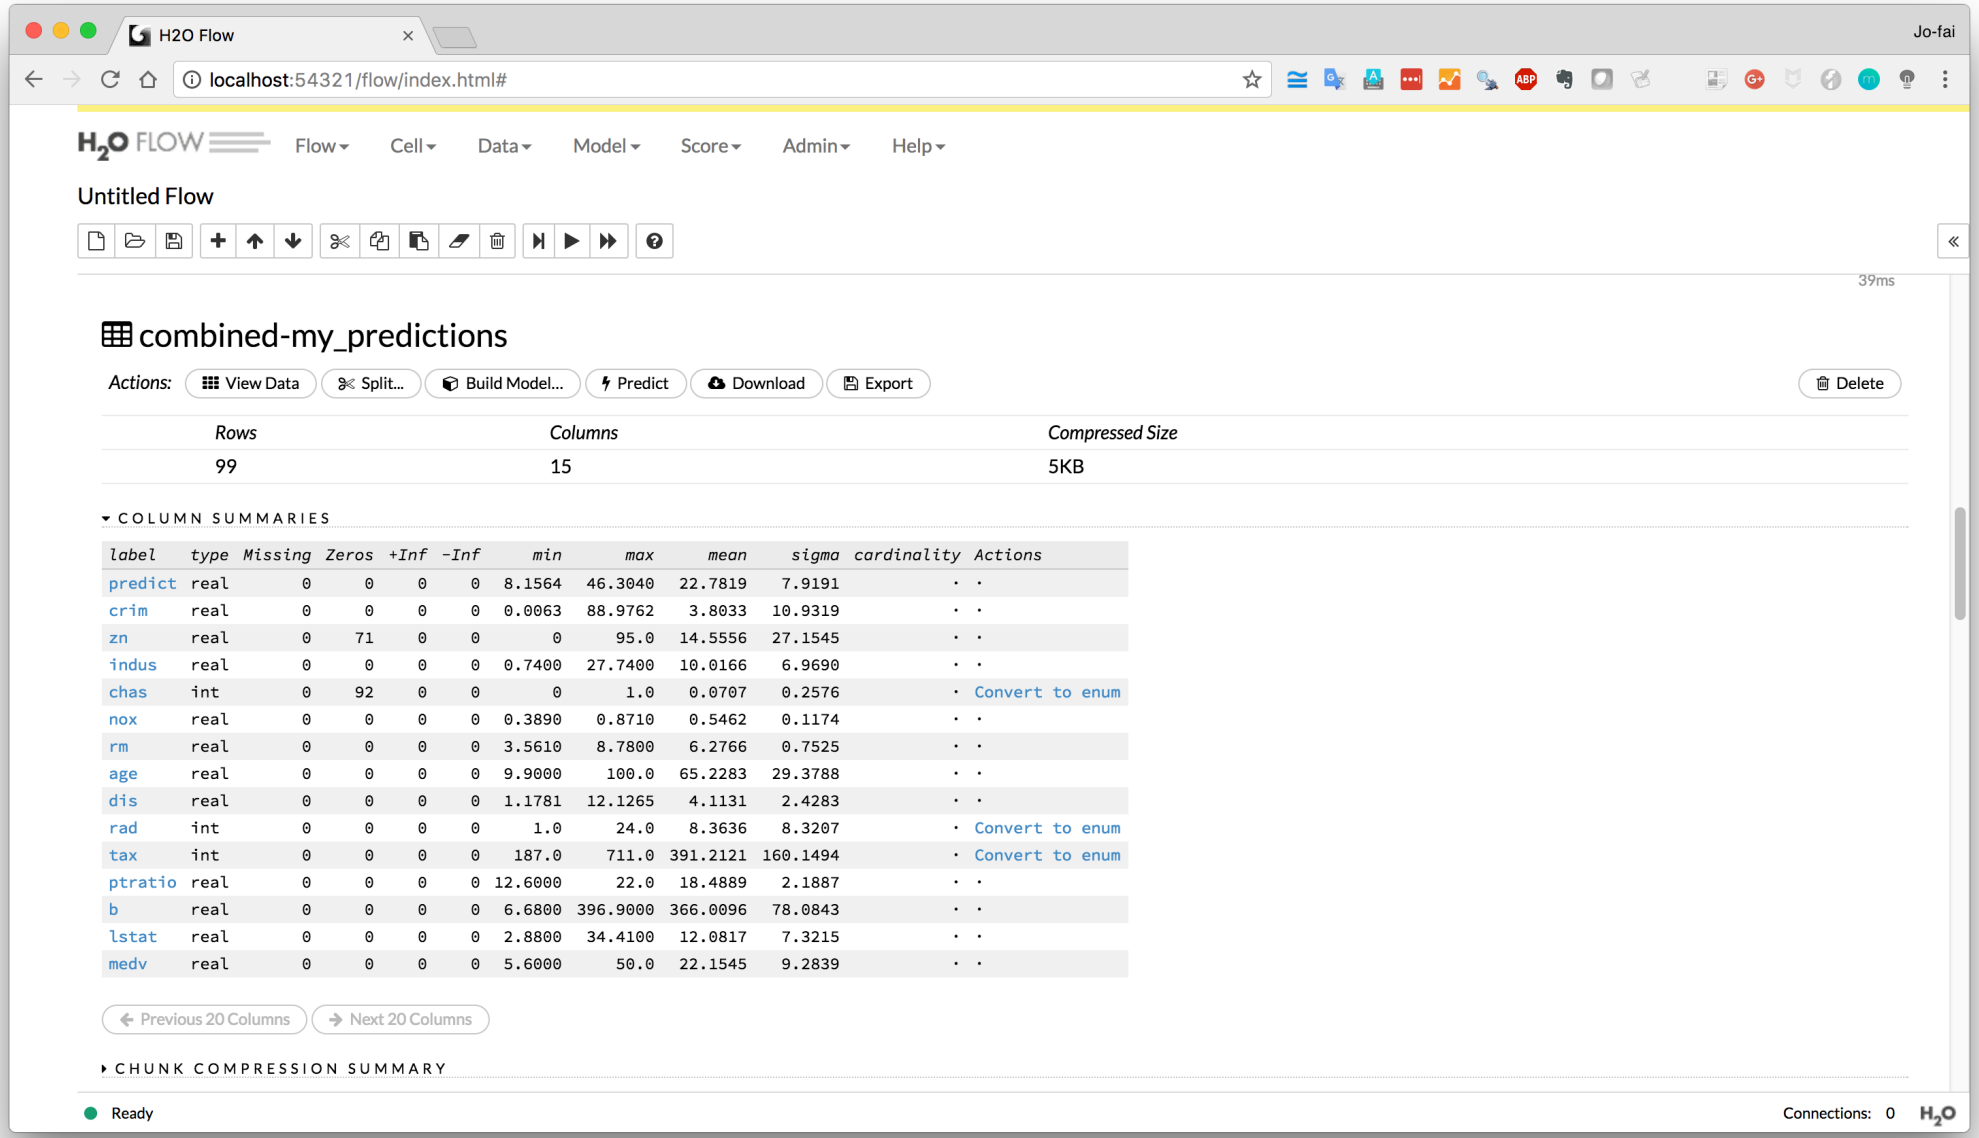
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
39ms
combined-my_predictions
Actions:
View Data
Split...
Build Model...
Predict
Download
Export
Delete
Rows
Columns
Compressed Size
99
15
5KB
COLUMN SUMMARIES
label type Missing Zeros +Inf -Inf min max mean sigma cardinality Actions
predict real 0 0 0 0 8.1564 46.3040 22.7819 7.9191 - -
crim real 0 0 0 0 0.0063 88.9762 3.8033 10.9319 - -
zn real 0 71 0 0 0 95.0 14.5556 27.1545 - -
indus real 0 0 0 0 0.7400 27.7400 10.0166 6.9690 - -
chas int 0 92 0 0 0 1.0 0.0707 0.2576 - Convert to enum
nox real 0 0 0 0 0.3890 0.8710 0.5462 0.1174 - -
rm real 0 0 0 0 3.5610 8.7800 6.2766 0.7525 - -
age real 0 0 0 0 9.9000 100.0 65.2283 29.3788 - -
dis real 0 0 0 0 1.1781 12.1265 4.1131 2.4283 - -
rad int 0 0 0 0 1.0 24.0 8.3636 8.3207 - Convert to enum
tax int 0 0 0 0 187.0 711.0 391.2121 160.1494 - Convert to enum
ptratio real 0 0 0 0 12.6000 22.0 18.4889 2.1887 - -
b real 0 0 0 0 6.6800 396.9000 366.0096 78.0843 - -
lstat real 0 0 0 0 2.8800 34.4100 12.0817 7.3215 - -
medv real 0 0 0 0 5.6000 50.0 22.1545 9.2839 - -
Previous 20 Columns
Next 20 Columns
CHUNK COMPRESSION SUMMARY
Ready
Connections: 0
H2O

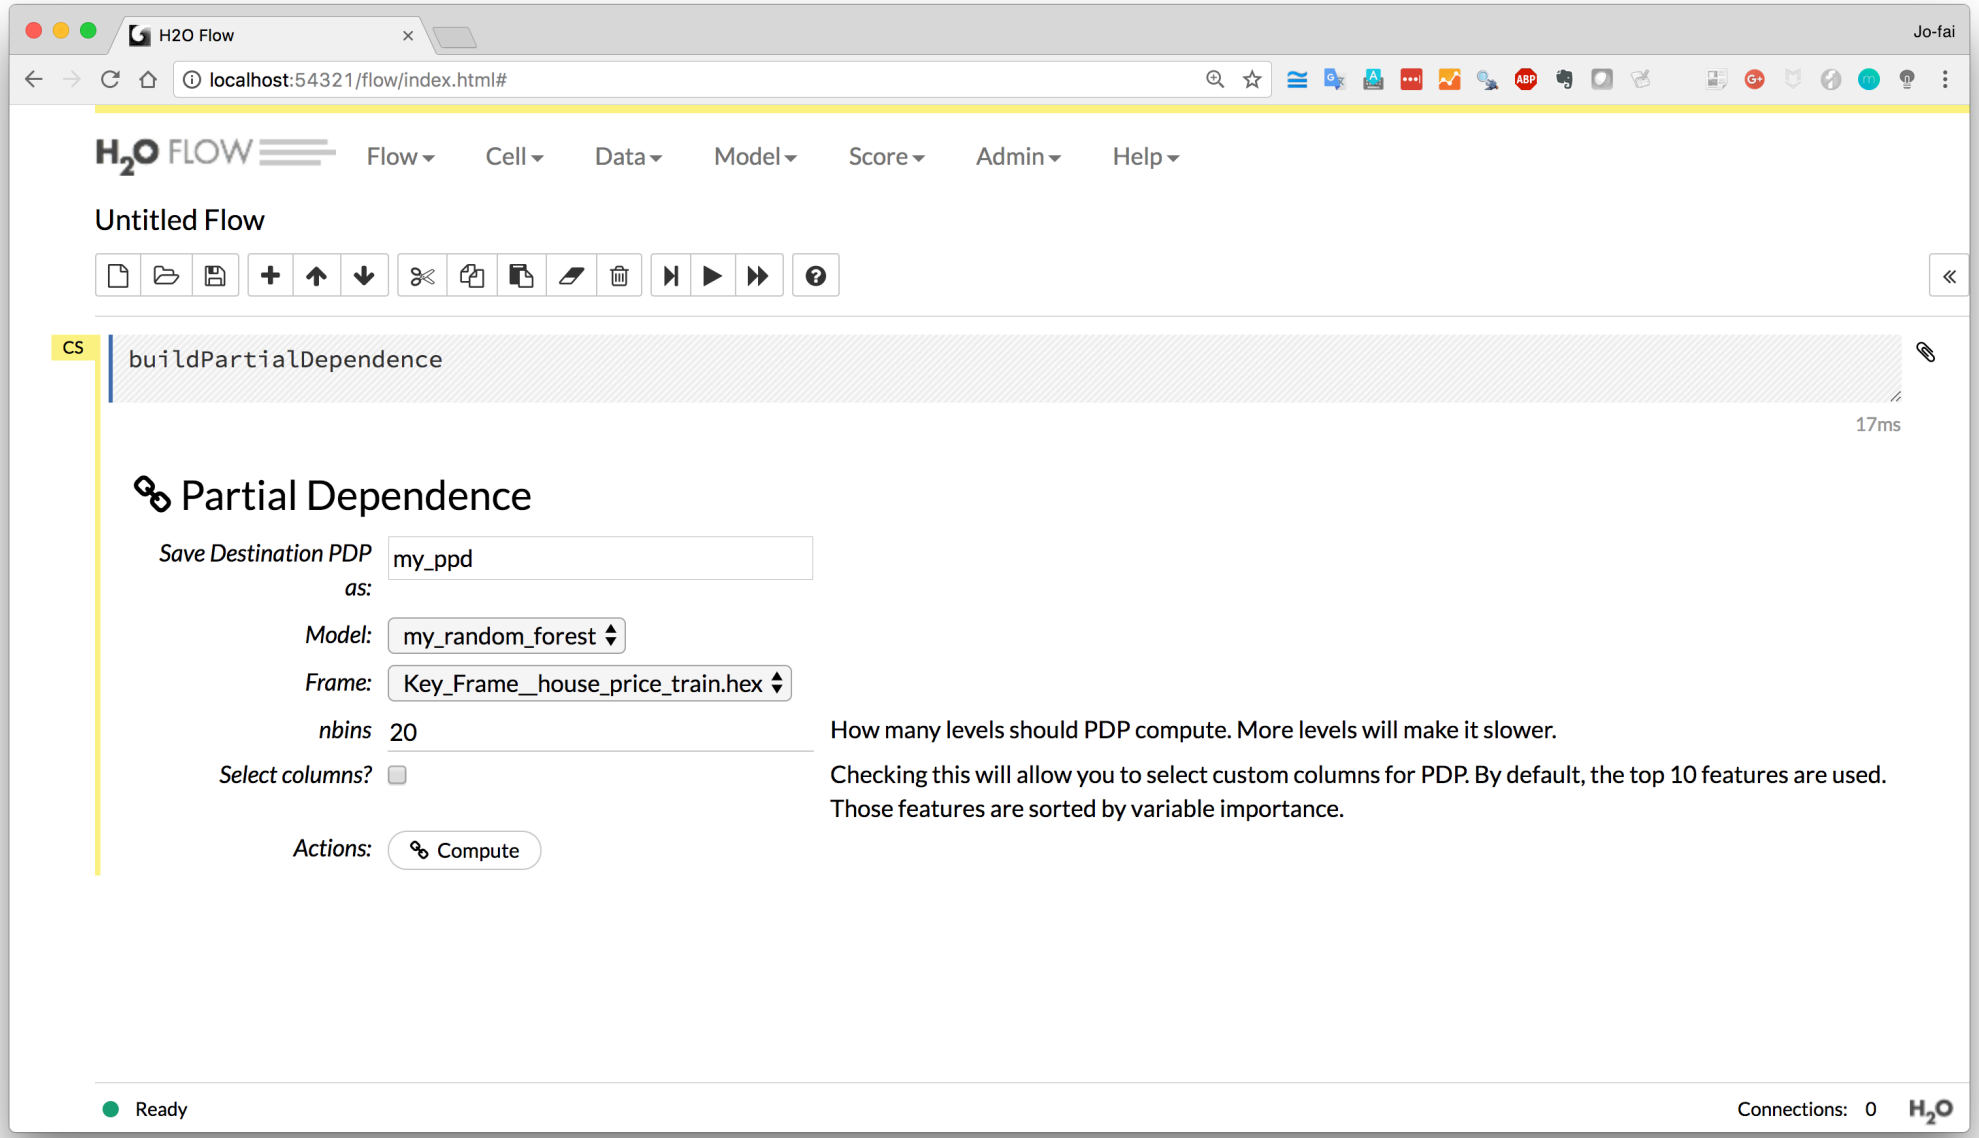
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
CS
buildPartialDependence
17ms
Partial Dependence
Save Destination PDP as:
my_ppd
Model:
my_random_forest
Frame:
Key_Frame__house_price_train.hex
nbins
20
How many levels should PDP compute. More levels will make it slower.
Select columns?
Checking this will allow you to select custom columns for PDP. By default, the top 10 features are used. Those features are sorted by variable importance.
Actions:
Compute
Ready
Connections: 0

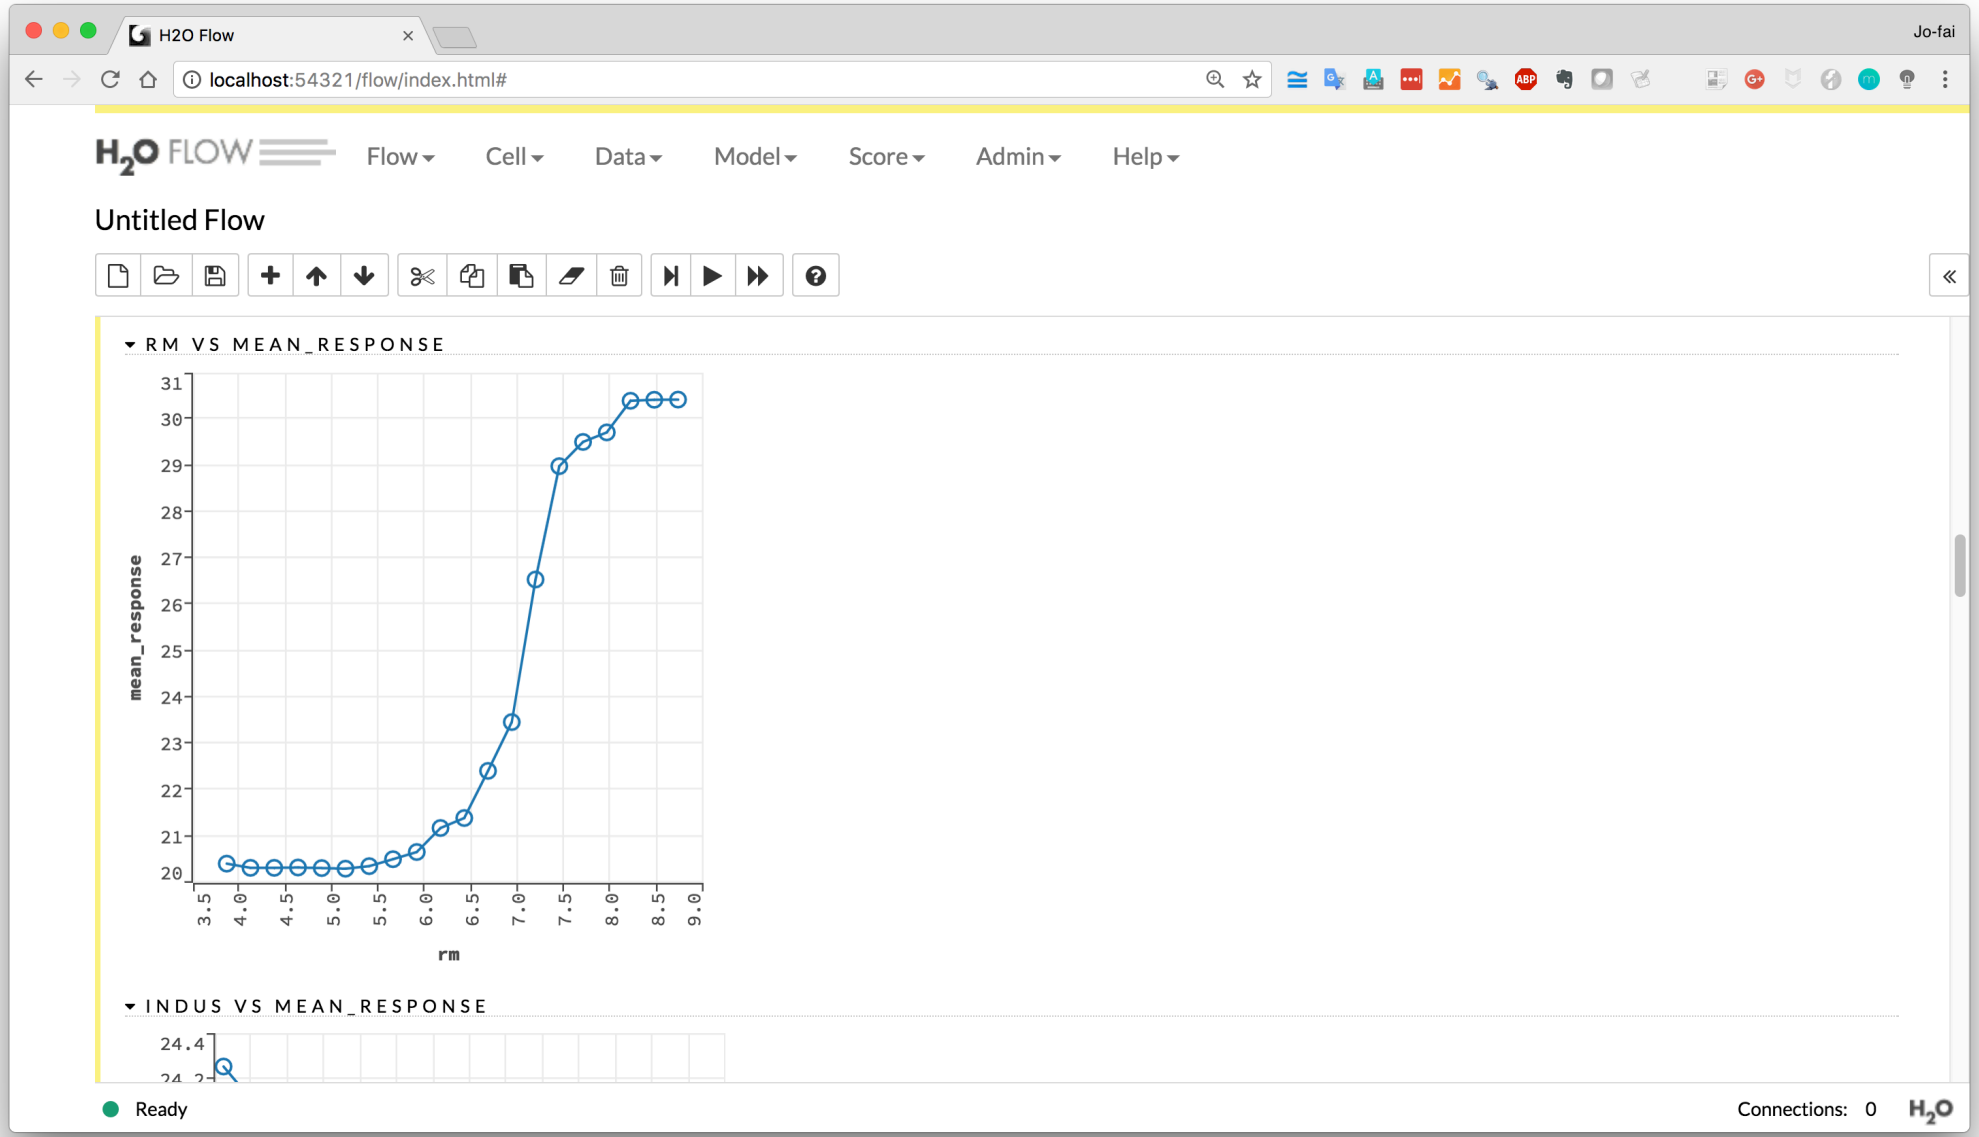
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
RM VS MEAN_RESPONSE
mean_response
31
30
29
28
27
26
25
24
23
22
21
20
3.5
4.0
4.5
5.0
5.5
6.0
6.5
7.0
7.5
8.0
8.5
9.0
rm
INDUS VS MEAN_RESPONSE
24.4
Ready
Connections: 0
H2O

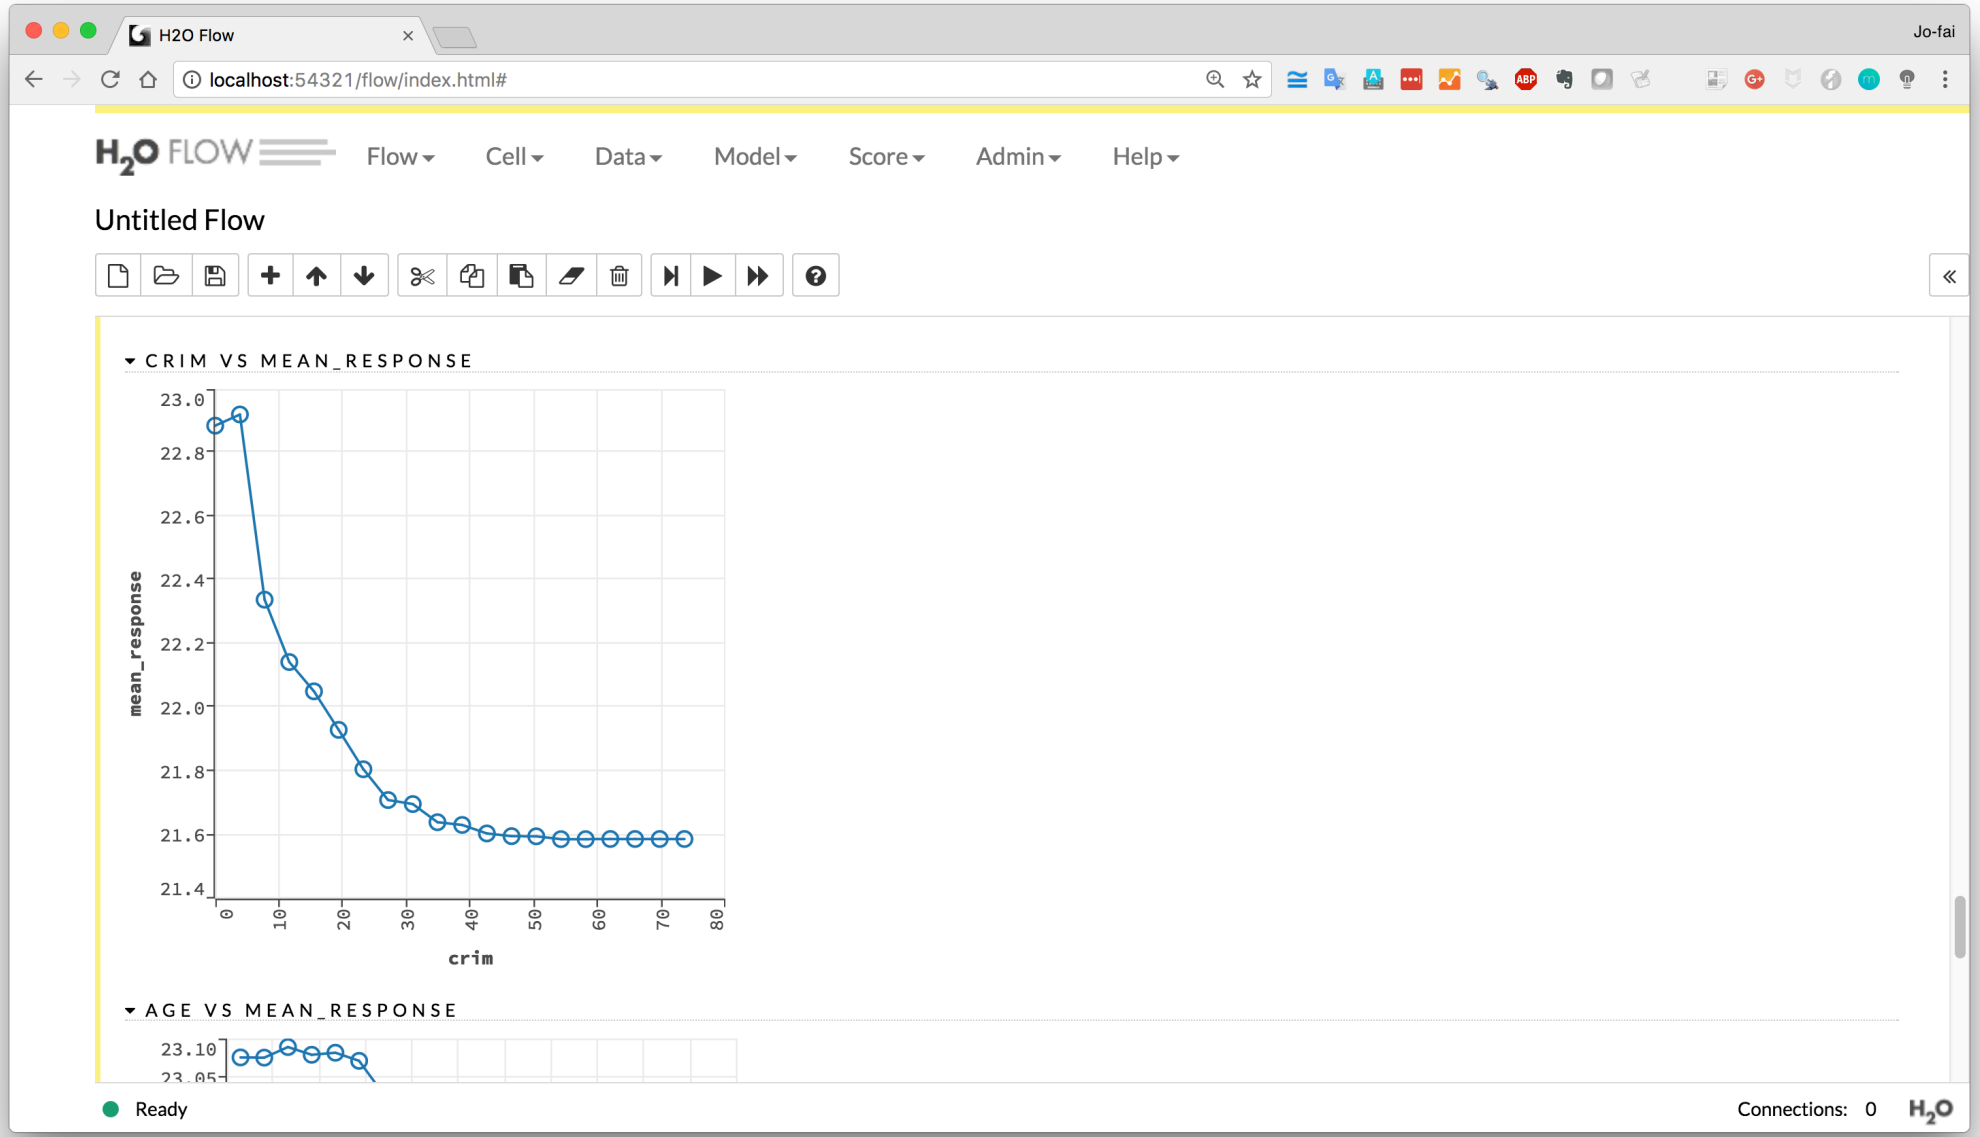
H2O Flow
Jo-fai
localhost:54321/flow/index.html#
H2O FLOW
Flow
Cell
Data
Model
Score
Admin
Help
Untitled Flow
CRIM VS MEAN_RESPONSE
mean_response
23.0
22.8
22.6
22.4
22.2
22.0
21.8
21.6
21.4
0
10
20
30
40
50
60
70
80
crim
AGE VS MEAN_RESPONSE
23.10
Ready
Connections: 0
H2O

Tutorial 2 - Classification

- **Data**: Human Activity Recognition Using Smartphone Sensors (2012)
- **Source**: https://archive.ics.uci.edu/ml/datasets/human+activity+recognition+using+smartphones

Tutorial 3 - Clustering

- **Data**: Water Treatment Plant (1993)
- **Source**: https://archive.ics.uci.edu/ml/datasets/water+treatment+plant

More About H_2O

- Running H_2O on Multi-Node Cluster
- Next-Gen H_2O Products

H_2O on Multi-Node Cluster

H_2O on Single Machine

H_2O on Multi-Node Cluster

H_2O with Distributed Data Storage Systems

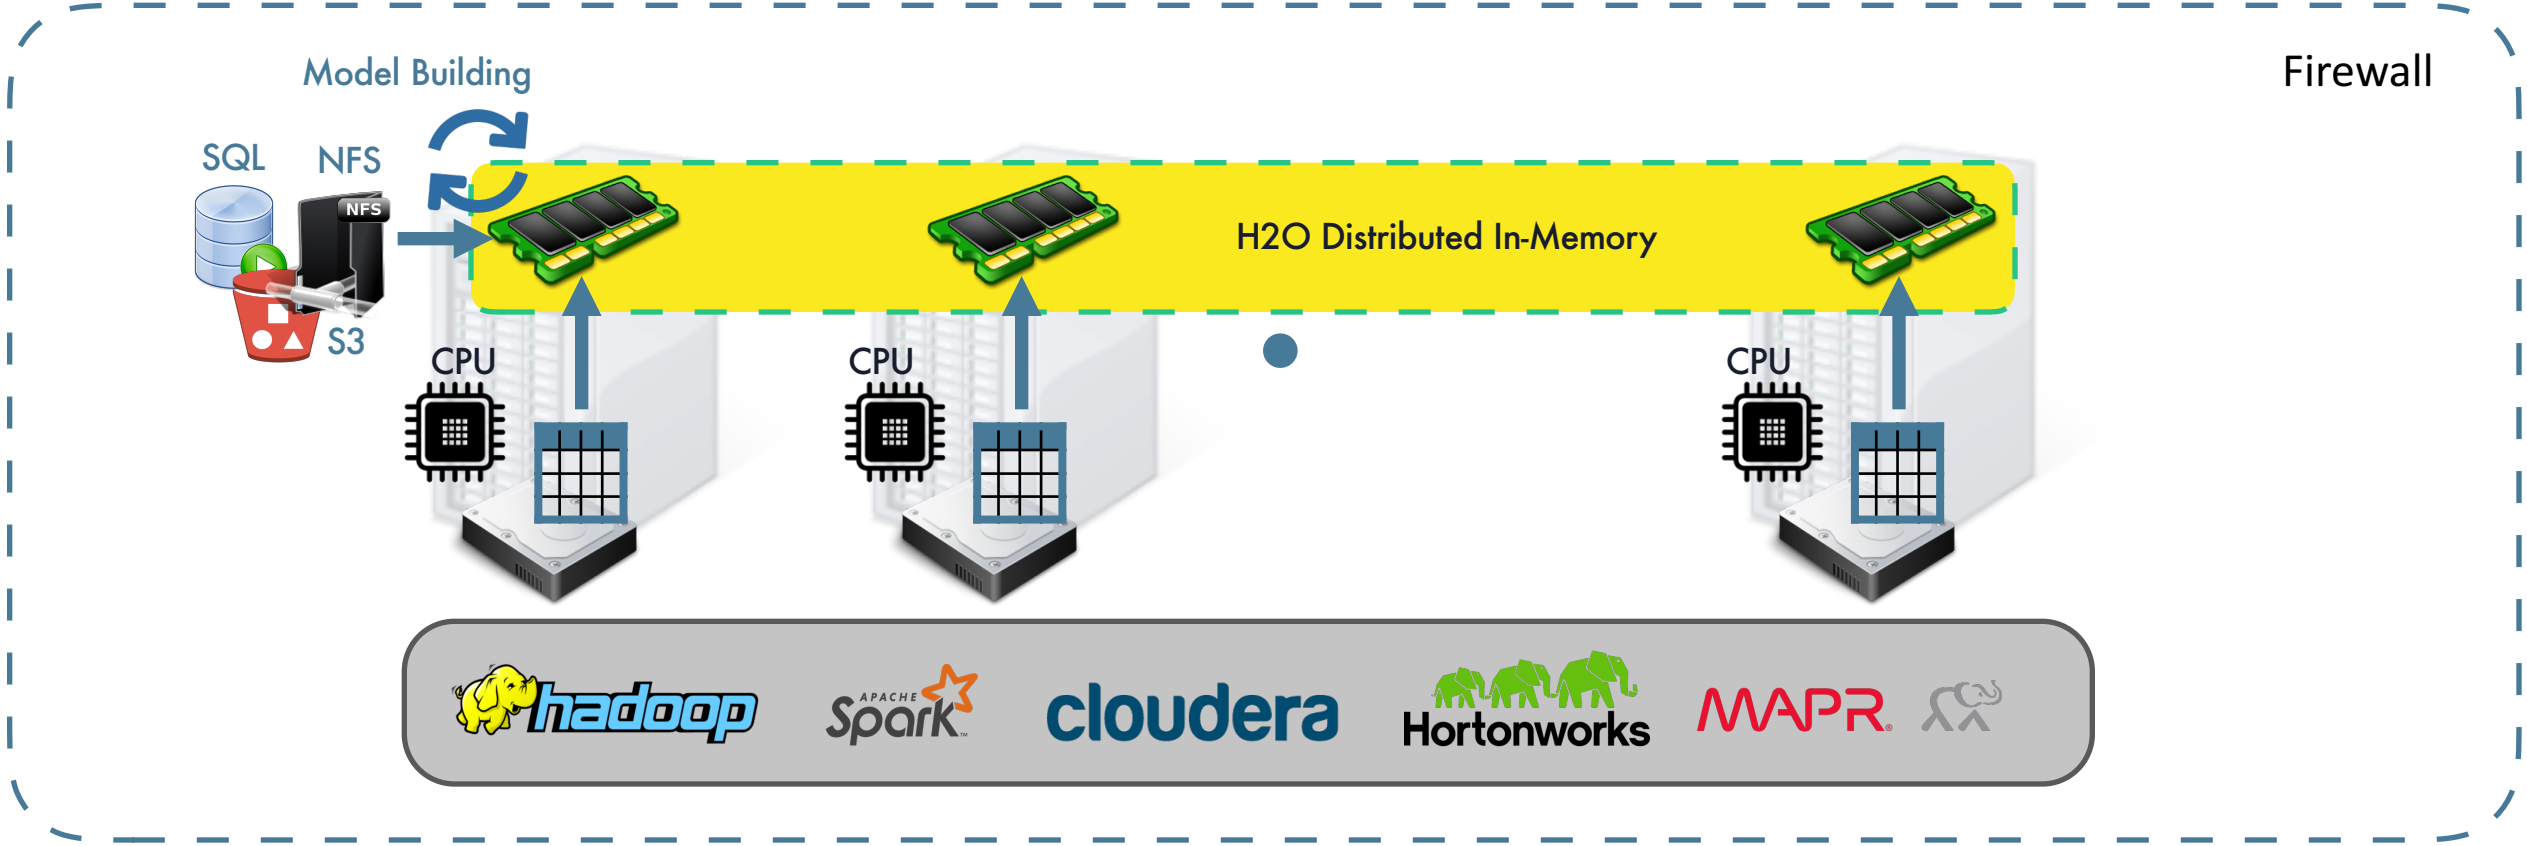

11M Rows

Size (Raw): **7.48 GB**
Compressed: **2.00 GB** (≈ 27% of Raw)

HIGGS.hex

Actions: View Data | Split... | Build Model... | Predict | Download | Export

Rows	Columns	Compressed Size
11000000	29	2GB

COLUMN SUMMARIES

label	type	Missing	Zeros	+Inf	-Inf	min	max	mean	sigma	cardinality	Actions
C1	enum	0	5170877	0	0	0	1.0	0.5299	0.4991	2	Convert to numeric
C2	real	0	0	0	0	0.2747	12.0989	0.9915	0.5654	·	·
C3	real	0	0	0	0	-2.4350	2.4349	-0.0	1.0088	·	·
C4	real	0	0	0	0	-1.7425	1.7432	-0.0	1.0063	·	·
C5	real	0	0	0	0	0.0002	15.3968	0.9985	0.6000	·	·
C6	real	0	0	0	0	-1.7439	1.7433	0.0	1.0063	·	·
C7	real	0	0	0	0	0.1375	9.9404	0.9909	0.4750	·	·
C8	real	0	0	0	0	-2.9697	2.9697	-0.0	1.0093	·	·
C9	real	0	0	0	0	-1.7412	1.7415	0.0	1.0059	·	·
C10	real	0	5394611	0	0	0	2.1731	1.0	1.0278	·	·
C11	real	0	0	0	0	0.1890	11.6471	0.9927	0.5000	·	·
C12	real	0	0	0	0	-2.9131	2.9132	-0.0	1.0093	·	·
C13	real	0	0	0	0	-1.7424	1.7432	-0.0	1.0062	·	·
C14	real	0	5523912	0	0	0	2.2149	1.0	1.0494	·	·
C15	real	0	0	0	0	0.2636	14.7090	0.9923	0.4877	·	·
C16	real	0	0	0	0	-2.7297	2.7300	0.0	1.0087	·	·
C17	real	0	0	0	0	-1.7421	1.7429	0.0	1.0063	·	·
C18	real	0	6265240	0	0	0	2.5482	1.0	1.1937	·	·
C19	real	0	0	0	0	0.3654	12.8826	0.9861	0.5058	·	·
C20	real	0	0	0	0	-2.4973	2.4980	-0.0	1.0077	·	·

H2O.ai

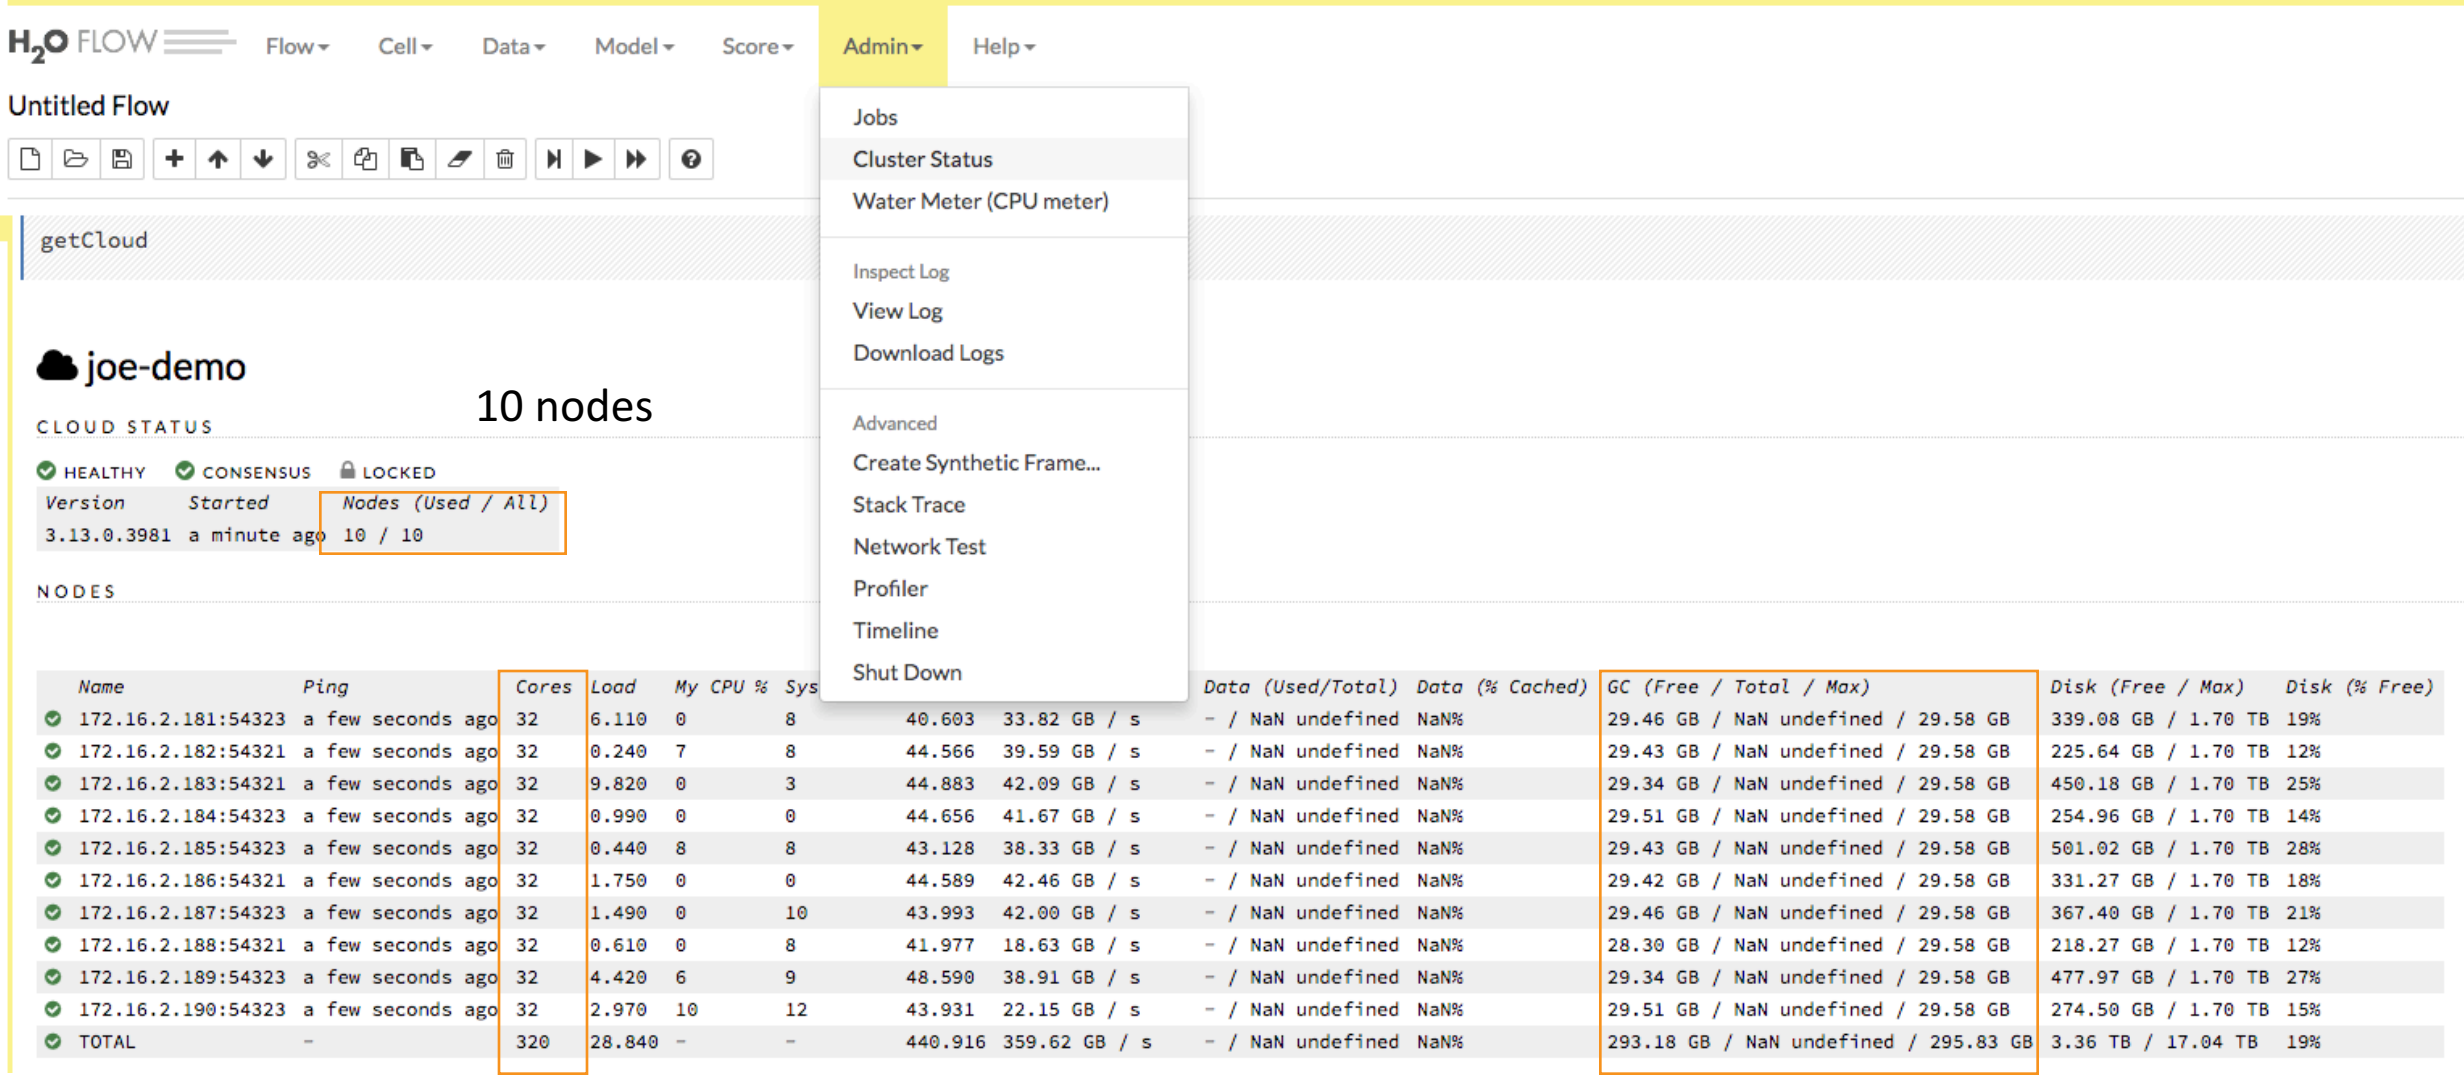

10 x 32 =
320 Cores

10 x 29.6 = 296
GB Memory

Higgs Boson Machine Learning Challenge

Use the ATLAS experiment to identify the Higgs boson

$13,000 · 1,785 teams · 3 years ago

Overview | Data | Discussion | Leaderboard | Rules | Team | My Submissions | Late Submission

https://www.kaggle.com/c/higgs-boson

Overview

- Description
- Evaluation
- Prizes
- About The Sponsors
- Timeline
- Winners

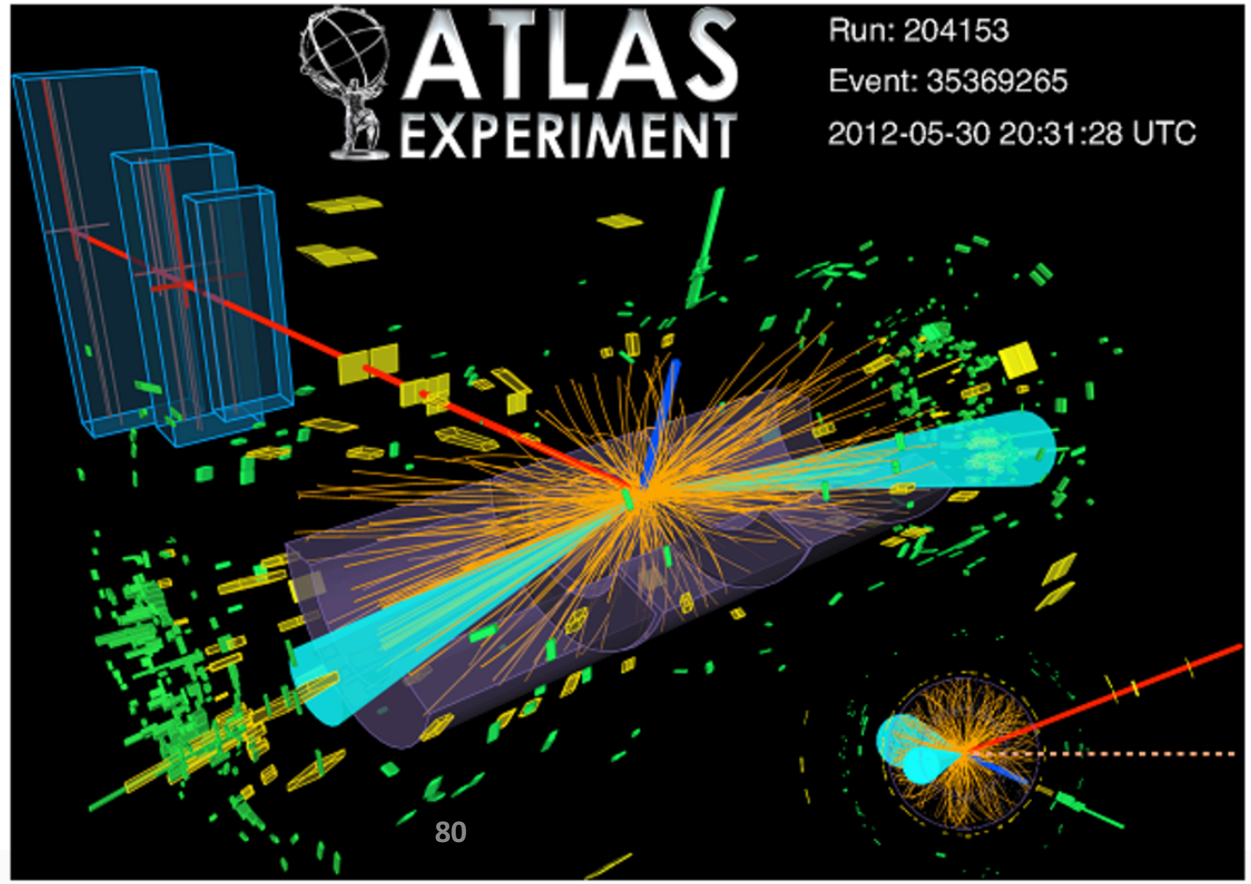

Learning from Higgs Boson Dataset

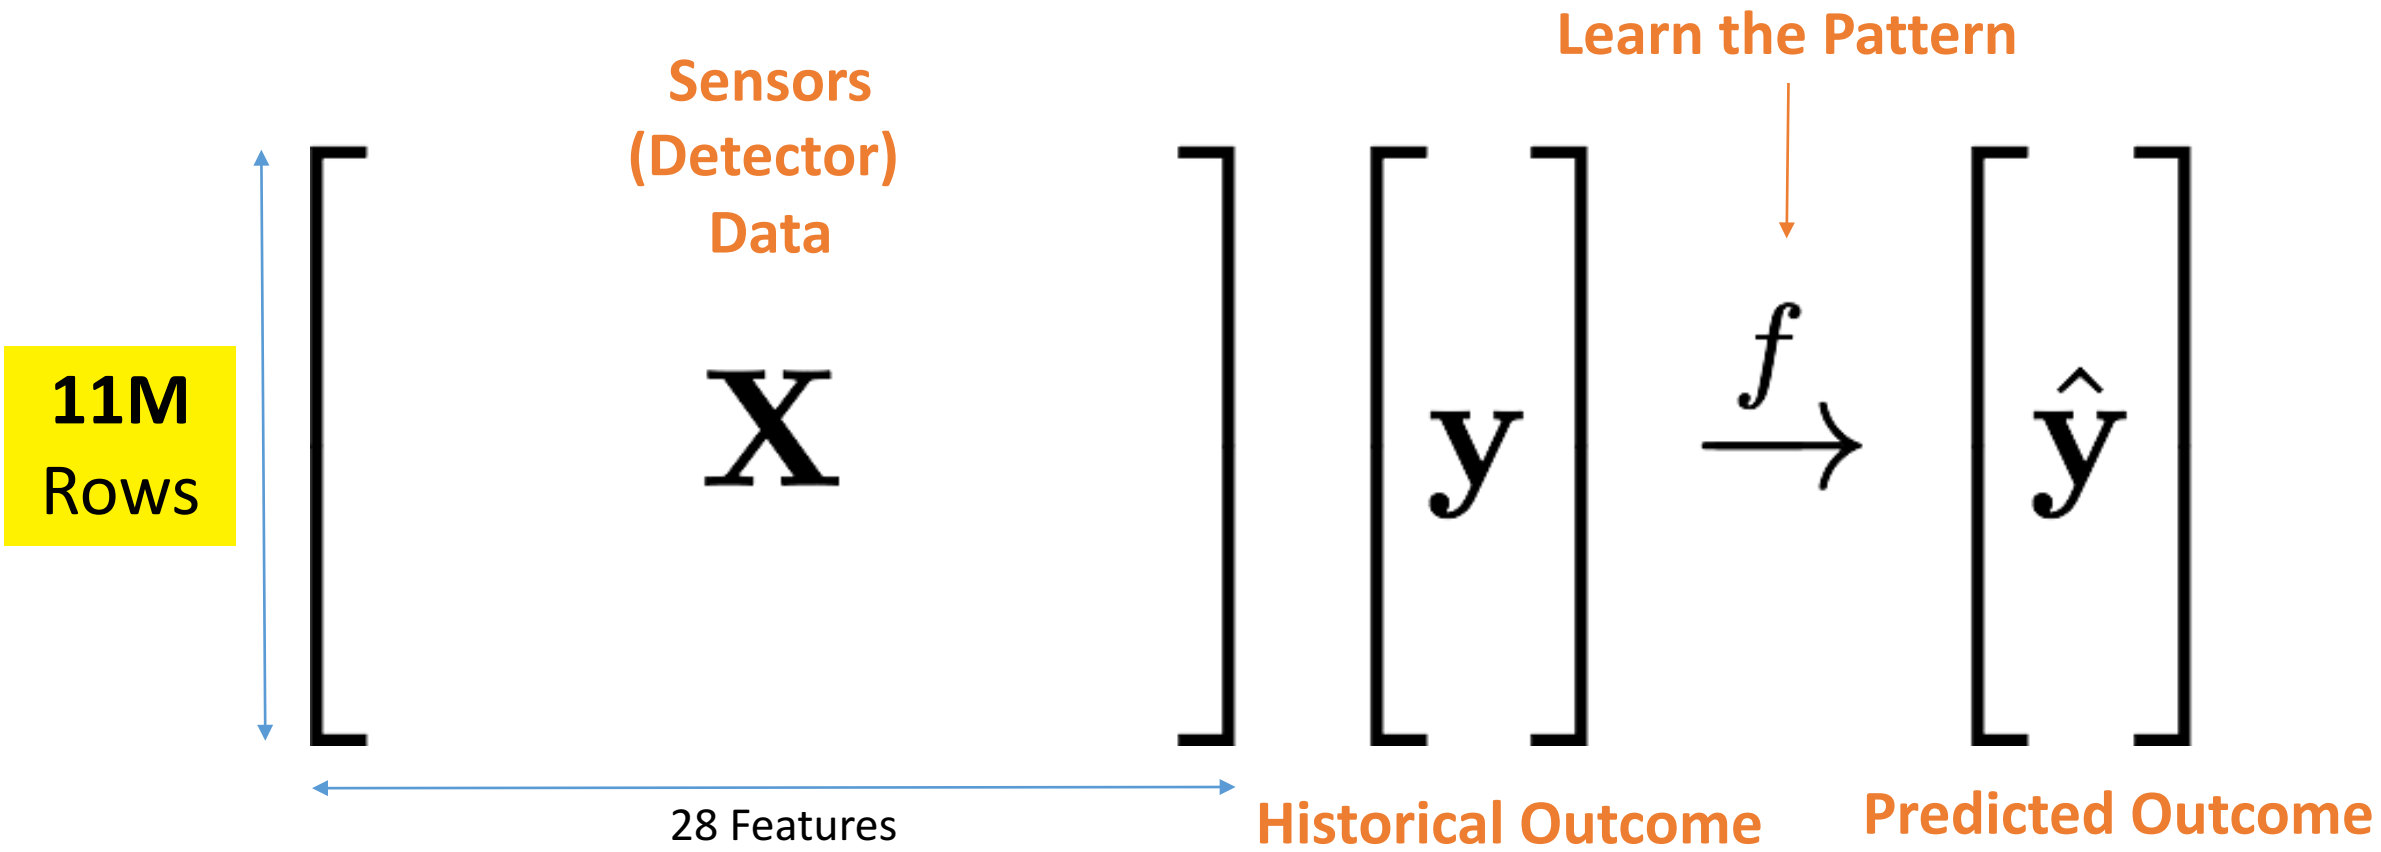

Raw Data Size: **7.48 GB**

H_2O Water Meter (CPU Usage Monitor)

10 x 32 = 320 Cores

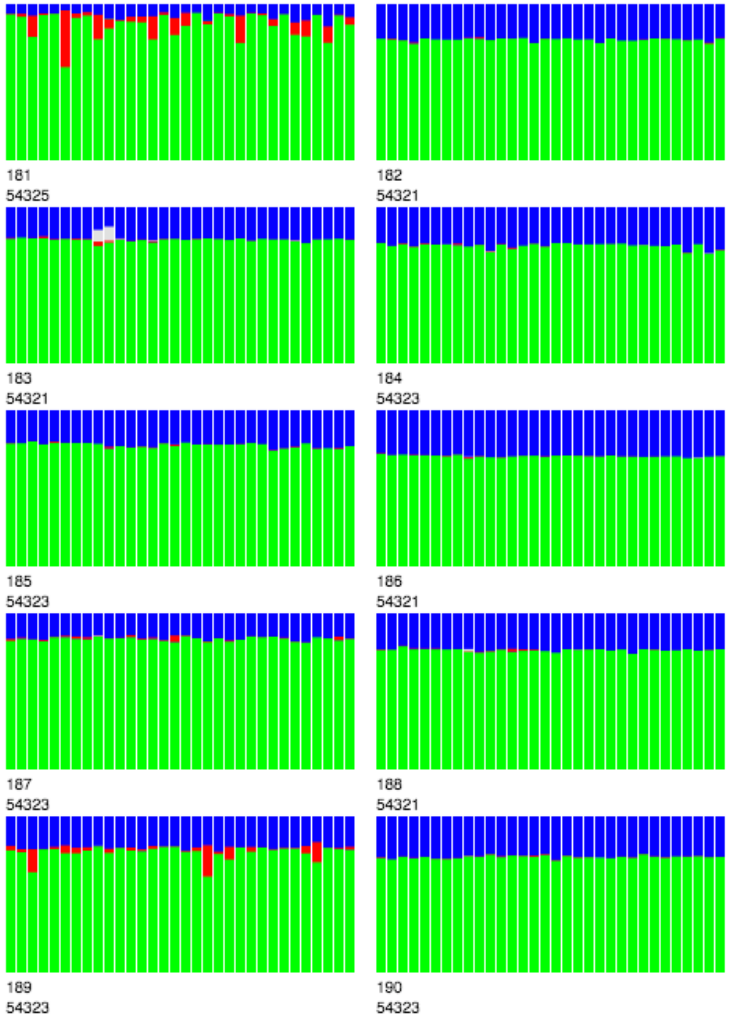

H_2O4GPU

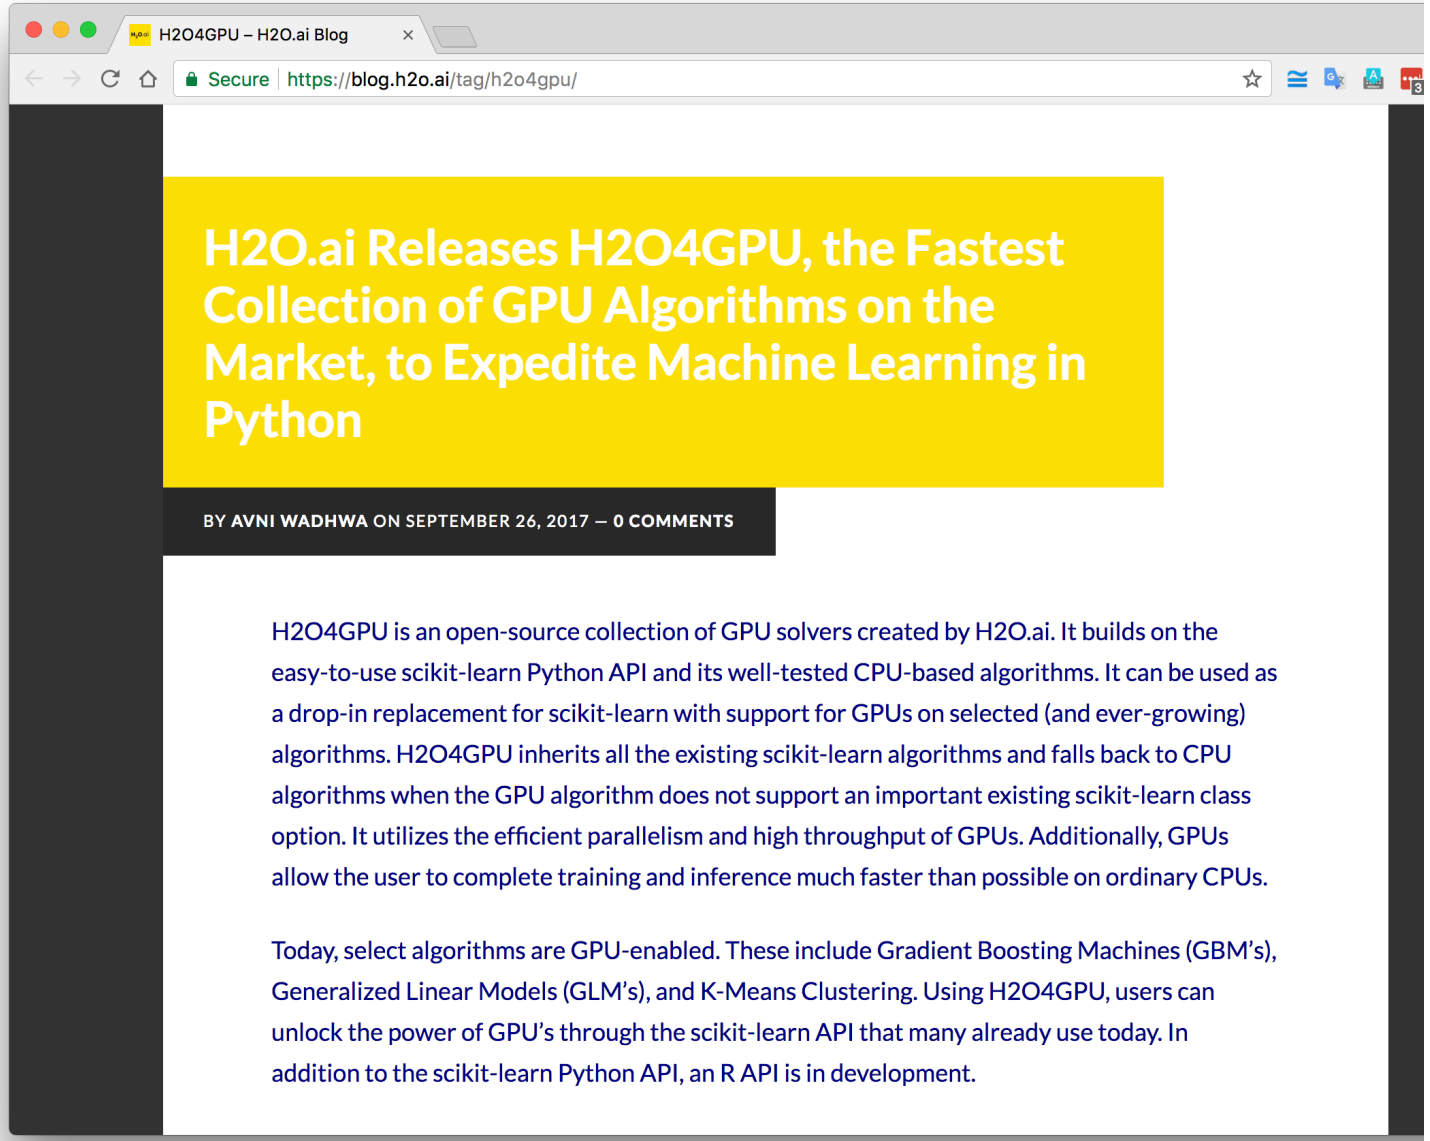

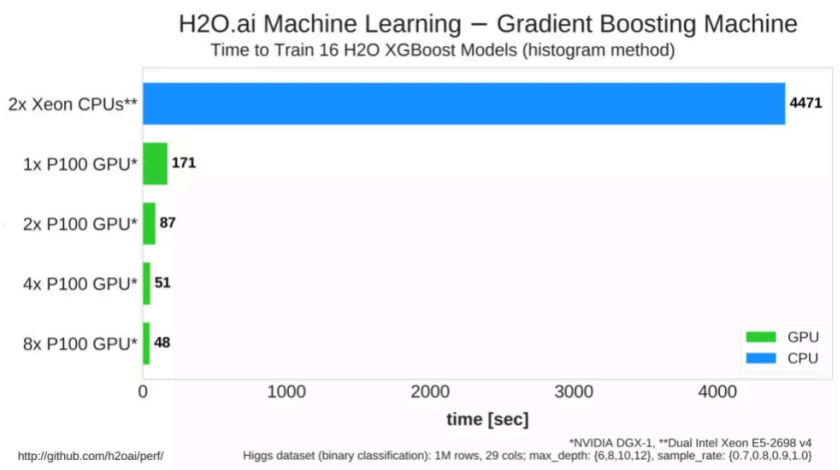

k-Means Clustering

- Based on NVIDIA prototype of k-Means algorithm in CUDA
- Improvements to original implementation:
  - Significantly faster than scikit-learn implementation (50x) and other GPU implementations (5-10x)
  - Supports multiple GPUs

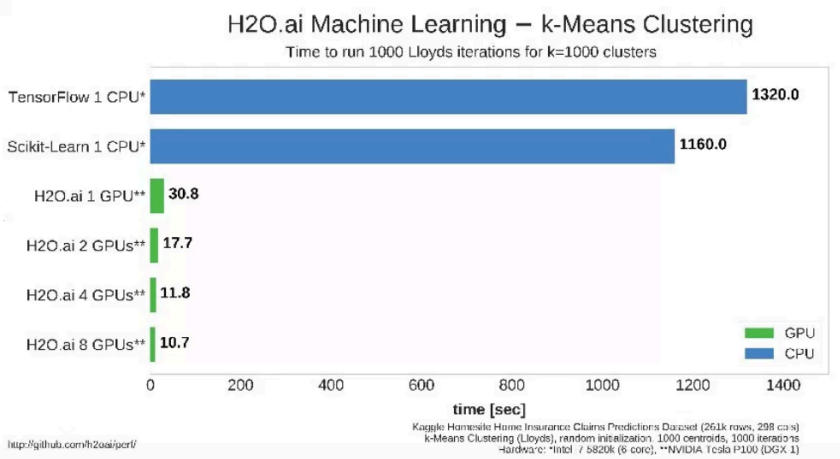

https://blog.h2o.ai/tag/h2o4gpu/

Driverless AI

H2O.ai and NVIDIA Bring Fast, ×

Jo-fai

www.businesswire.com/news/home/20170926006769/en/H2O.ai-NVIDIA-Bring-Fast-Accurate-Interpretable...

BusinessWire
A Berkshire Hathaway Company

HOME SERVICES NEWS EDUCATION ABOUT US

Search

Log In Sign Up

H2O.ai and NVIDIA Bring Fast, Accurate and Interpretable Driverless AI with Automated Machine Learning and Feature Engineering

Businesses Can Leapfrog Data Scientist Shortage and Accelerate Adoption of AI with H2O.ai Driverless AI on NVIDIA DGX Systems

September 26, 2017 04:57 PM Eastern Daylight Time

NEW YORK--(BUSINESS WIRE)--**Strata Data Conference —** H2O.ai today announced an offering built on NVIDIA DGX Systems to further democratize machine learning and address the growing demands placed on a limited number of trained data scientists.

"Interpretable AI builds trust in AI and automates report generation for regulatory purposes. Driverless AI on GPU Computing Platforms accelerates learning and data products with AI in enterprises."

Tweet this

Traditional approaches to machine learning require organizations to rely on highly skilled data scientists with years of domain knowledge. These experts continuously experiment to discover features that make machine learning algorithms more accurate and interpretable for business teams to consume. With most of the world's elite data scientists flocking to Software AI giants, there is a vacuum in the diverse talent needed to transform businesses with AI.

H2O.ai's Driverless AI platform, fully integrated on NVIDIA DGX systems, lets business users, analysts and data scientists use an incredibly fast, intuitive computing platform. Now customers can apply Automatic Feature Engineering and quickly develop hundreds of machine learning models to help businesses mitigate risks and maximize revenue

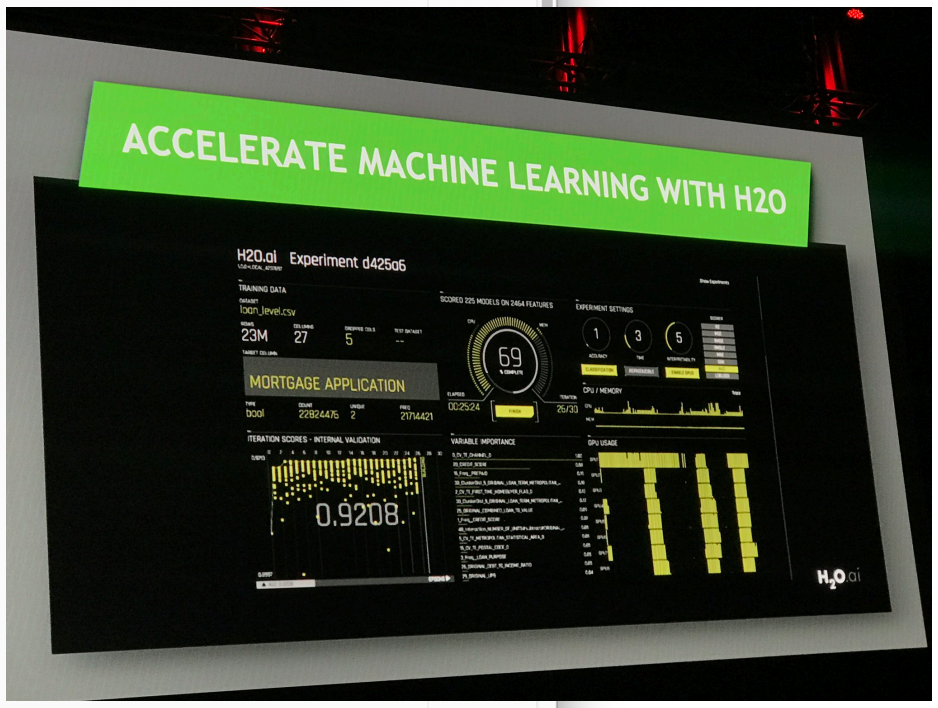

DK0KfLaUMAEp....jpg-large

Show All

http://www.businesswire.com/news/home/20170926006769/en/H2O.ai-NVIDIA-Bring-Fast-Accurate-Interpretable-Driverless

H2O.ai

Thank you!

- Prof. Dragan Savić
- Prof. Zoran Kapelan

- Code, Slides & Documents
  - bit.ly/h2o_meetups
  - docs.h2o.ai
- Contact
  - joe@h2o.ai
  - @matlabulous
  - github.com/woobe
- Please search/ask questions on **Stack Overflow**
  - Use the tag `**h2o**` (not H2 zero)